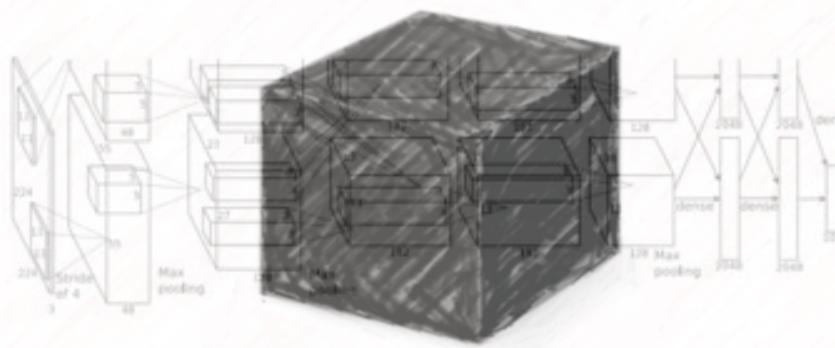


Last lesson

recap, final thought and final project



Input



Deep neural network

Dog

Output

$$\mathbf{x} \longrightarrow f(\mathbf{x}; \mathbf{W}) \longrightarrow \hat{\mathbf{y}}$$

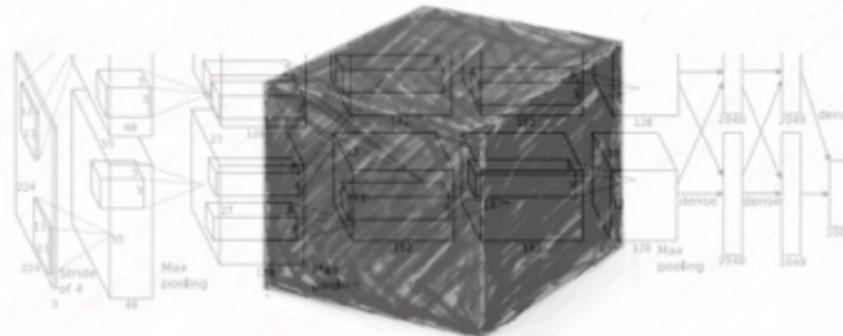
Optimization problem:

$$\min_{\mathbf{W}} \left[\sum_i loss(\hat{y}_i, y_i) \right] \longrightarrow \mathbf{W}_{t+1} = \mathbf{W}_t - \alpha \cdot \frac{\partial \sum_i loss(\hat{y}_i, y_i)}{\partial \mathbf{W}_t}$$

Solved using gradient descent:



Input



Deep neural network

Dog

Output

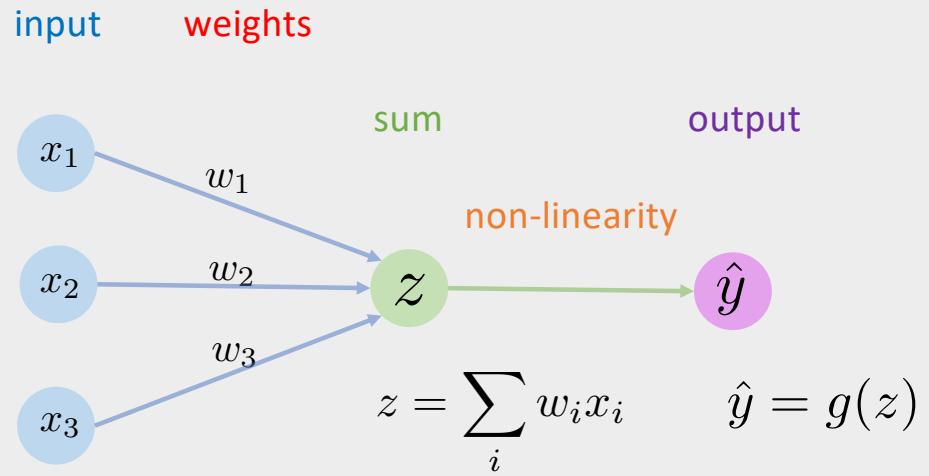
desired
output

\mathbf{x}

$f(\mathbf{x} ; \mathbf{W})$

$\hat{y} \leftarrow \text{loss} \rightarrow y$

Reminder: The Perceptron



output

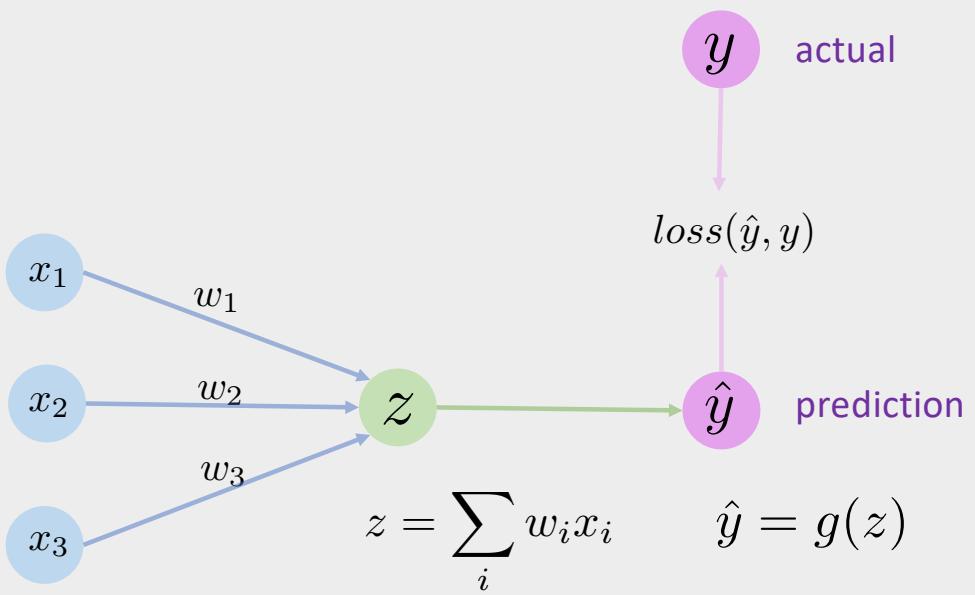
$$\hat{y} = g \left(\sum_i w_i x_i \right)$$

weights

input

non-linearity

sum



$$g(z_i) = \frac{1}{1 + e^{-z_i}}$$

$$l(y, \hat{y}) = \begin{cases} -\log(\hat{y}) & y = 1 \\ -\log(1 - \hat{y}) & y = 0 \end{cases}$$

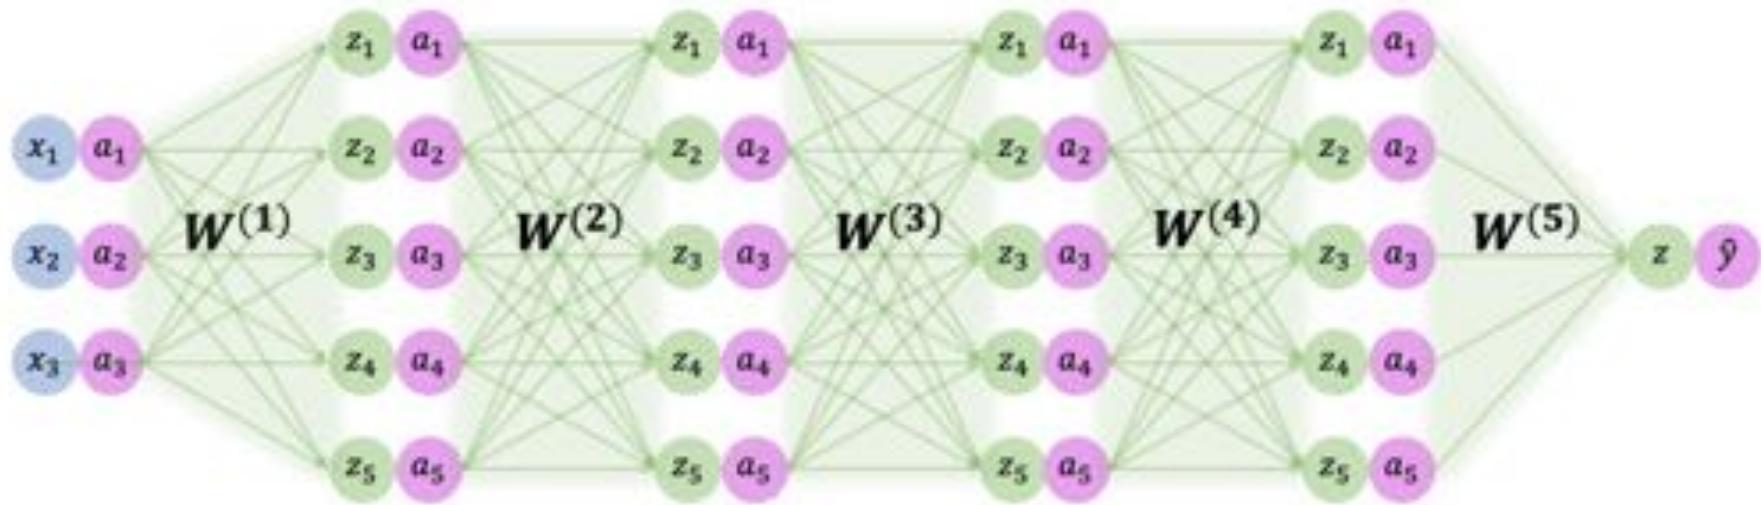
$$J(\vec{w}) = \sum_i l(y_i, \hat{y}_i)$$

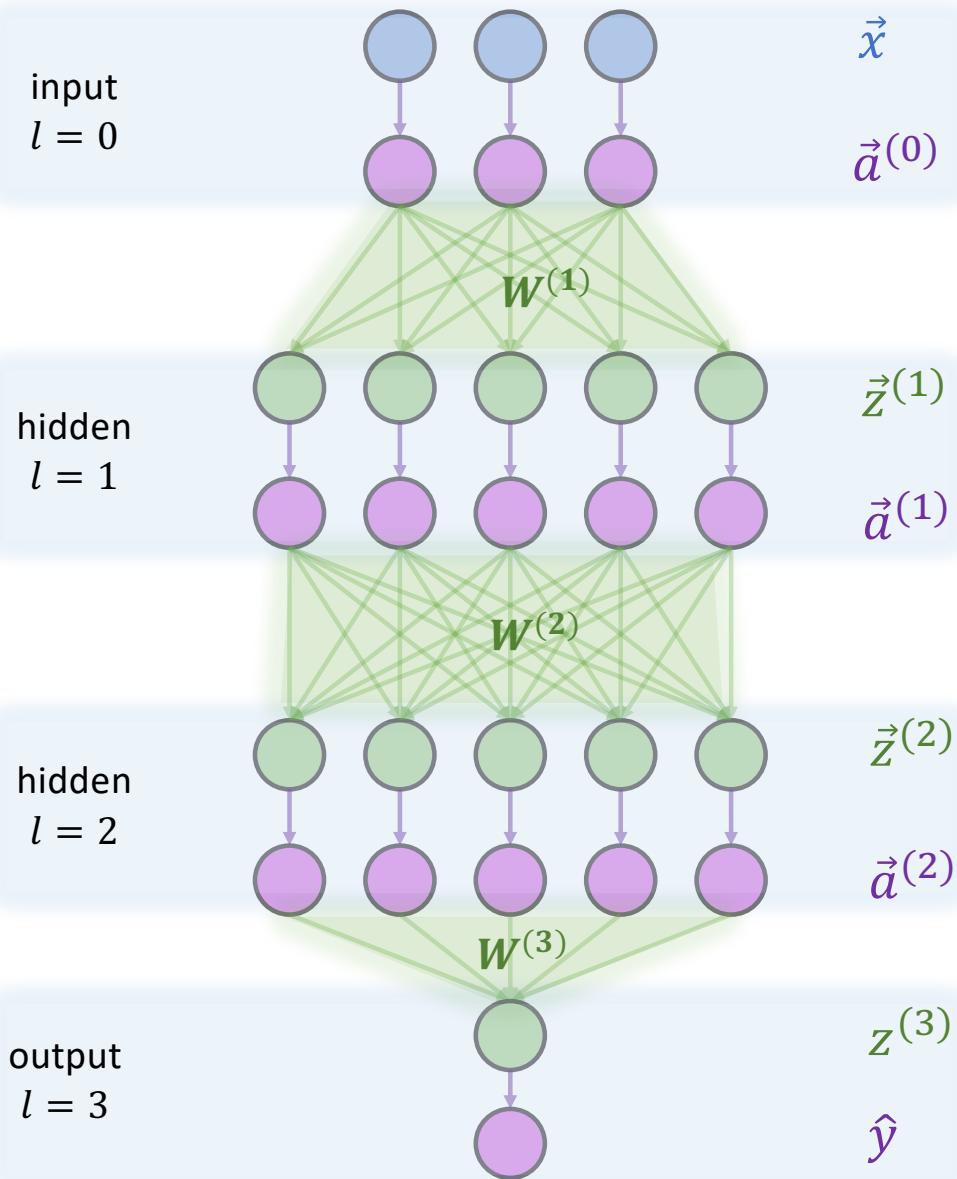
$$\min_{\vec{w}} [J(\vec{w})]$$

Gradient Descent
→

$$\vec{w} := \vec{w} - \alpha \frac{\partial J}{\partial \vec{w}}$$

Multilayer Perceptron (fully connected / dense layers)





Forward Propagation

$$\vec{a}^{(0)} = \vec{x}$$

$$\vec{z}^{(1)} = \vec{a}^{(0)} \cdot W^{(1)}$$

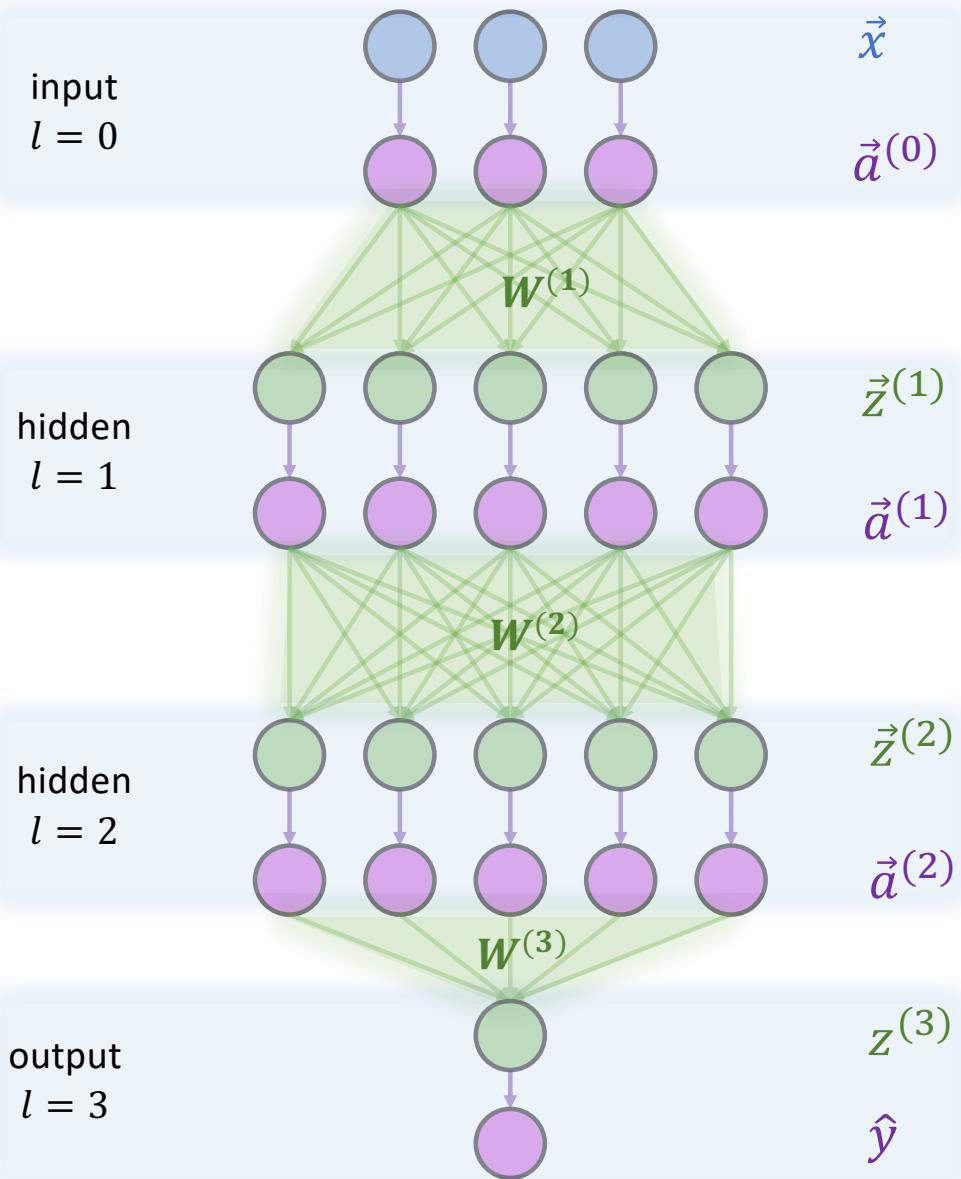
$$a_i^{(1)} = g(z_i^{(1)})$$

$$\vec{z}^{(2)} = \vec{a}^{(1)} \cdot W^{(2)}$$

$$a_i^{(2)} = g(z_i^{(2)})$$

$$\vec{z}^{(3)} = \vec{a}^{(2)} \cdot W^{(3)}$$

$$\hat{y} = g(z^{(3)})$$



Back Propagation

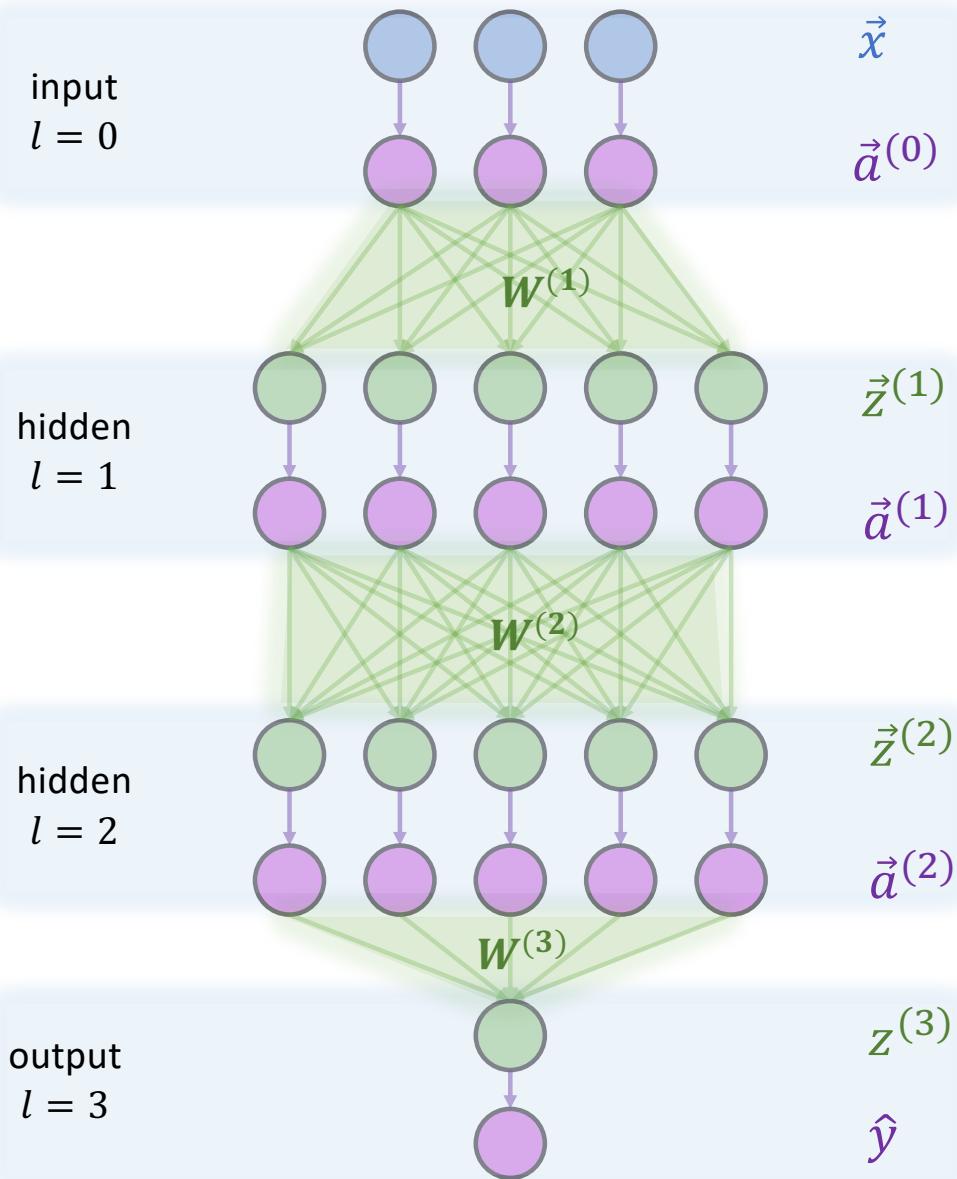
input: y , \hat{y} , $\{W^{(l)}\}$, $\{\vec{z}^{(l)}\}$

$$\vec{\delta}^{(l_{max})} = \hat{y} - y$$

for $l = l_{max} - 1$ to 1:

$$\vec{\delta}^{(l-1)} = \vec{\delta}^{(l)} \cdot W^{T(l)} \odot g'(\vec{z}^{(l-1)})$$

return: $\{\vec{\delta}^{(l)}\}$



Gradient Descent (1-step)

input: $\{W^{(l)}\}$, $\{\vec{a}^{(l)}\}$, $\{\vec{\delta}^{(l)}\}$

for $l = 1$ **to** l_{max} :

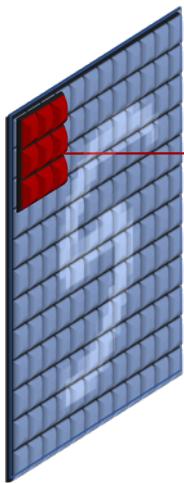
$$\frac{\partial J}{\partial W^{(l)}} = \vec{a}^{T(l-1)} \cdot \vec{\delta}^{(l)}$$

$$W^{(l)} = W^{(l)} - \alpha \cdot \frac{\partial J}{\partial W^{(l)}}$$

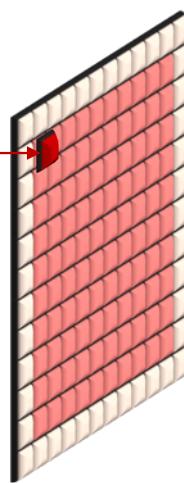


3x3 kernel

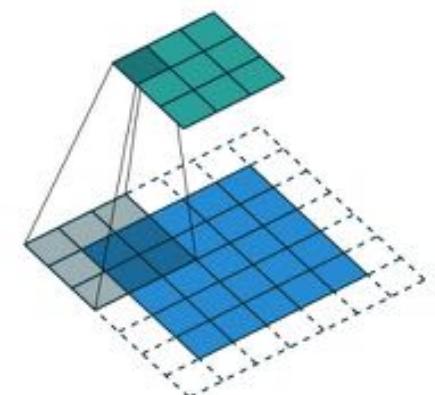
Let's connect each neuron to only a local region (3x3) of the input volume:

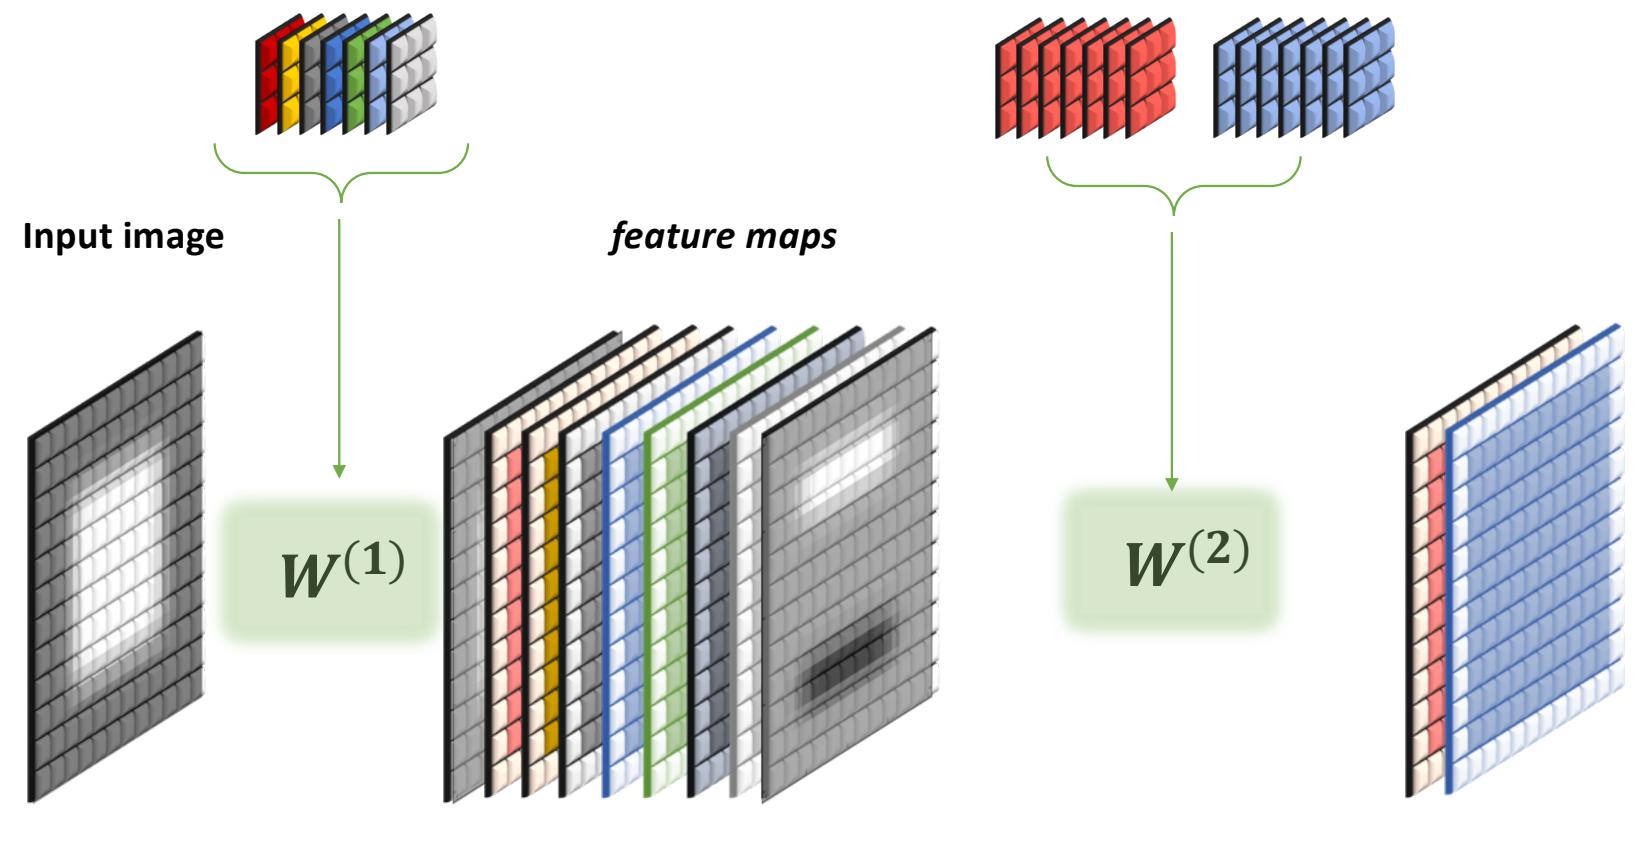


12 x 12



10 x 10





Input layer

1st conv layer

2nd conv layer

$12 \times 12 \times 1$

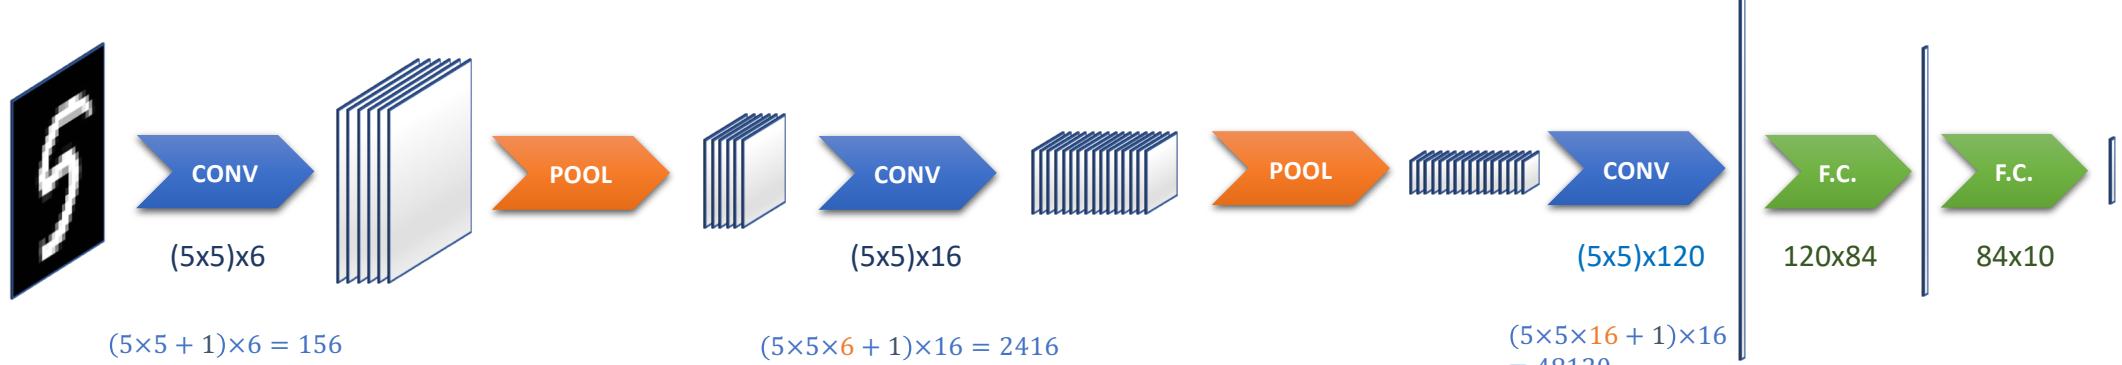
(depth of 7)

$10 \times 10 \times 7$

(depth of 2)

$8 \times 8 \times 2$

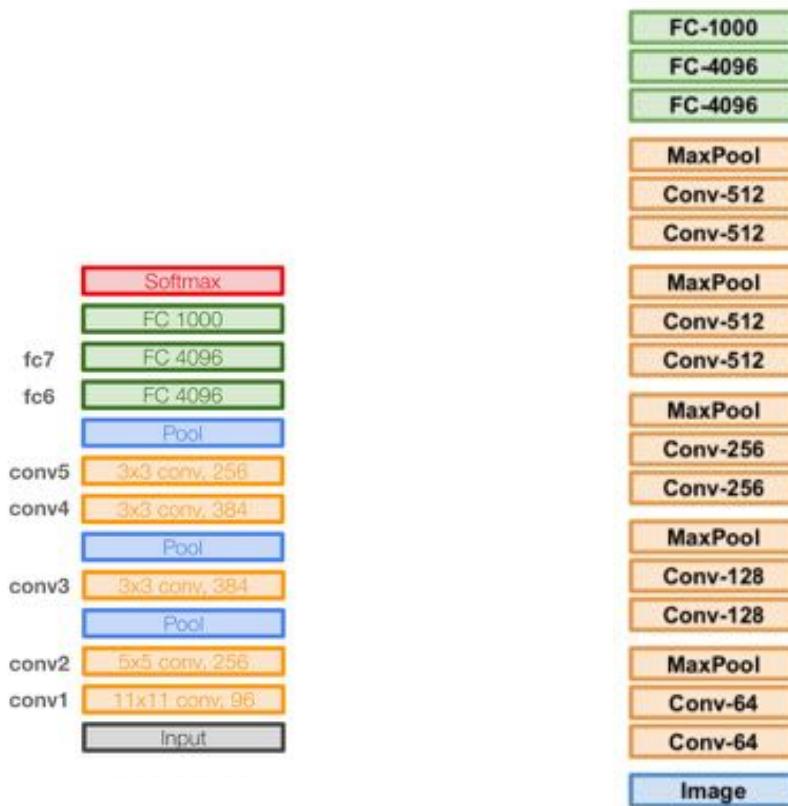
LeNet-5 (LeCun et al., 1998)



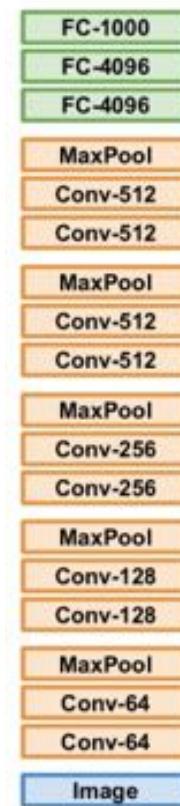
```
# LeNet-5 Model
model = Sequential()
model.add(Conv2D(6, kernel_size=(5, 5),
                 activation='relu',
                 input_shape=(32,32,1)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(16, kernel_size=(5, 5),
                 activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(120, kernel_size=(5, 5),
                 activation='relu'))
model.add(Flatten())
model.add(Dense(84, activation='relu'))
model.add(Dense(10, activation='softmax'))
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_19 (Conv2D)	(None, 28, 28, 6)	156
max_pooling2d_11 (MaxPooling)	(None, 14, 14, 6)	0
conv2d_20 (Conv2D)	(None, 10, 10, 16)	2416
max_pooling2d_12 (MaxPooling)	(None, 5, 5, 16)	0
conv2d_21 (Conv2D)	(None, 1, 1, 120)	48120
flatten_7 (Flatten)	(None, 120)	0
dense_13 (Dense)	(None, 84)	10164
dense_14 (Dense)	(None, 10)	850
=====		
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		

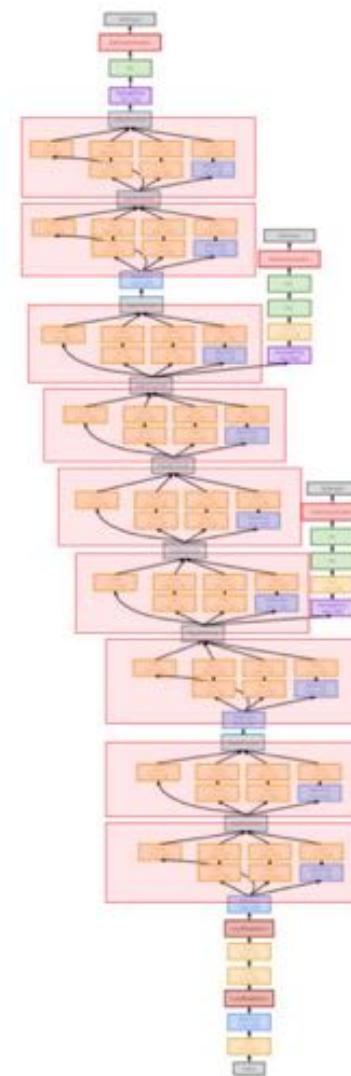
AlexNet



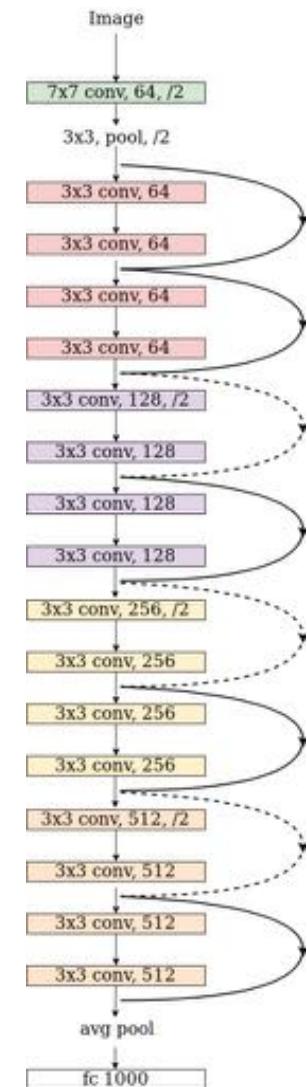
VGG

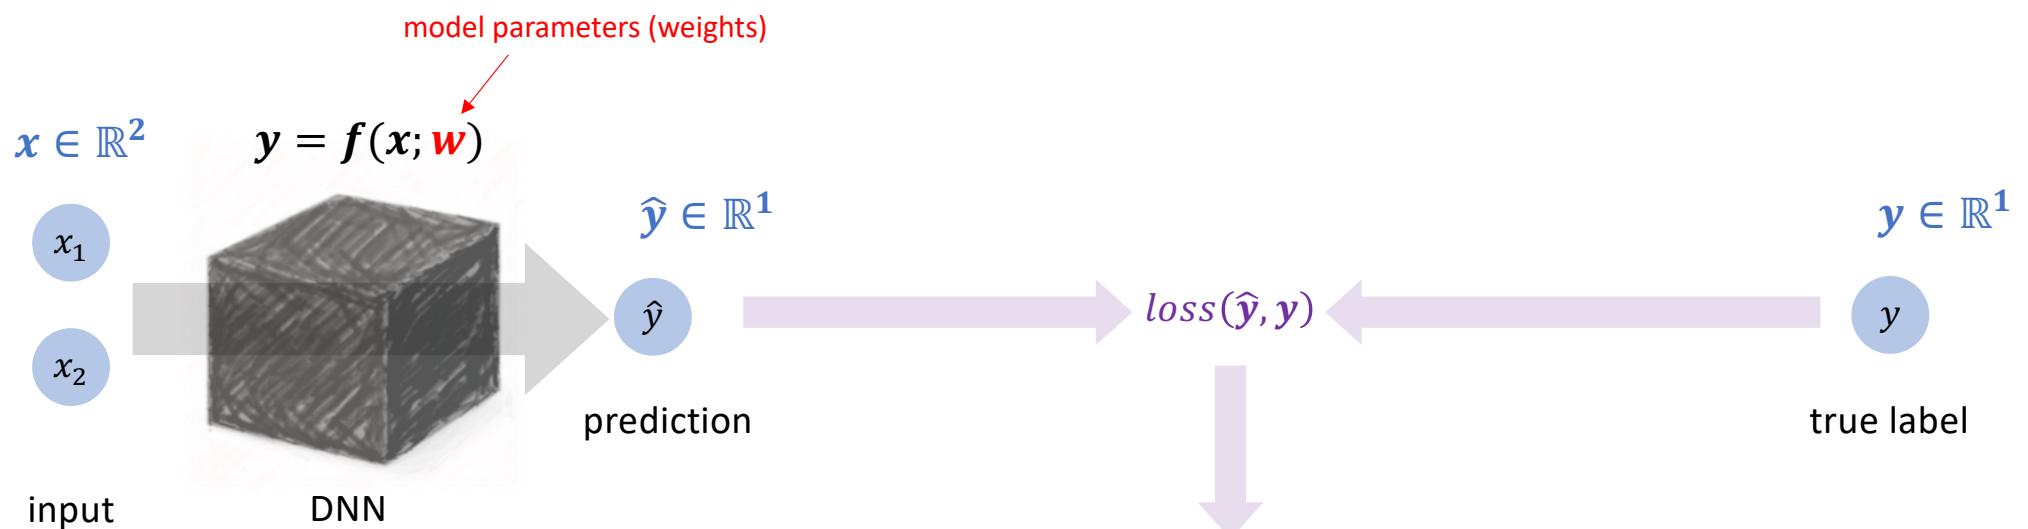


Inception



ResNet

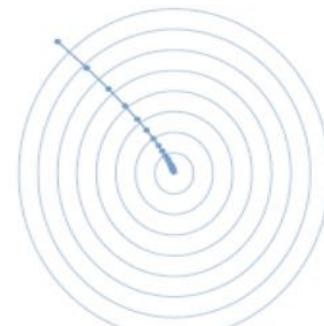




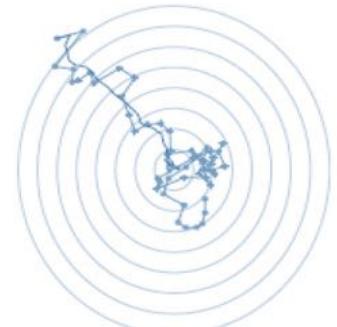
Optimization problem: $\mathbf{w}^* = \arg \min_{\mathbf{w}} \left[\sum_i loss(\hat{y}_i, y_i) \right]$

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \cdot \frac{\partial J(\mathbf{w}_t)}{\partial \mathbf{w}_t}$$

Gradient descent algorithm

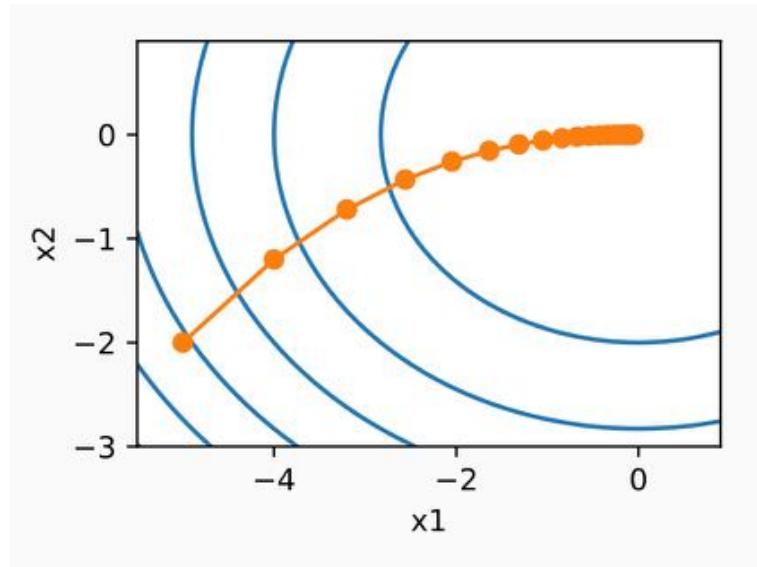


gradient-descent

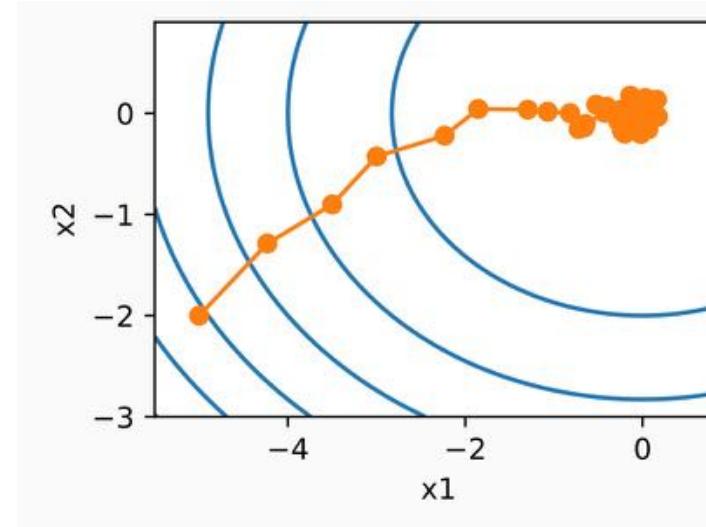


Stochastic gradient-descent

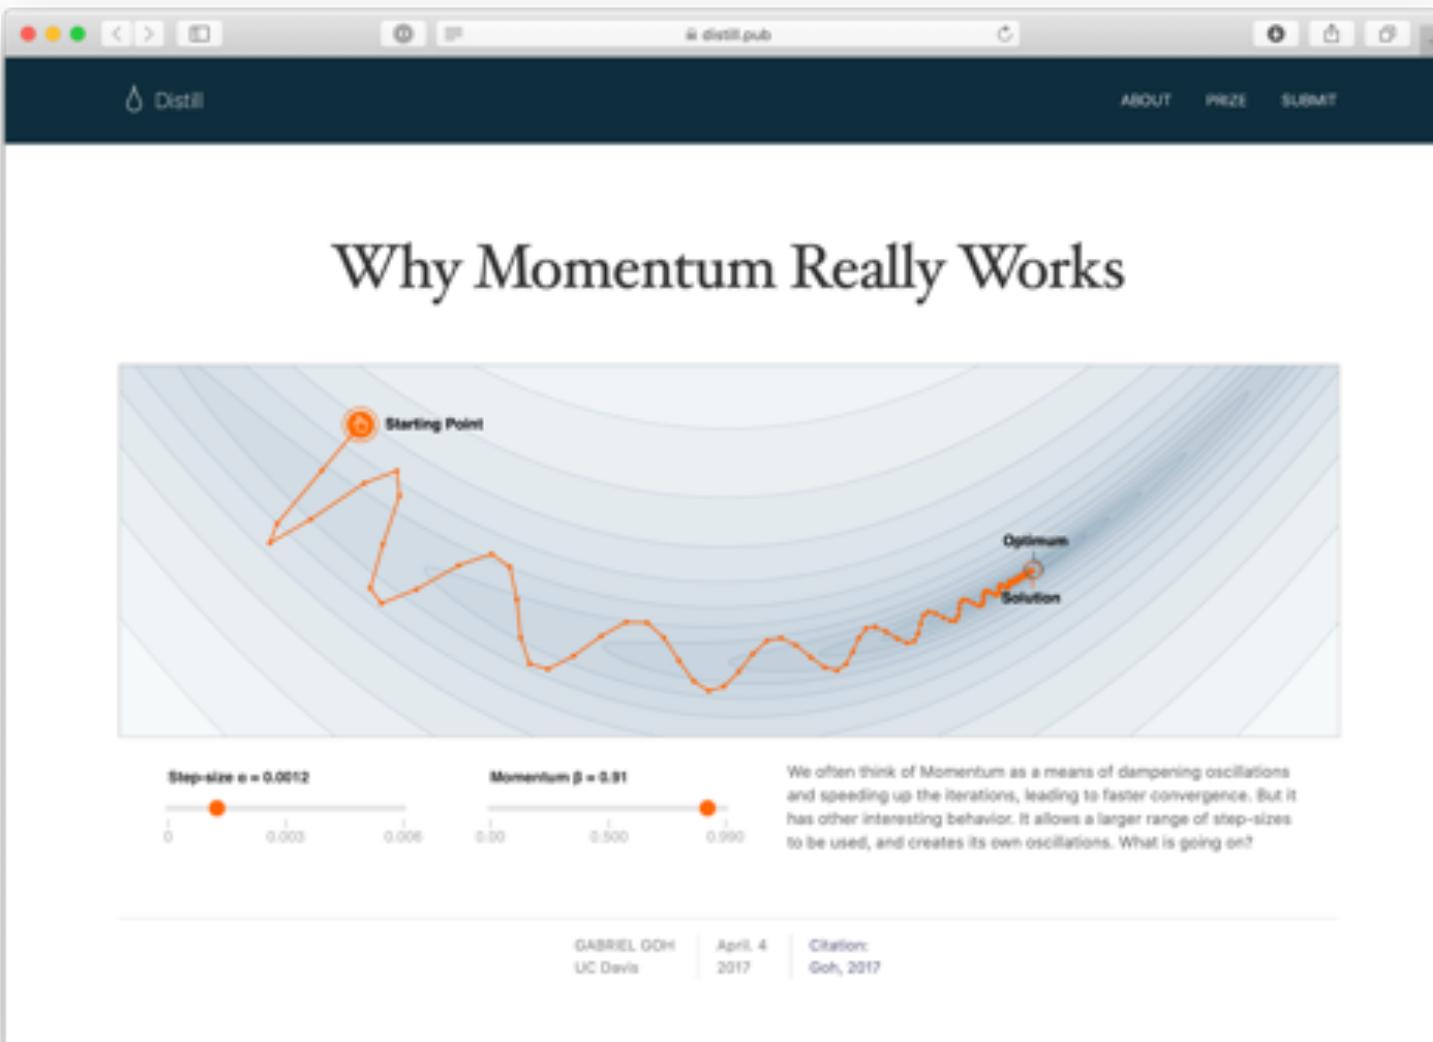
Reminder: 2-D Gradient Descent



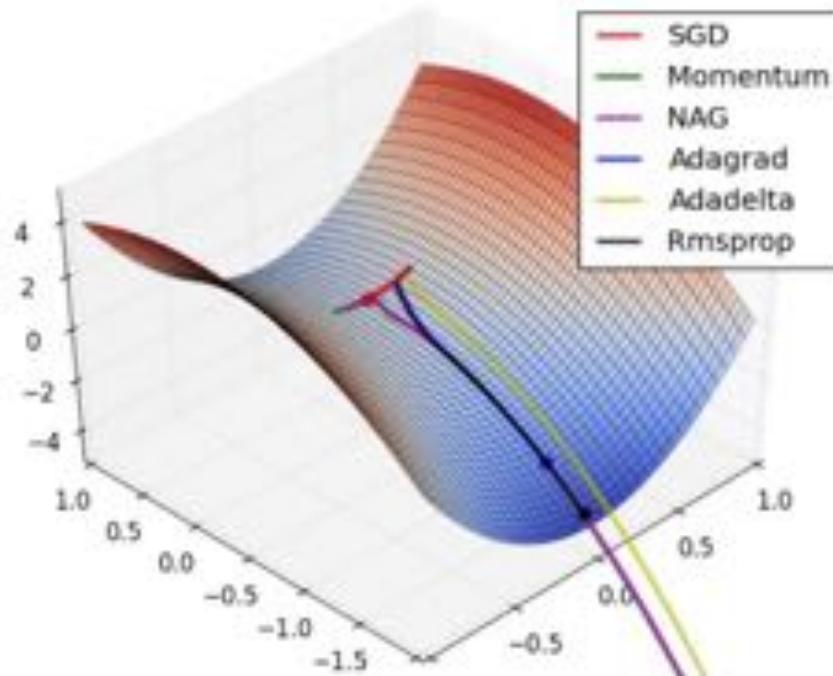
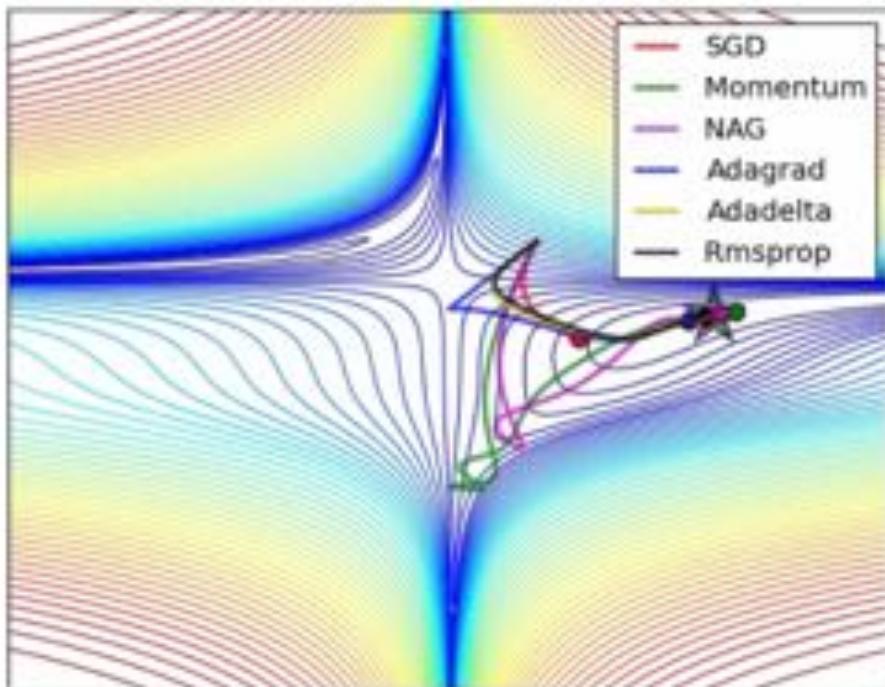
Gradient Descent



Stochastic Gradient Descent (SGD)



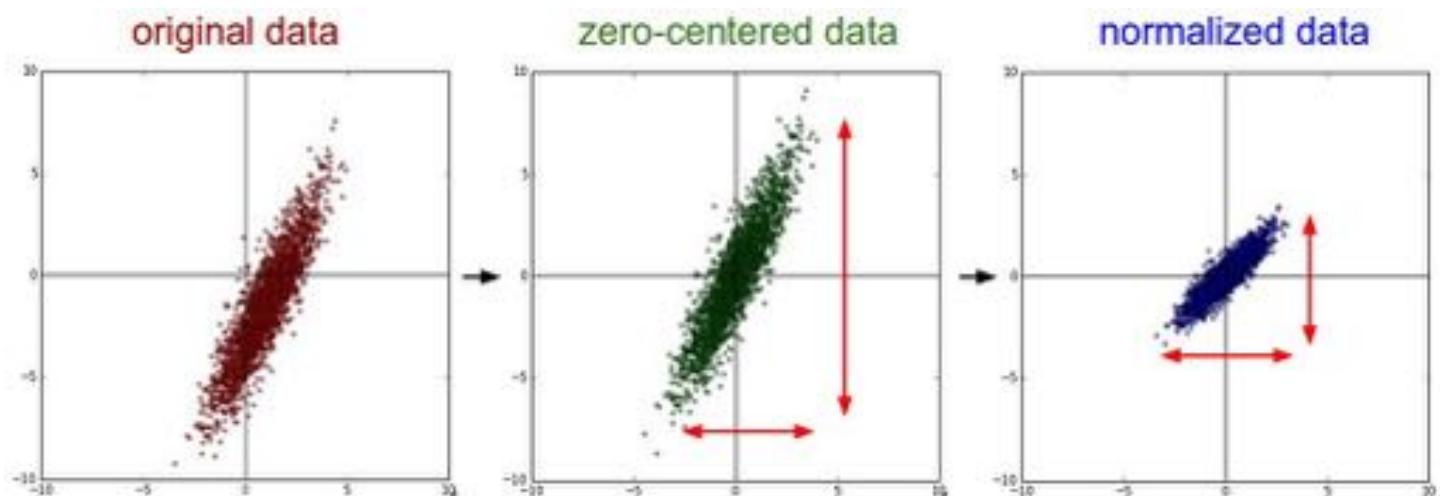
<https://distill.pub/2017/momentum/>



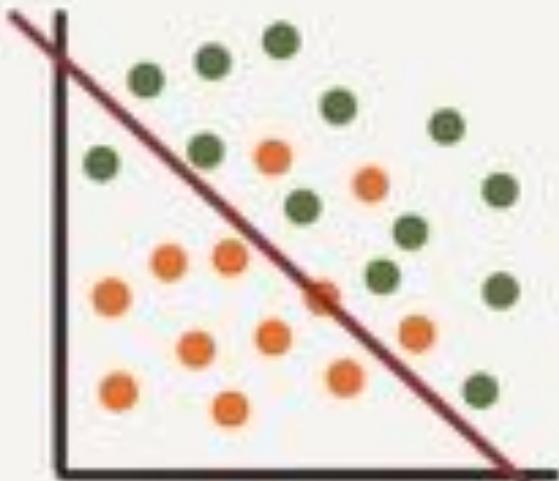
Animations that may help your intuitions about the learning process dynamics. **Left:** Contours of a loss surface and time evolution of different optimization algorithms. Notice the "overshooting" behavior of momentum-based methods, which make the optimization look like a ball rolling down the hill. **Right:** A visualization of a saddle point in the optimization landscape, where the curvature along different dimension has different signs (one dimension curves up and another down). Notice that SGD has a very hard time breaking symmetry and gets stuck on the top. Conversely, algorithms such as RMSprop will see very low gradients in the saddle direction. Due to the denominator term in the RMSprop update, this will increase the effective learning rate along this direction, helping RMSProp proceed. Images credit: [Alec Radford](#).

Data preprocessing

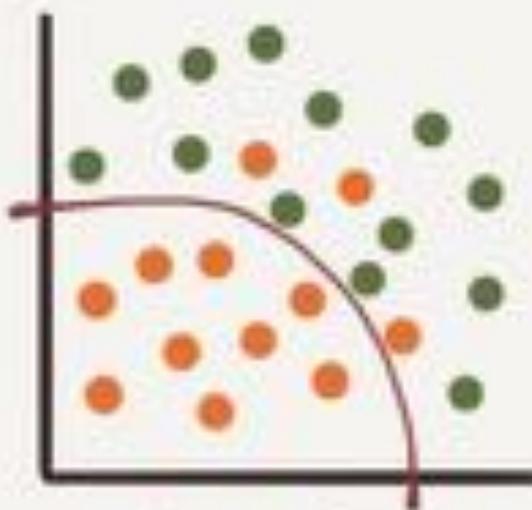
- **Mean subtraction** (most common): `X -= np.mean(X, axis = 0)`
(subtracting the mean across every individual feature in the data)
- **Normalization**: `X /= np.std(X, axis = 0)`
(normalizing the data so that all features are approximately the same scale)



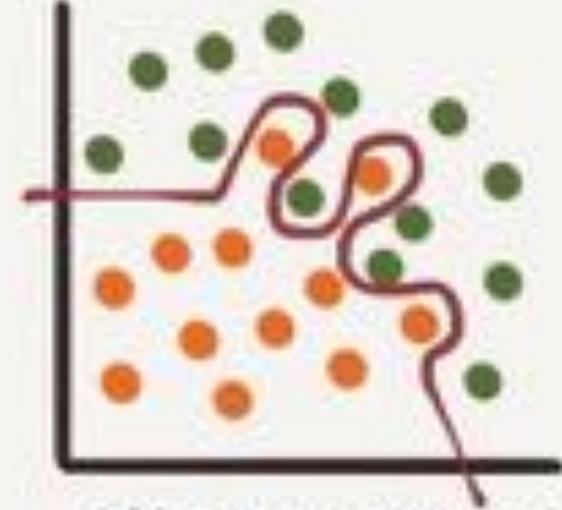
Over-fitting



HIGH BIAS
"UNDERFIT"



JUST RIGHT

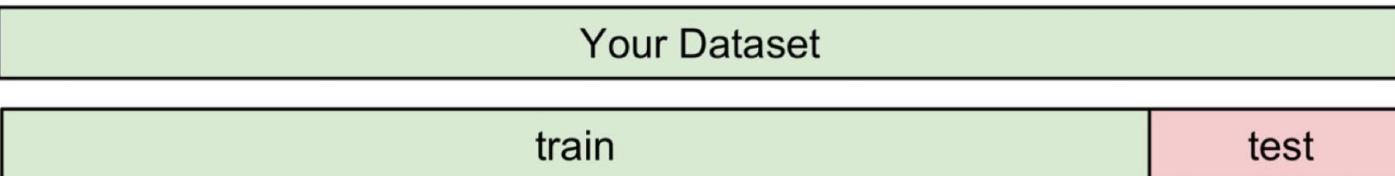
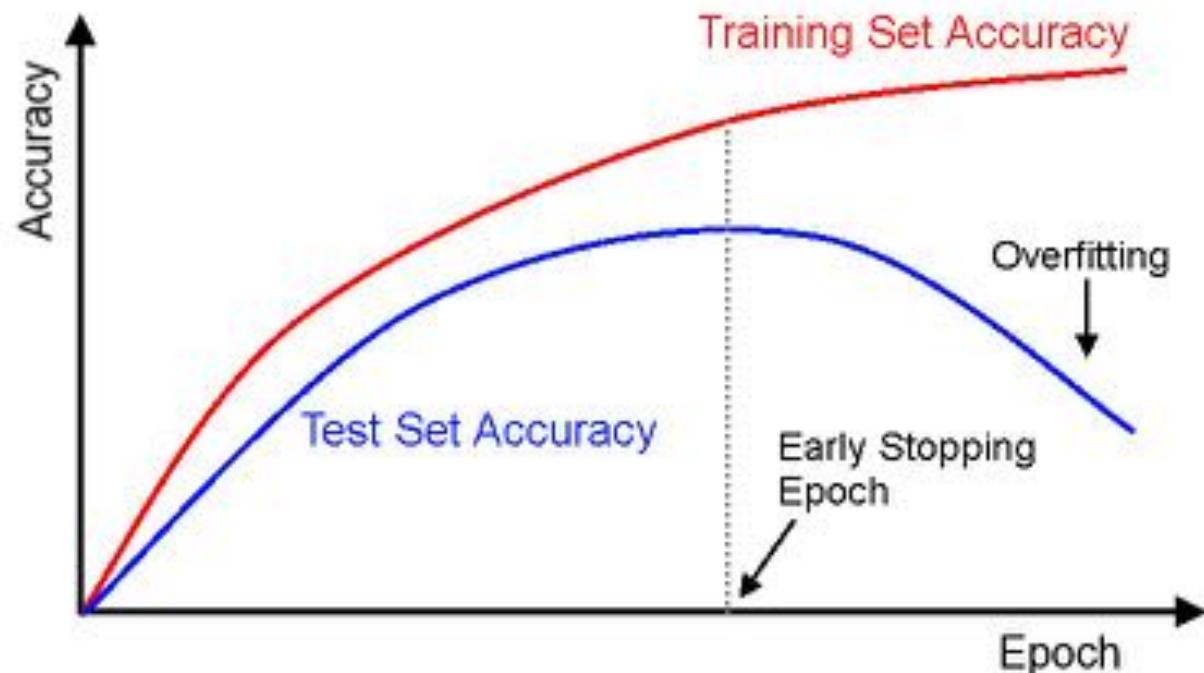


HIGH VARIANCE
"OVERFIT"

Even when we have far more examples than features, deep neural networks are capable of over-fitting.

Overfitting

- Early stopping
- Regularization
- Dropout

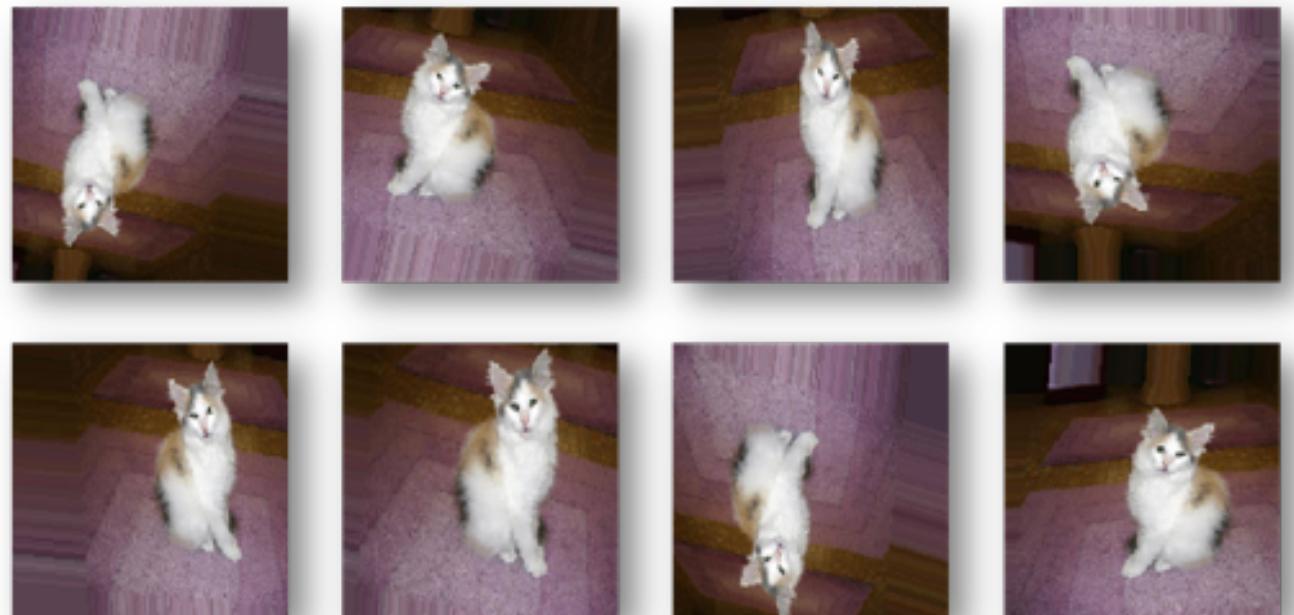


$$\text{Accuracy} = \frac{\text{\# correct predictions}}{\text{\# total predictions}}$$

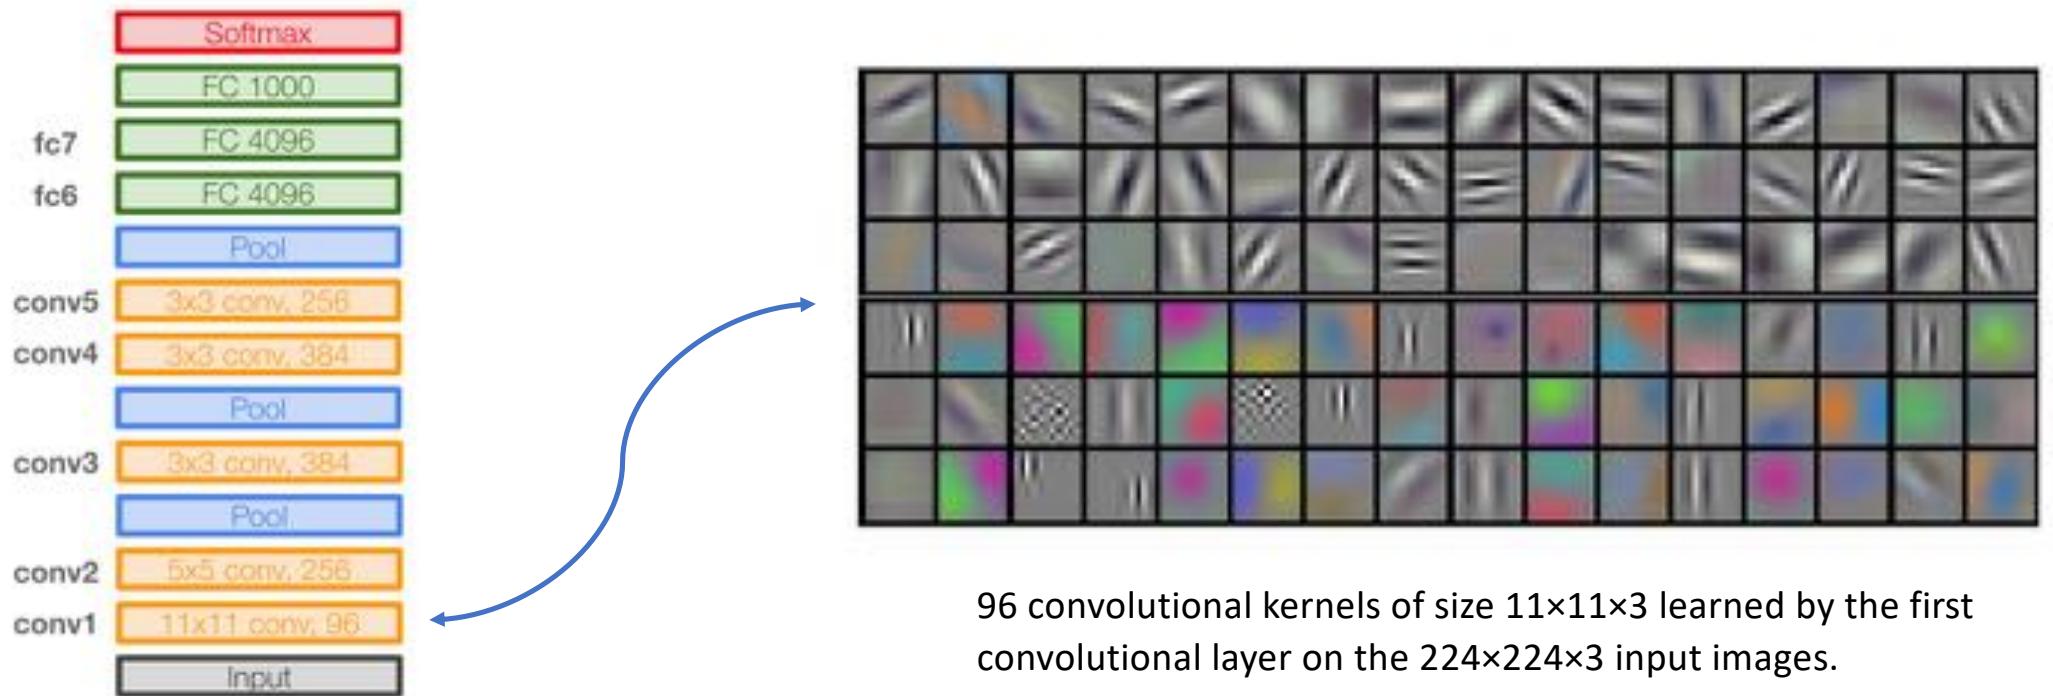
Data Augmentation

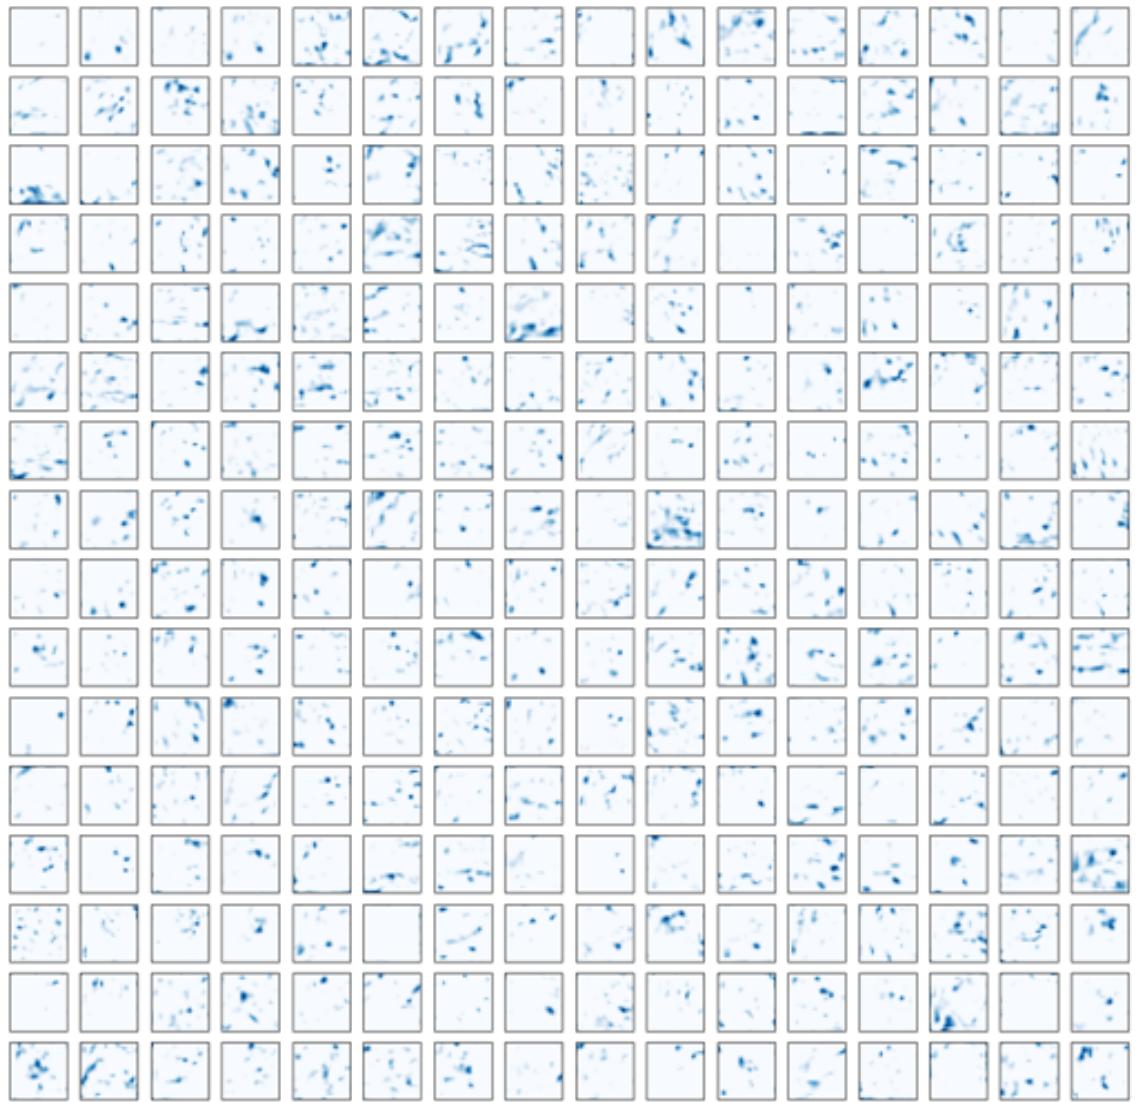
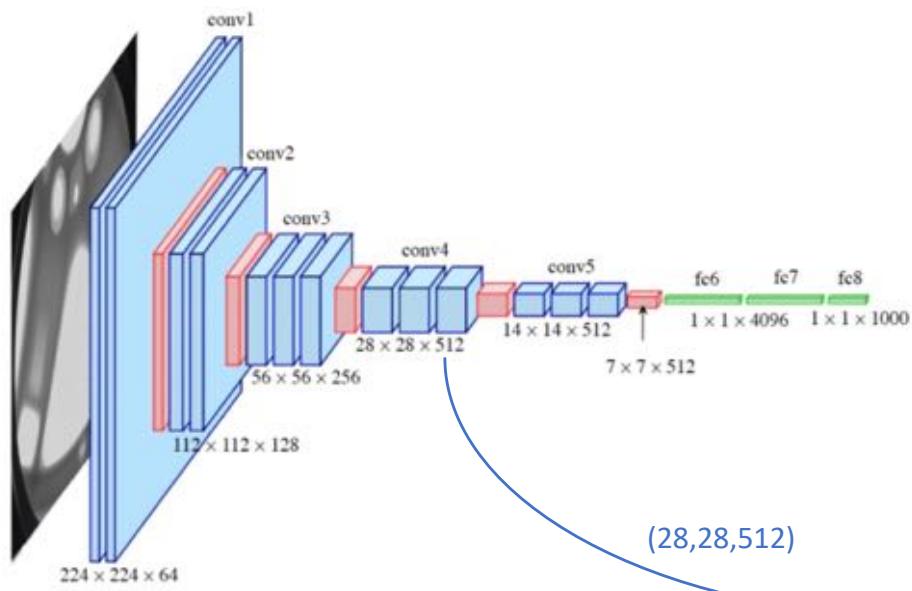
- Artificially creating more images from images you already have by changing the size, orientation etc. of the image

```
from keras.preprocessing.image import ImageDataGenerator  
  
ImageDataGenerator(  
    rotation_range=10.,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    shear_range=0.,  
    zoom_range=.1.,  
    horizontal_flip=True,  
    vertical_flip=True)
```



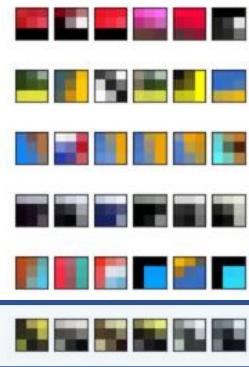
Features visualization



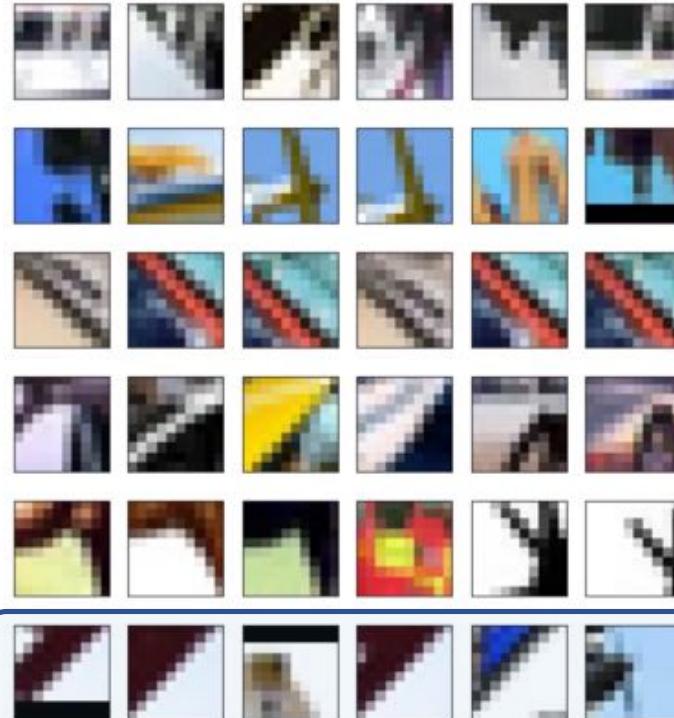


Activation maps

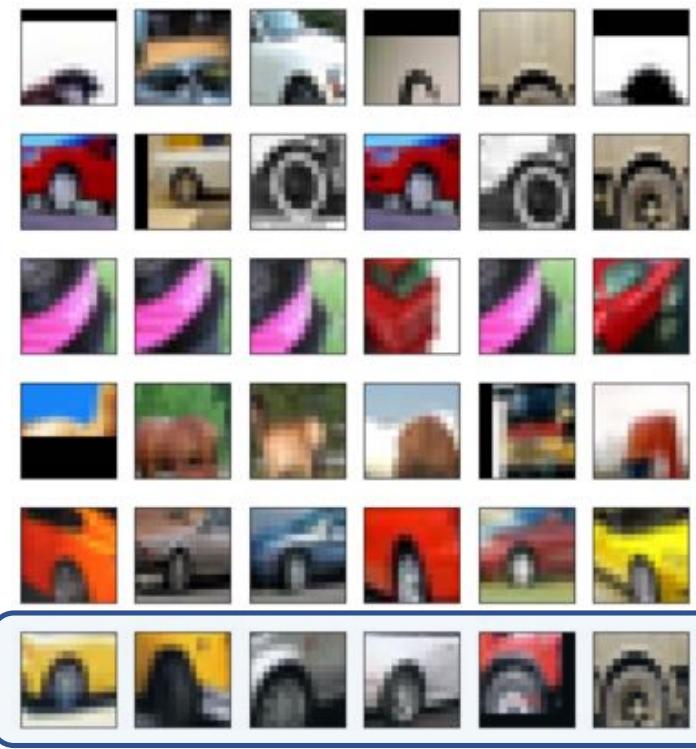
Maximally activating patches



a few patches that
excite one specific filter
in the 1st conv layer of
the network



a few patches that excite one specific filter
in the 2nd conv layer of the network



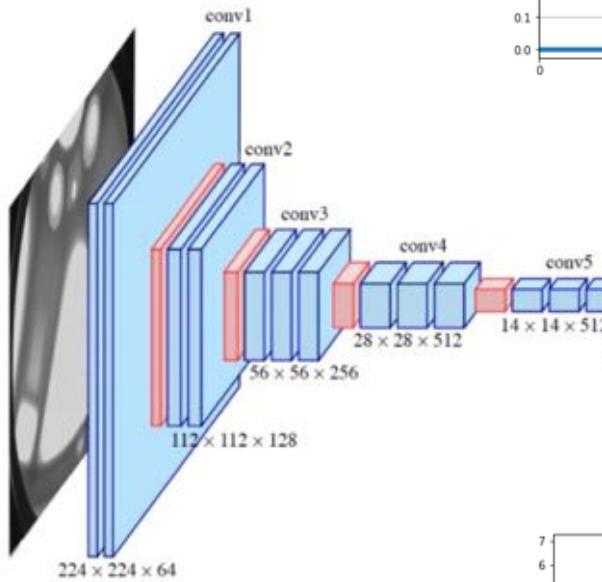
a few patches that excite one specific filter
in the 3rd conv layer of the network



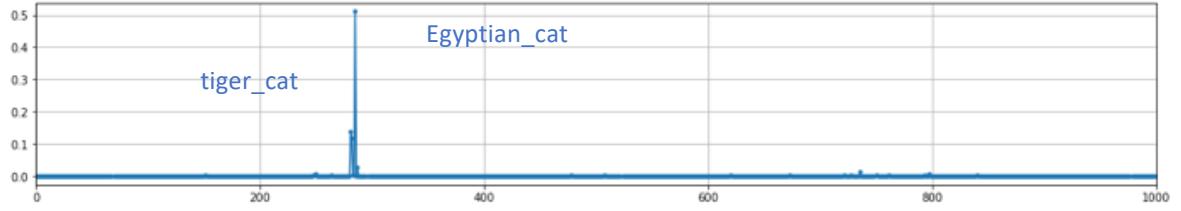
(224,224,3)



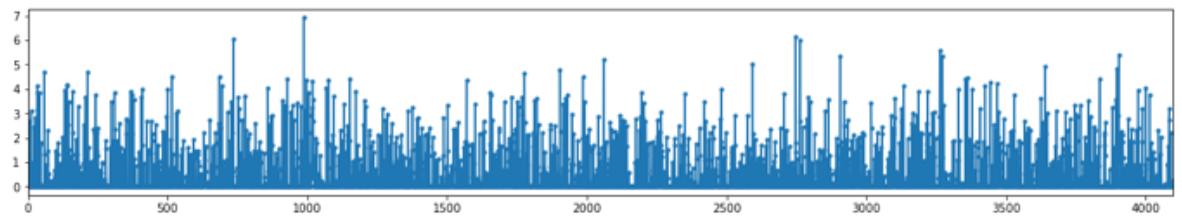
150528-D pixels space



VGG-16

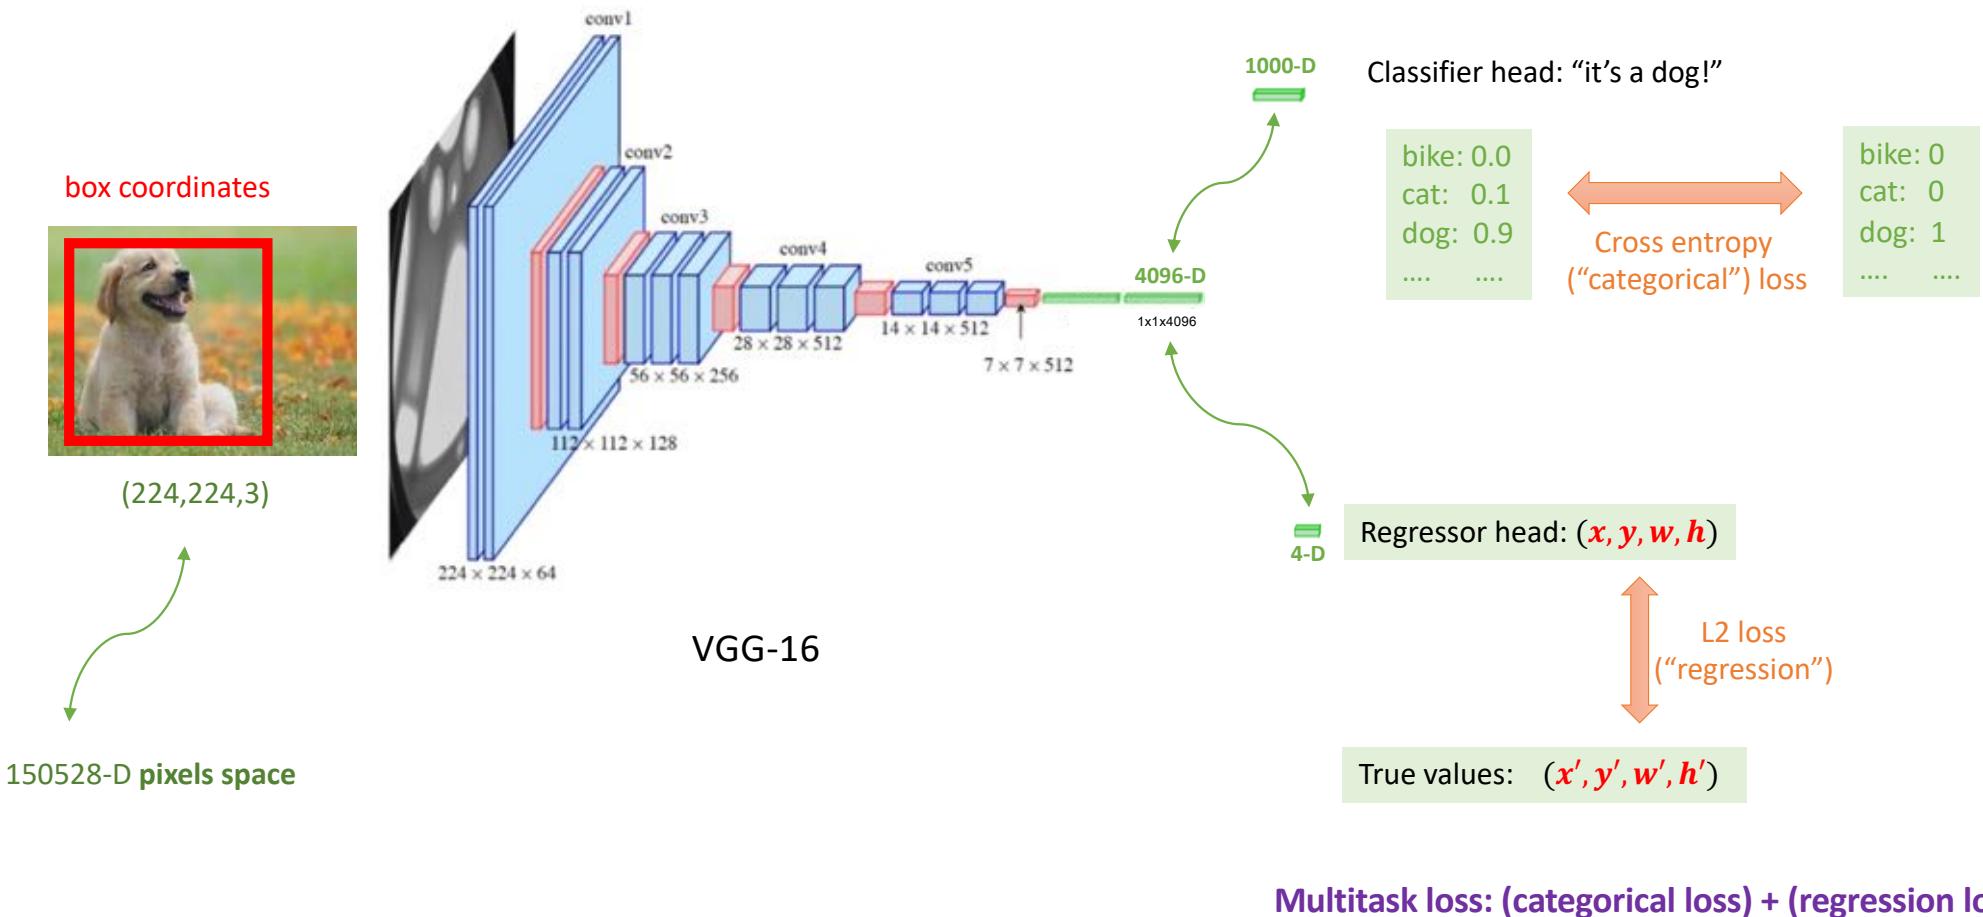


1000-D

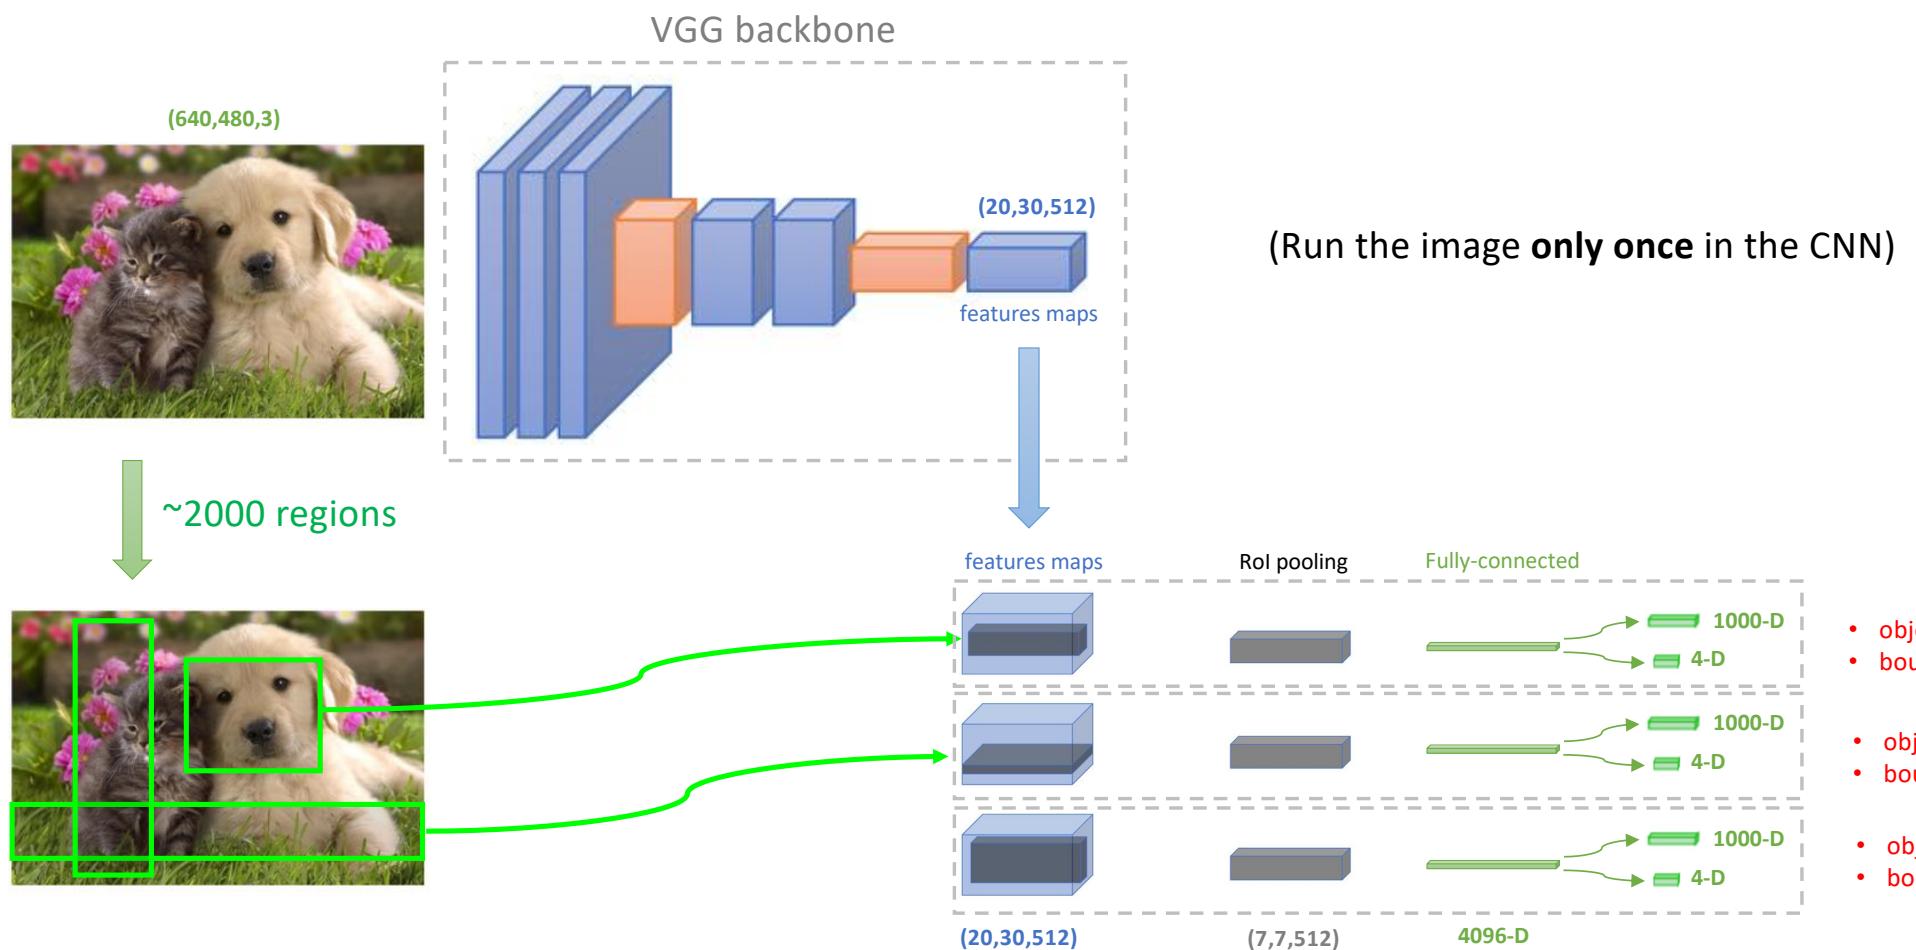


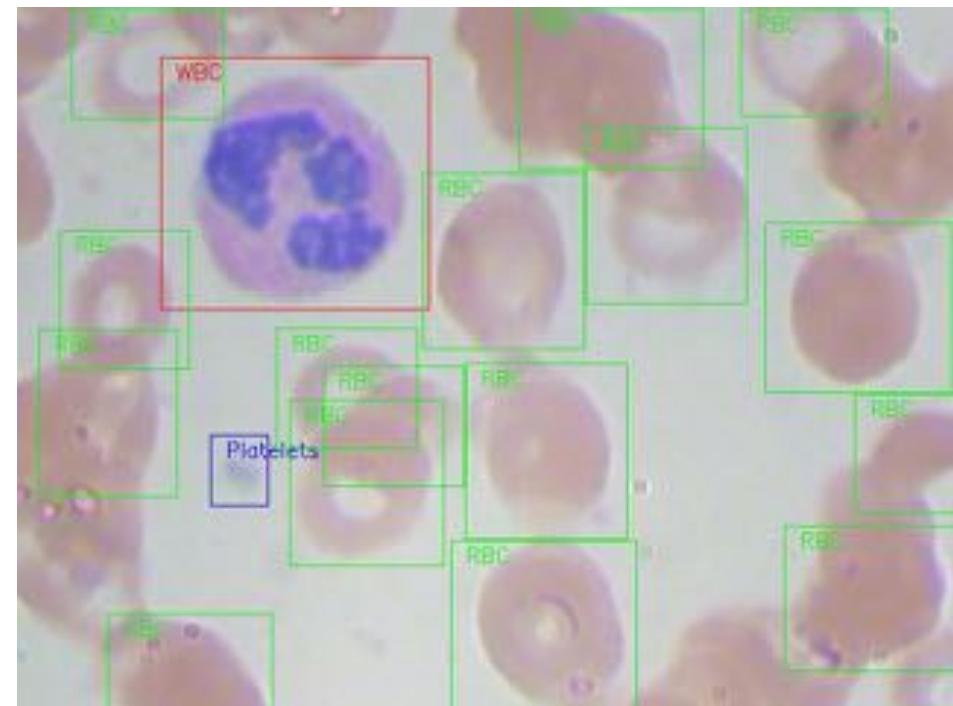
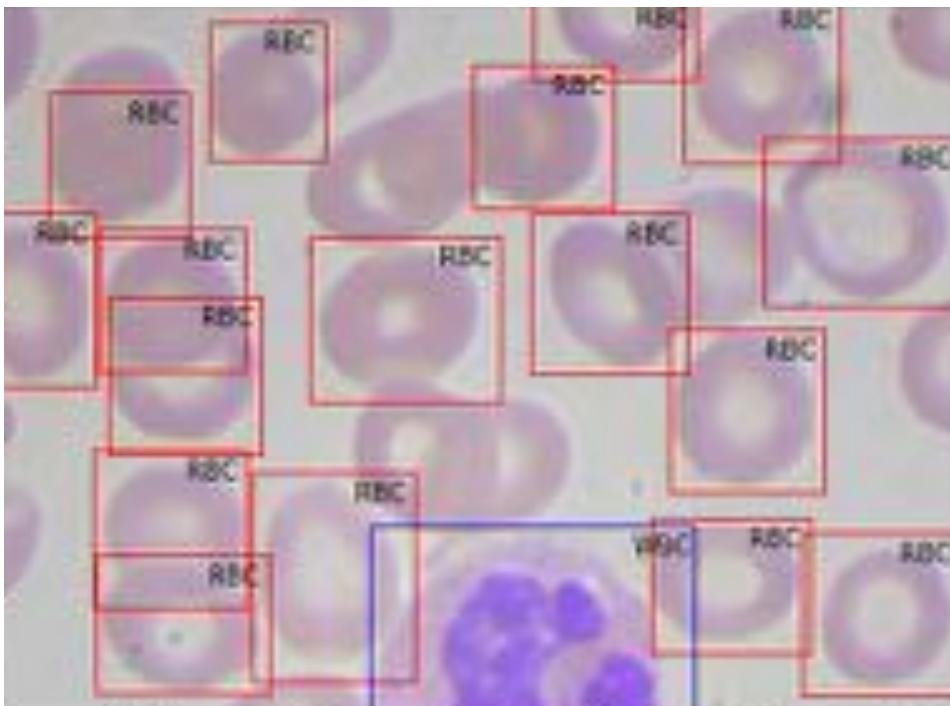
4096-D features space

Single object: Classification + Localization



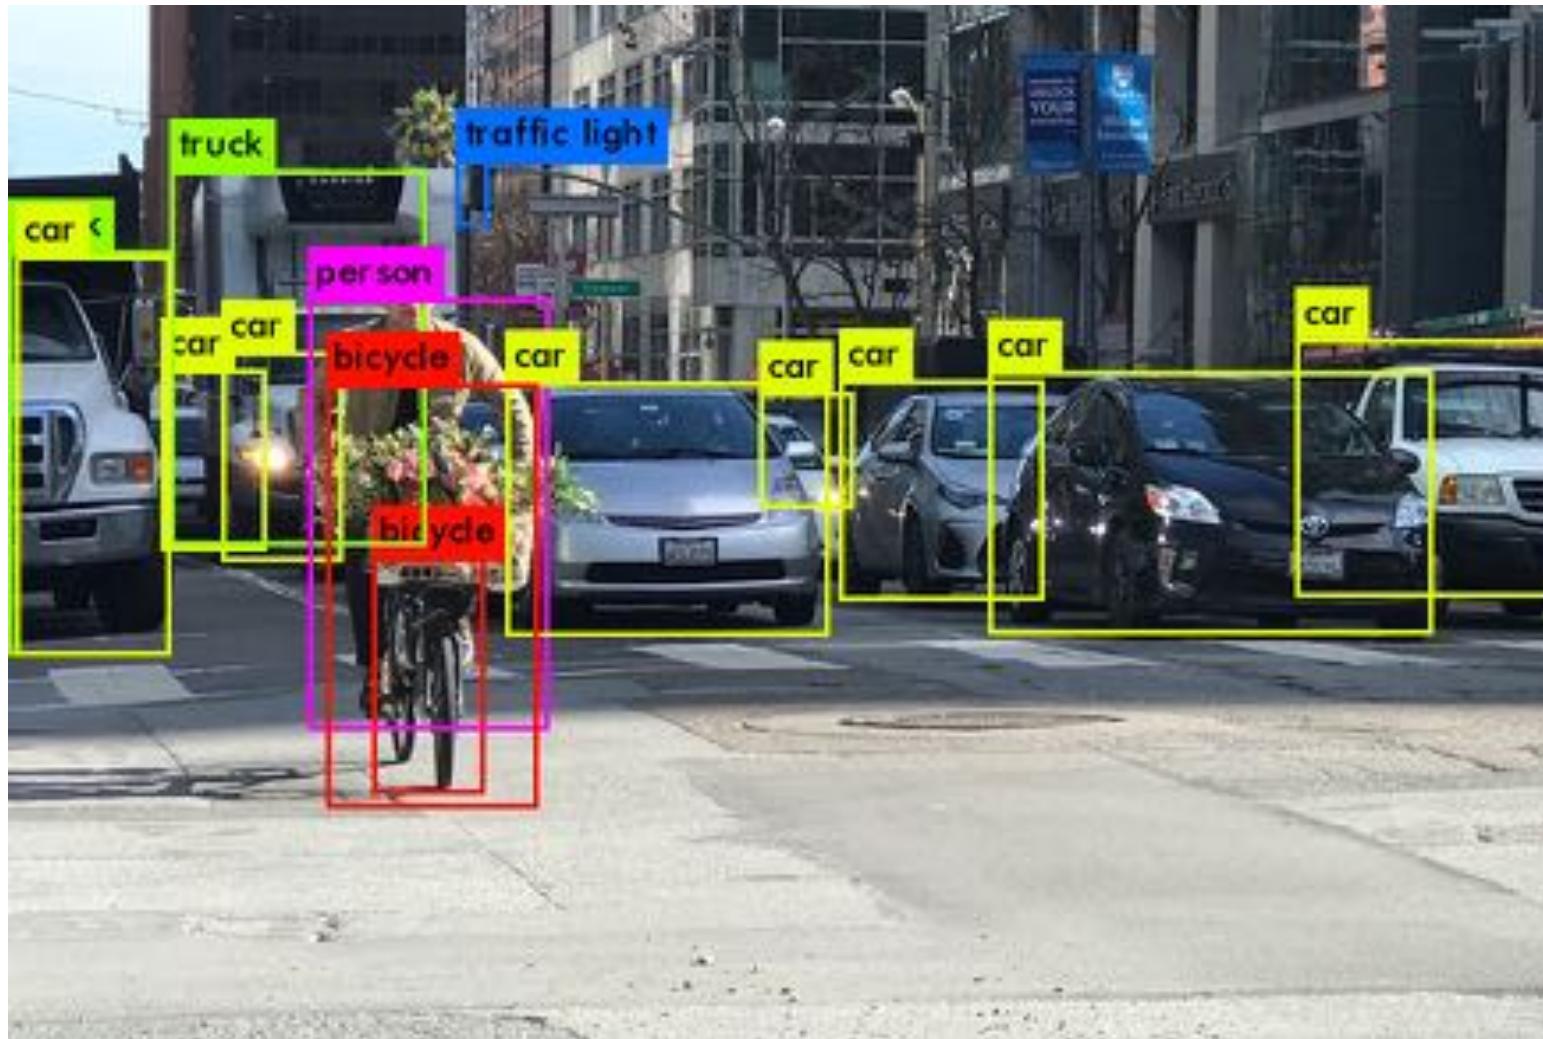
Fast R-CNN



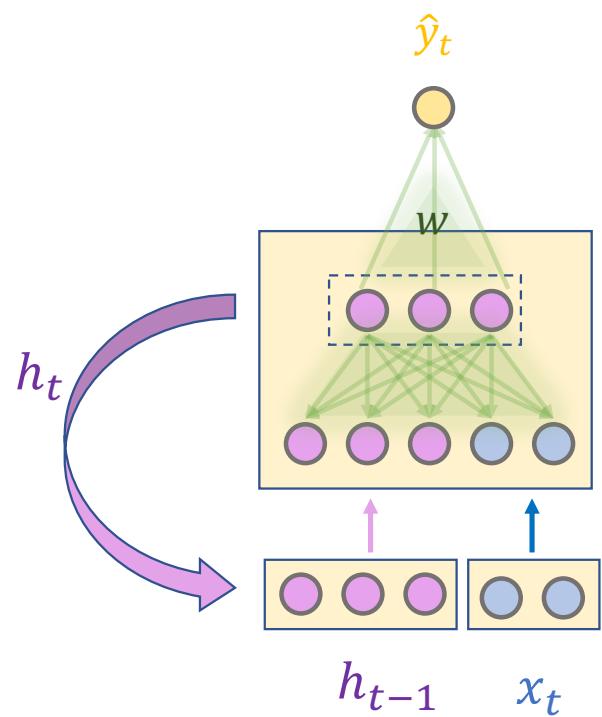


<https://www.analyticsvidhya.com/blog/2018/11/implementation-faster-r-cnn-python-object-detection/>

YOLO

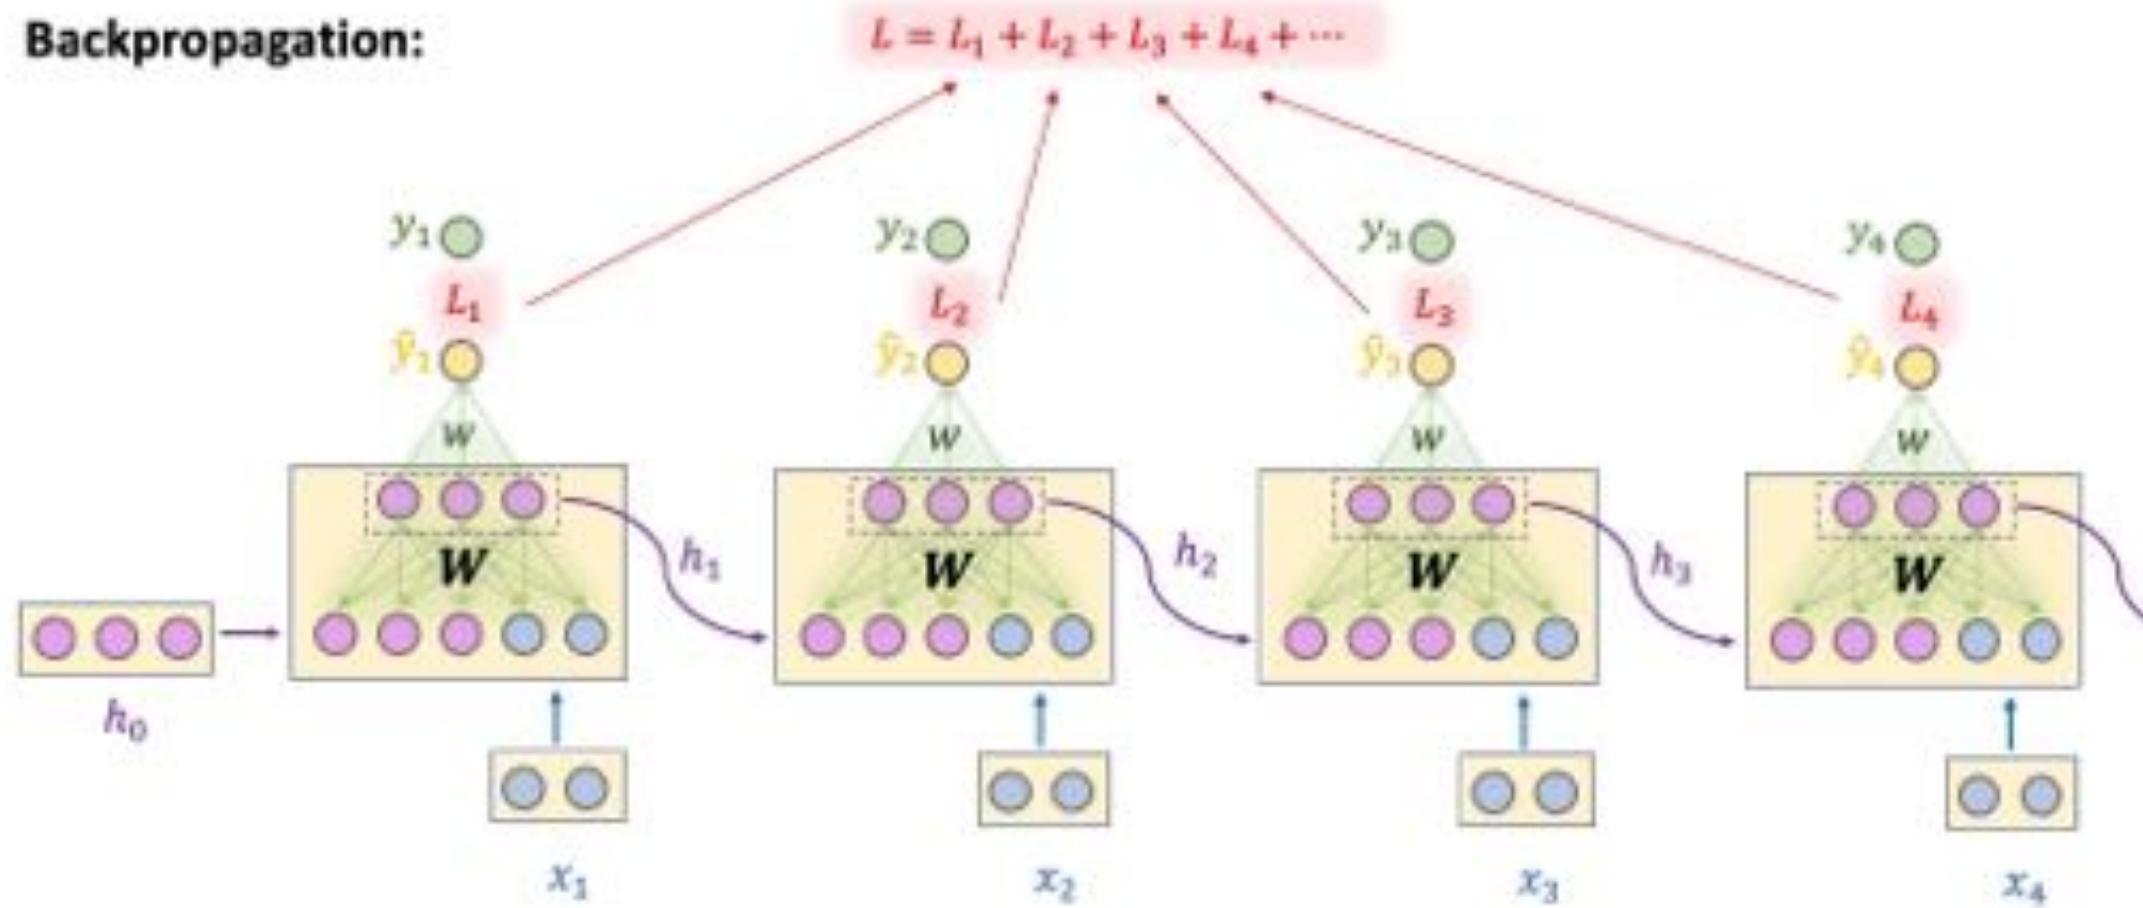


Recurrent neural network (RNN)



<https://www.tensorflow.org/guide/keras/rnn>

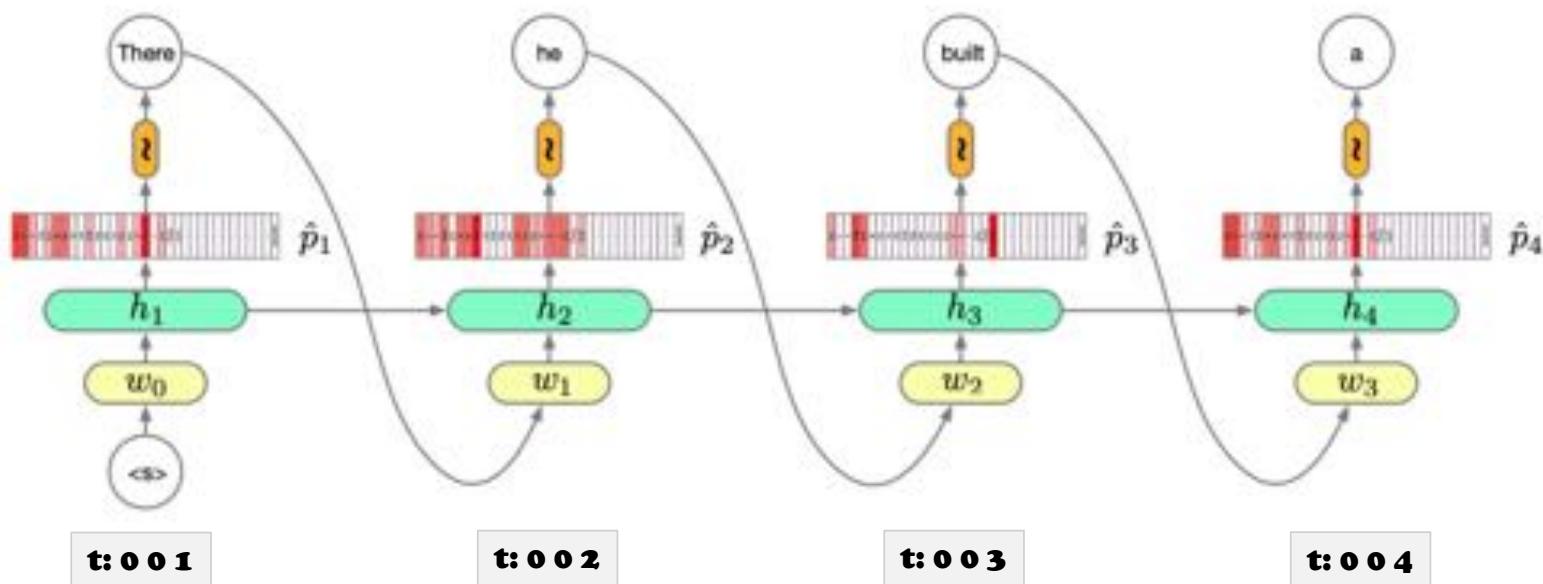
Backpropagation:



Predicting next word

Words representation (1-hot):

Man (5391)	Woman (9853)	King (4914)	Queen (7157)	Apple (456)	Orange (6257)
$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ \vdots \\ 0 \\ 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$



at first:

tyntd-iafhatawiaoihrdemot lytdws e ,tfti, astai f ogoh eoase rrranbyne 'nhthnee e
plia tkldrgd t o idoe ns,smtt h ne etie h,hregtrs nigtike,aoaenns lng

↓ train more

"Tmont thithey" fomesscerliund
Keushey. Thom here
sheulke, anmerenith ol sivh I lalterthend Bleipile shuwyl fil on aseterlome
coaniogennc Phe lism thond hon at. MeiDimorotion in ther thize."

↓ train more

Aftair fall unsuch that the hall for Prince Velzonski's that me of
her hearly, and behs to so arwage fiving were to it beloge, pavu say falling misfort
how, and Gogition is so overelical and ofter.

↓ train more

"Why do what that day," replied Natasha, and wishing to himself the fact the
princess, Princess Mary was easier, fed in had oftened him.
Pierre aking his soul came to the packs and drove up his father-in-law women.

- **Sentiment Analysis:**
(many-to-one)

“There is nothing to like
in this movie.” → 
- **Translation:**

Voulez-vous chanter avec
moi? → Do you want to sing with
me?
- **Name entity recognition:**
(many-to-many)

Yesterday, Harry Potter
met Hermione Granger. → Yesterday, **Harry Potter**
met **Hermione Granger**.

- **Video activity recognition:**
(many-to-one)



Running

- **Image captioning:**
(one-to-many)



A dog is running in the grass with a frisbee



2014

The OG
GAN



2015

Deep
Convolutional
GAN



2016

Coupled
GAN



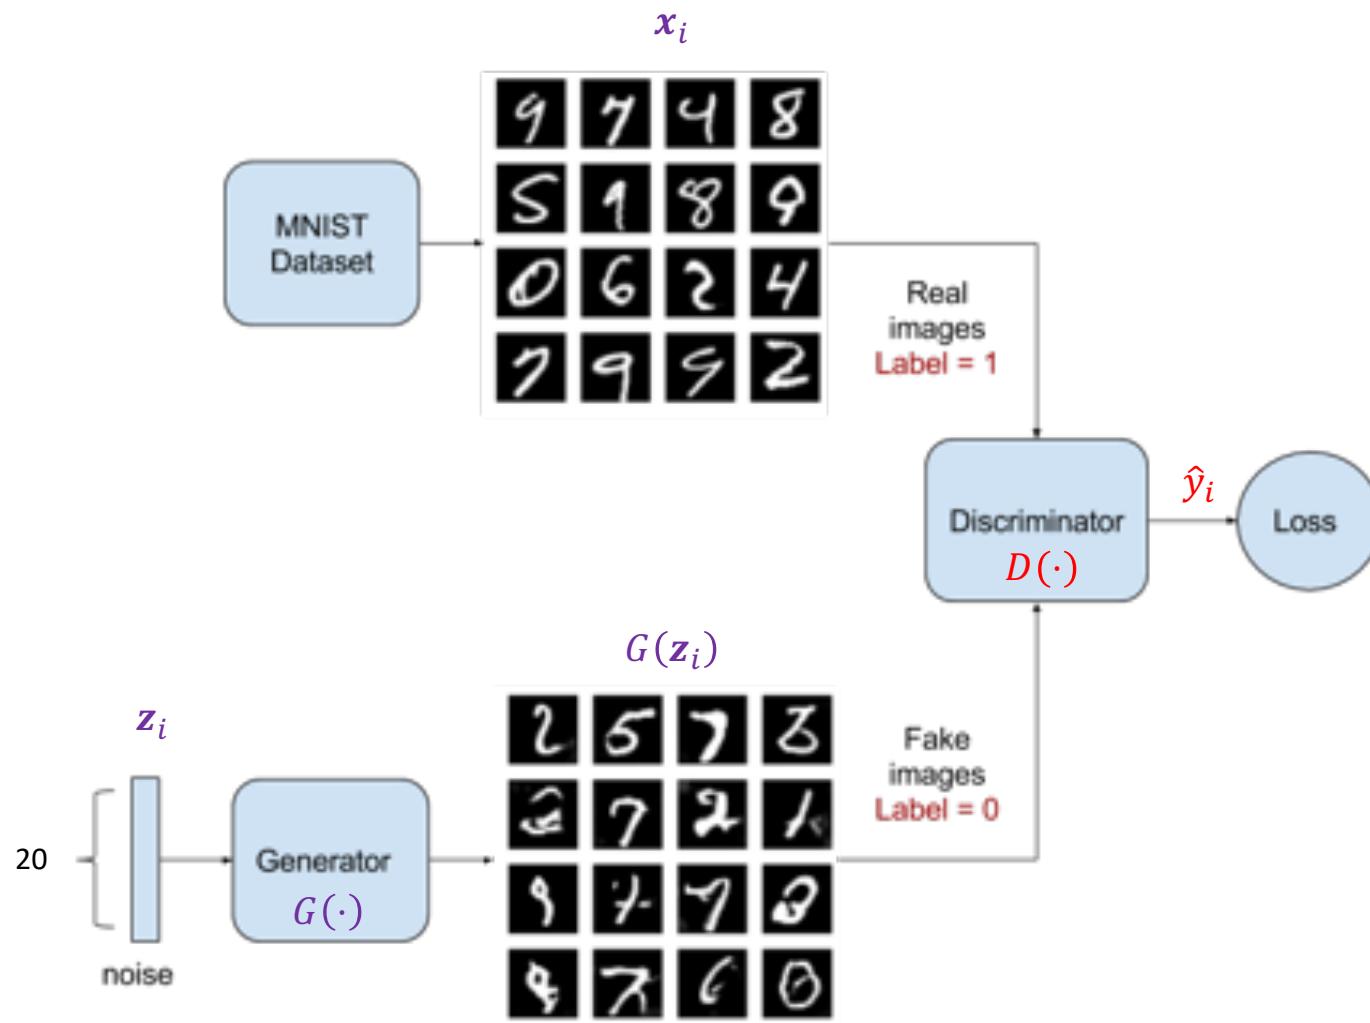
2017

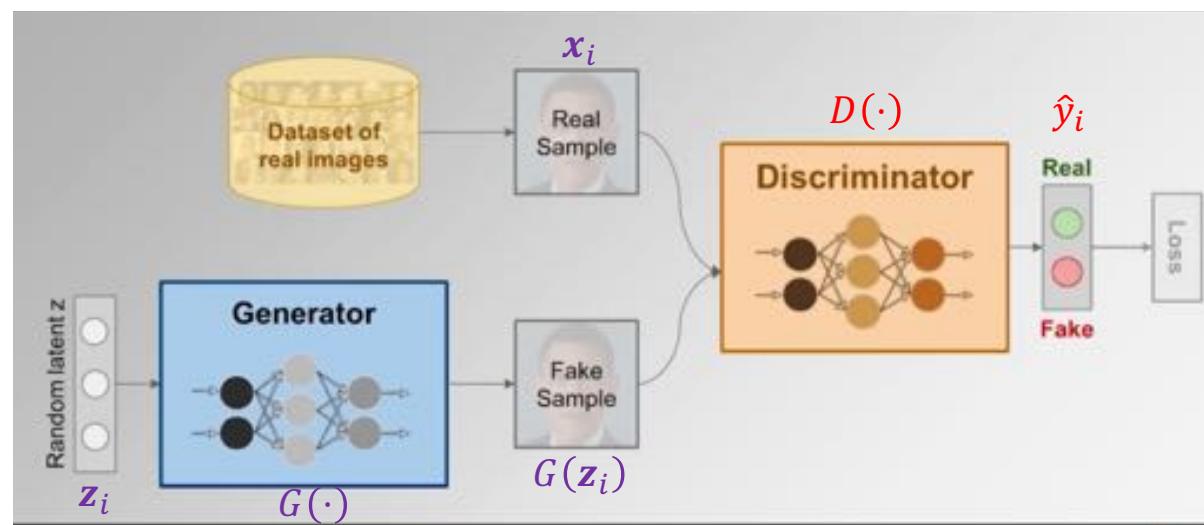
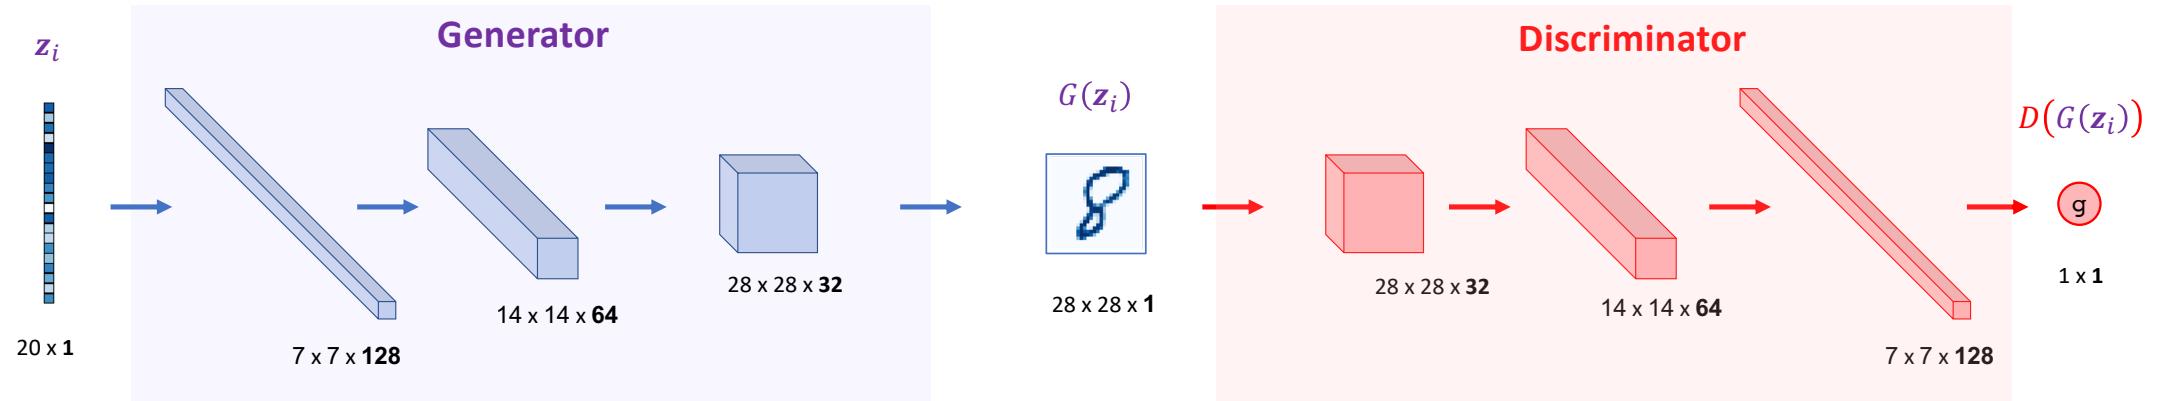
Progressively
Growing
GAN

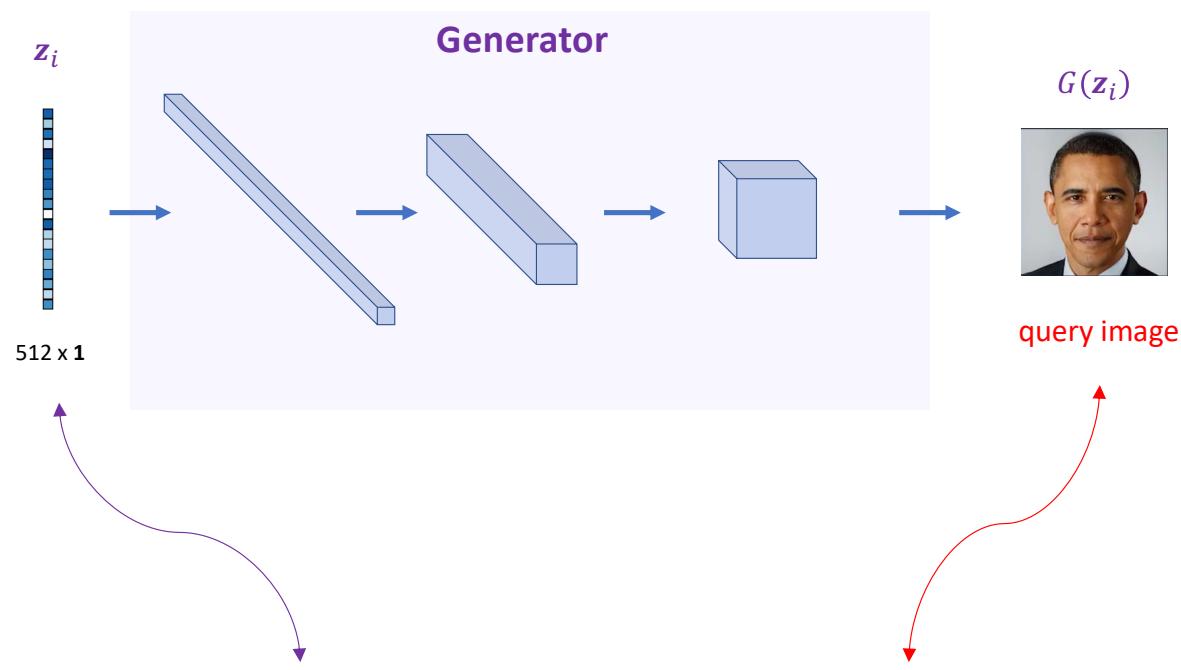


2018

Style-based
GAN

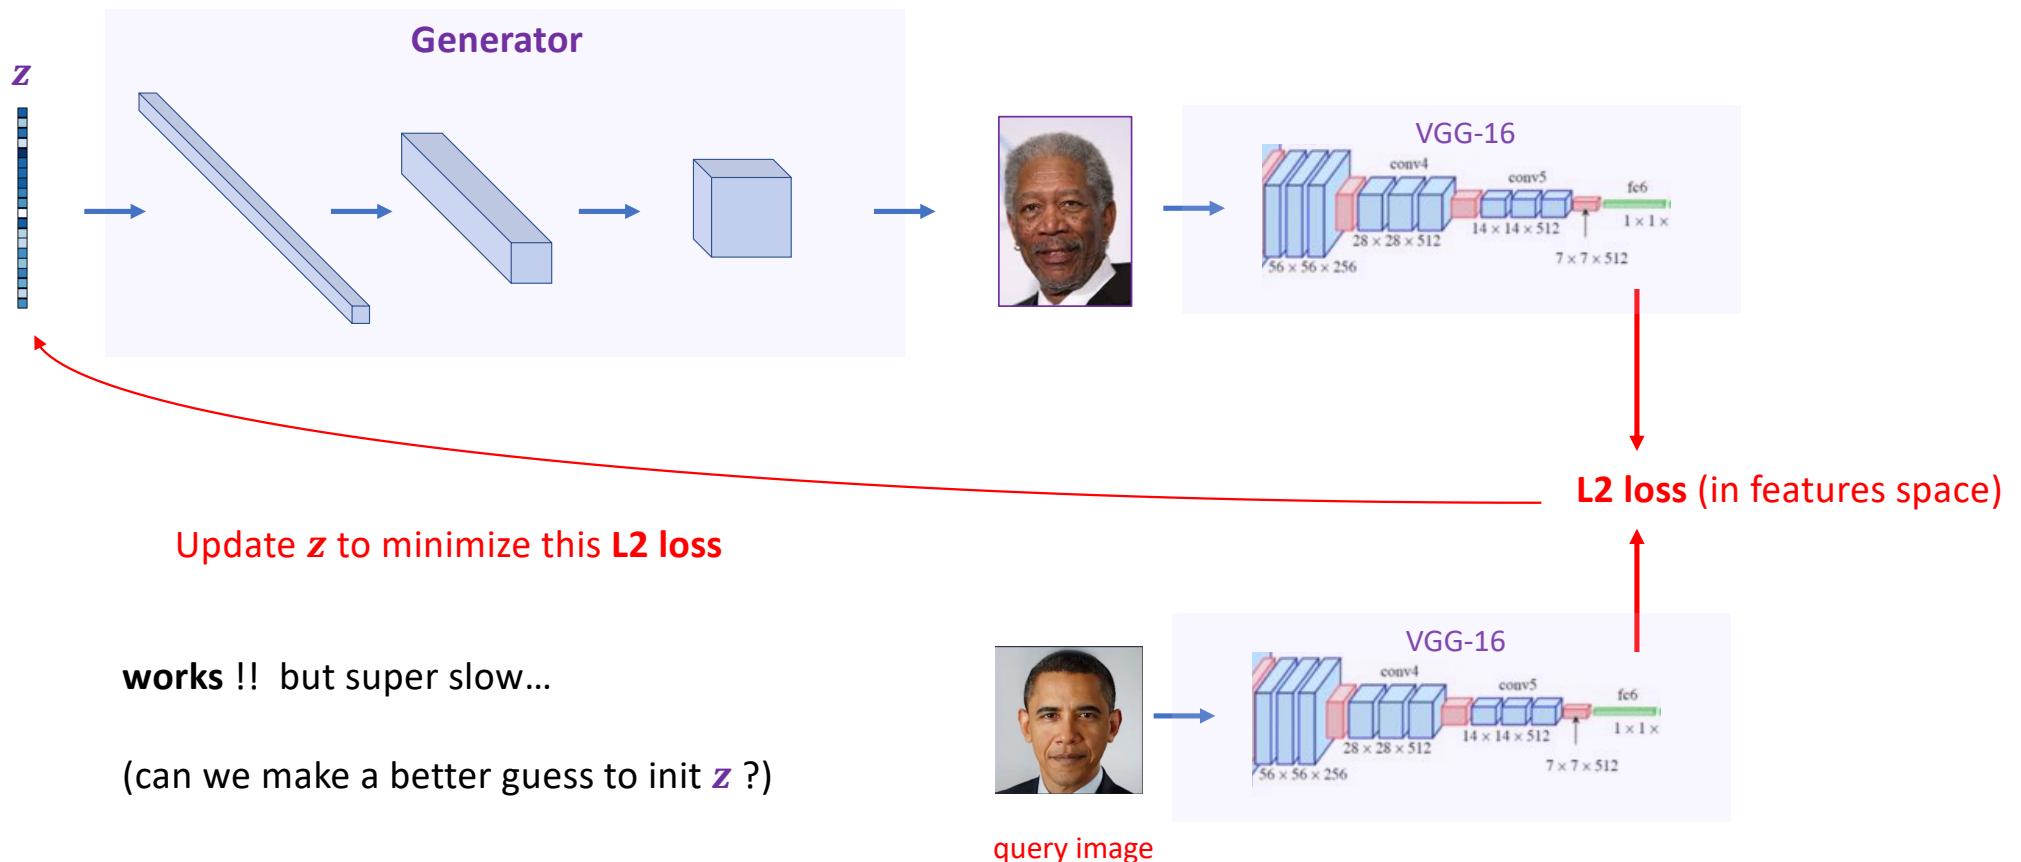






- How can we find \mathbf{z}_i that will produce a specific **query image**? (random search..?)

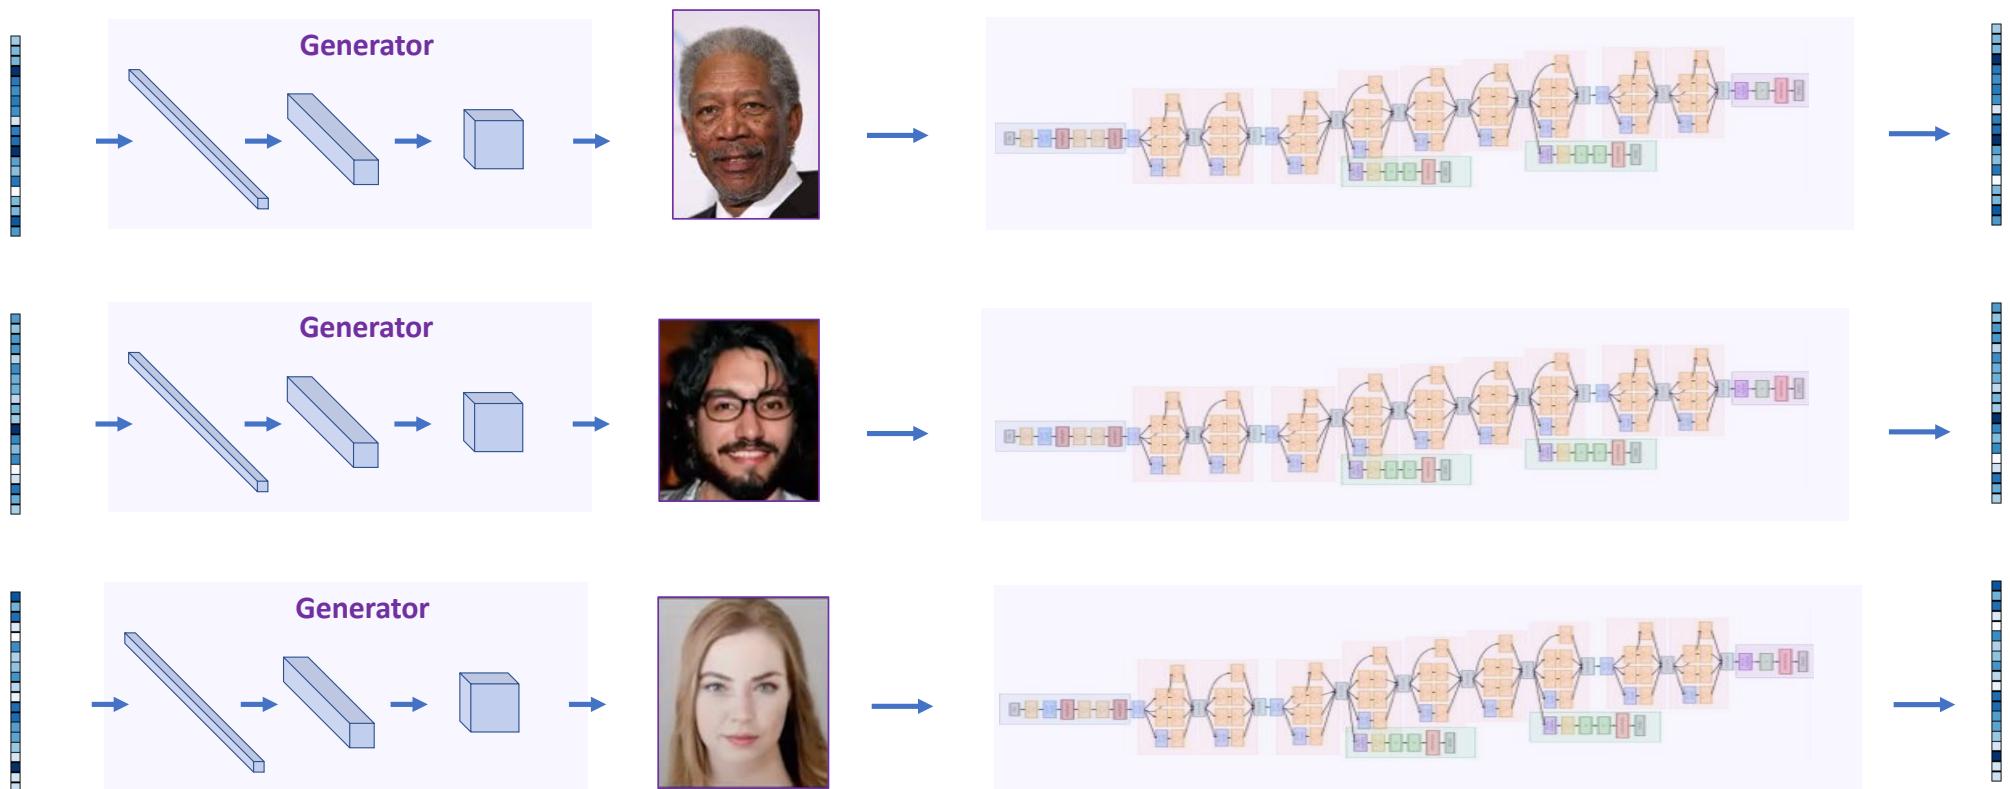
- How can we find \mathbf{z}_i that will produce a specific **query image**?



Can we make a better guess to init \mathbf{z} ?

Let's generate some training set:

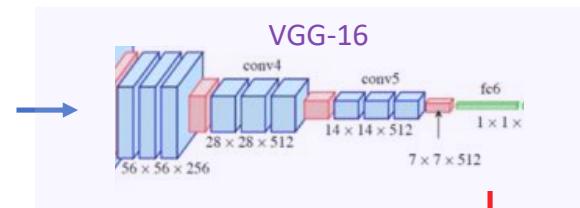
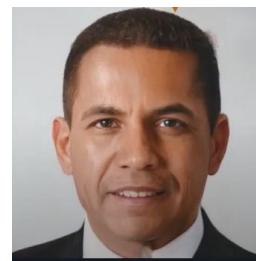
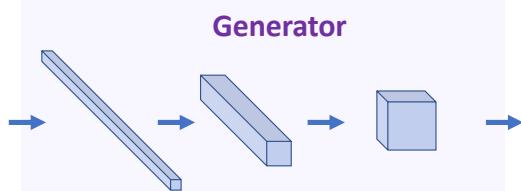
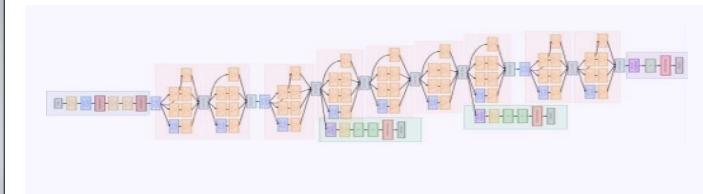
Then, **train a new model** to go from image to \mathbf{z} :



query image

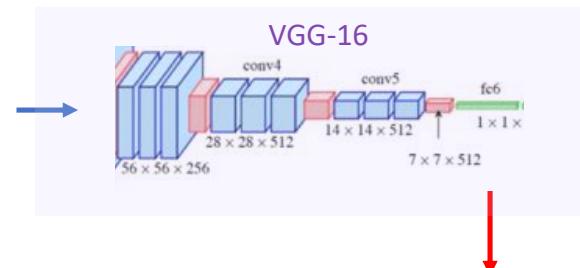
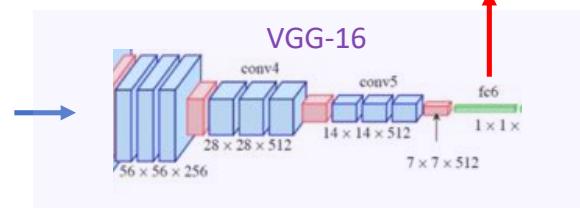


Initial z



z ← Update z to minimize this L2 loss

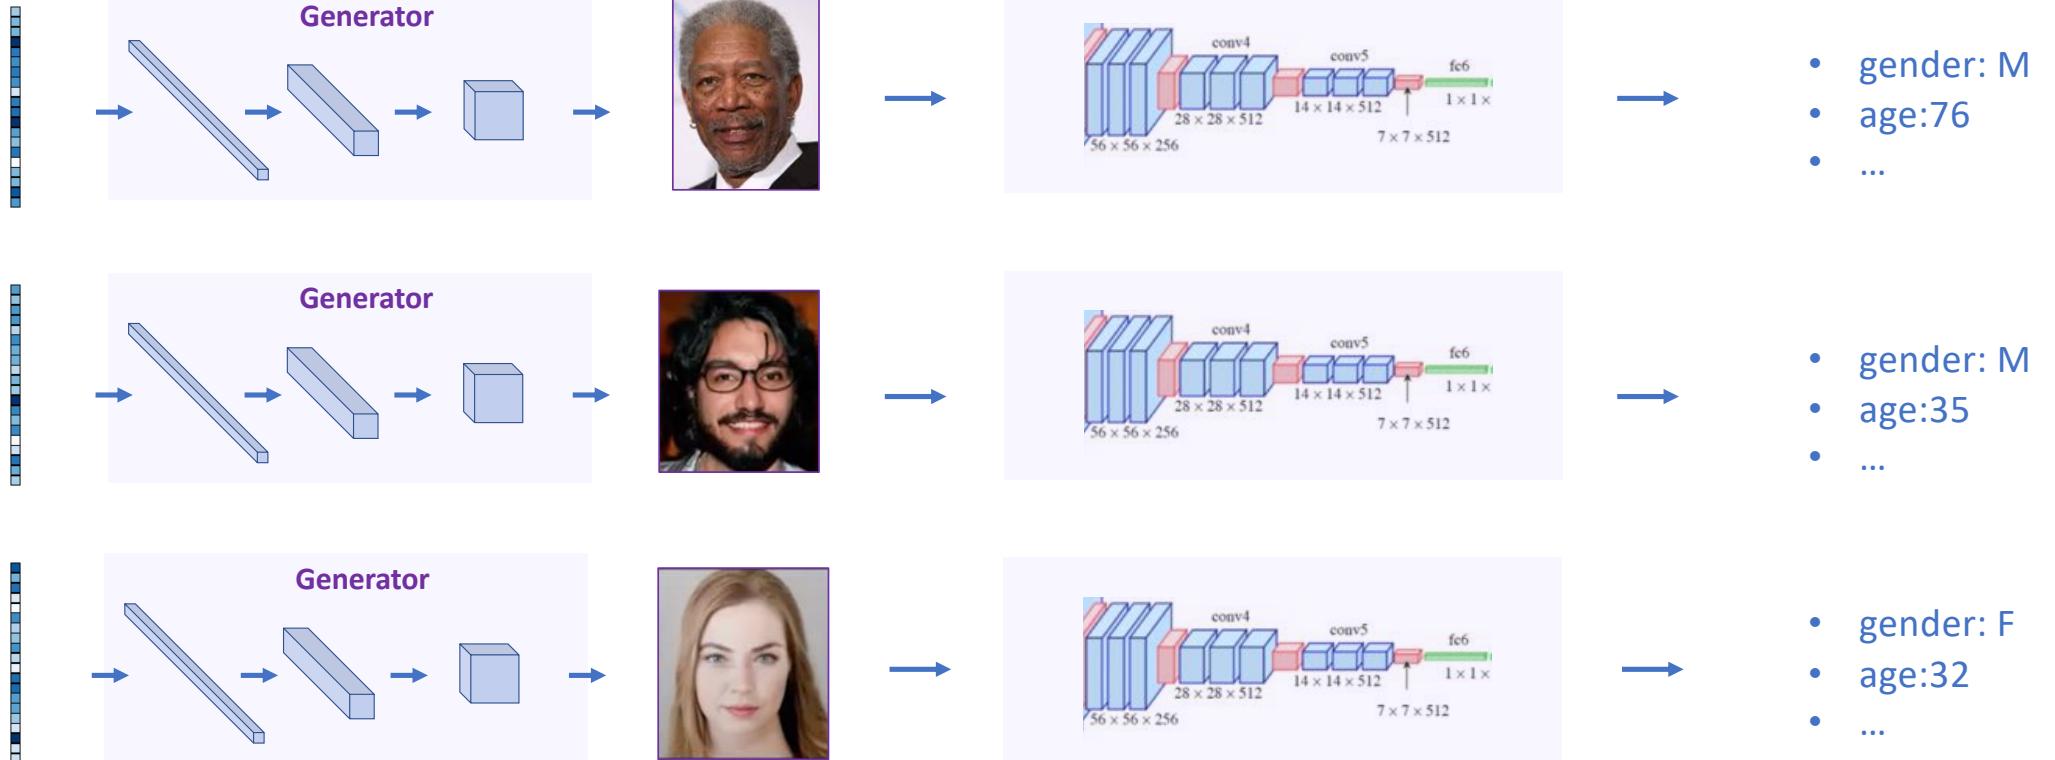
(Ok, so we found z_i , now what?)



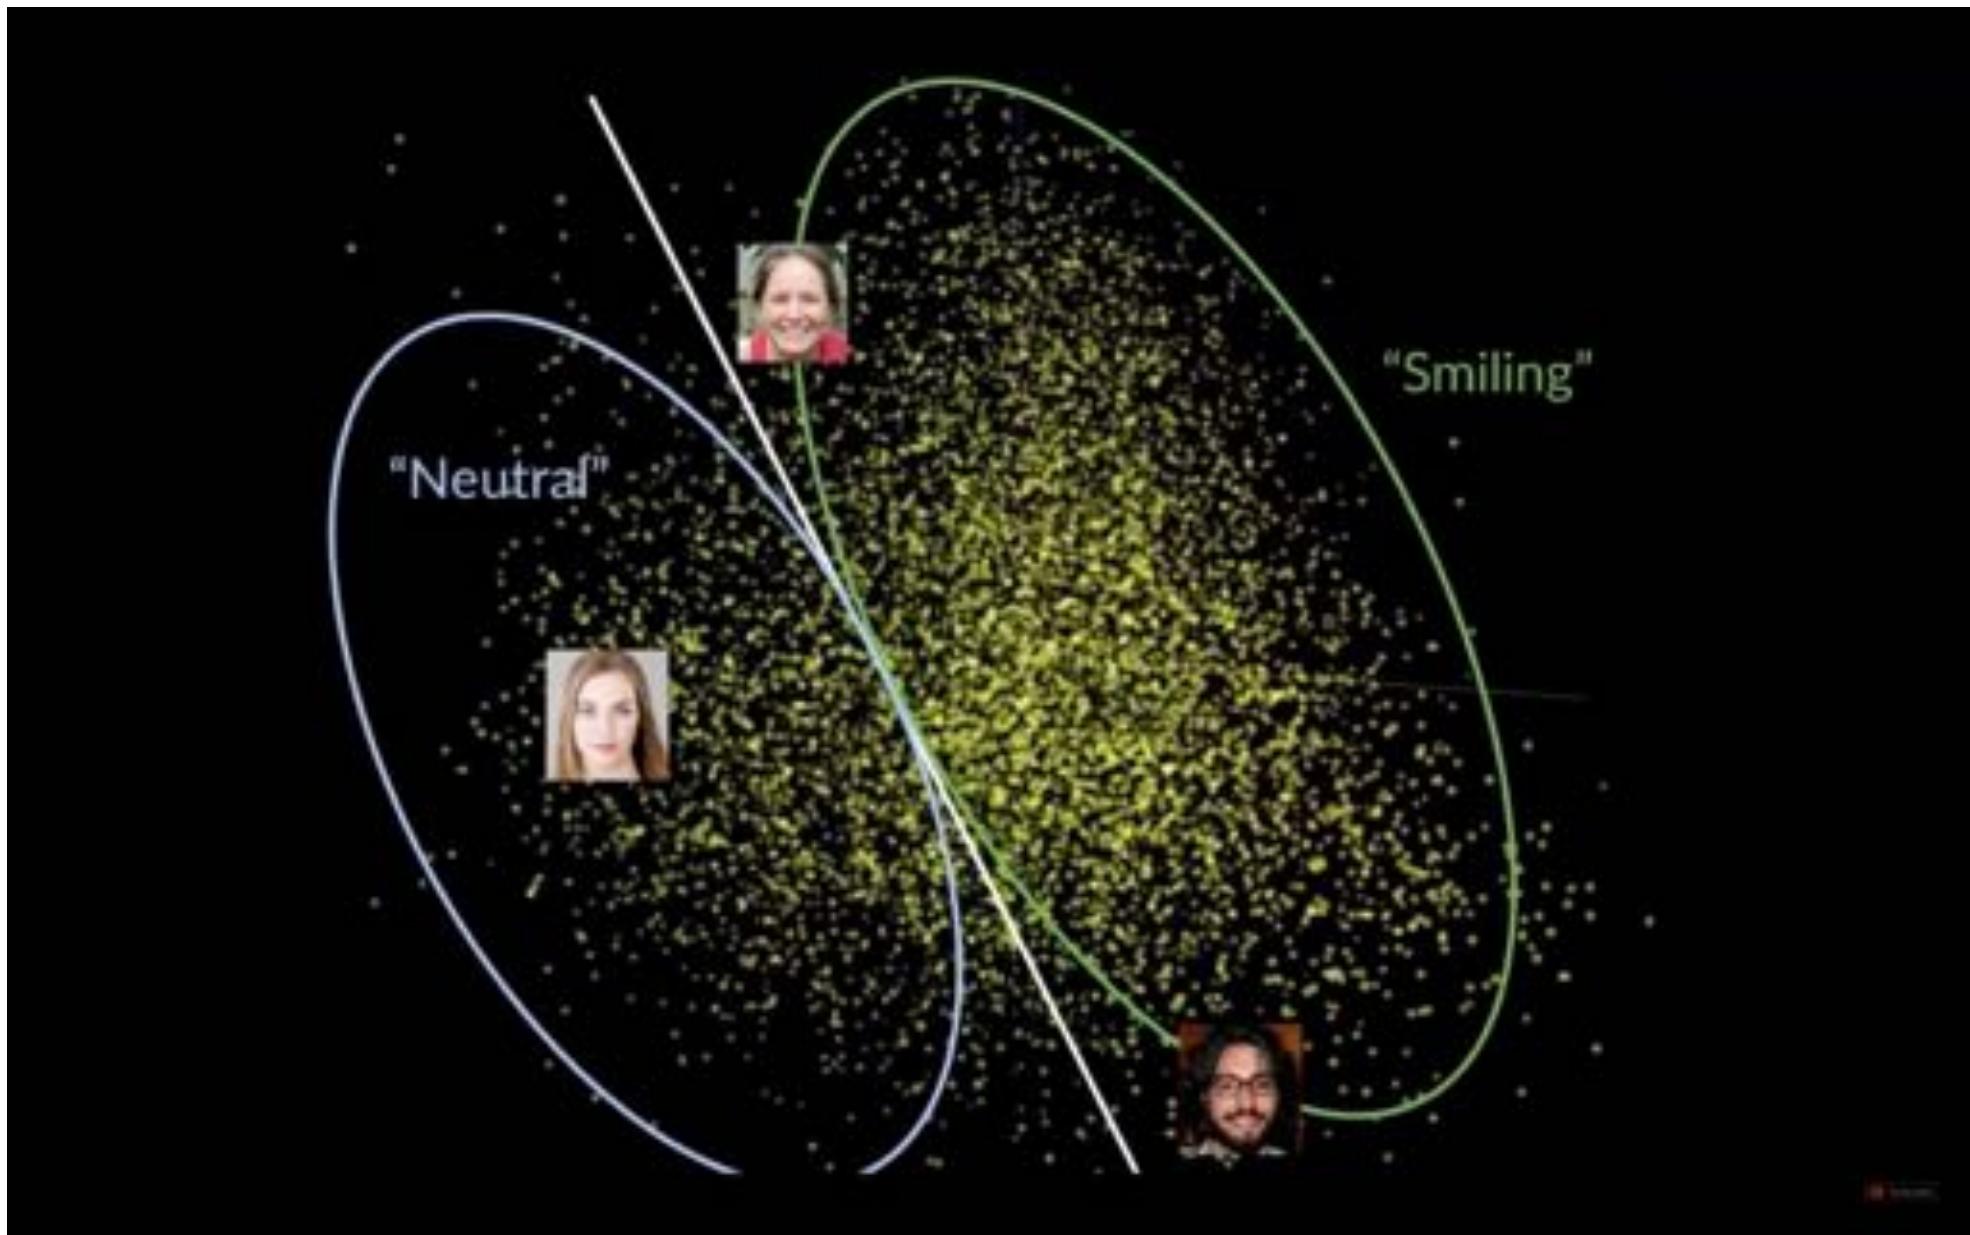
L2 loss (in features space)

Generate some
random vectors...

Use pre-trained classifier to extract some attributes:

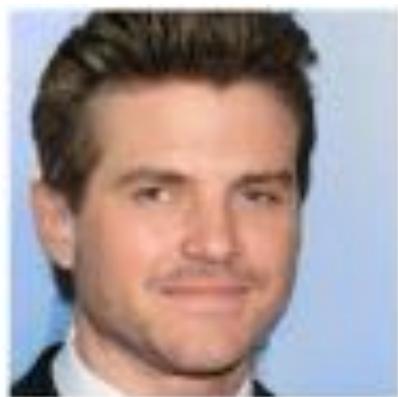








Original



Pose



Age



Expression



Eyeglasses





<https://youtu.be/0VH1Lim8gL8>

History of Deep Learning Ideas and Milestones*



Perspective:

- **Universe created**
13.8 billion years ago
- **Earth created**
4.54 billion years ago
- **Modern humans**
300,000 years ago
- **Civilization**
12,000 years ago
- **Written record**
5,000 years ago

- 1943: Neural networks
- 1957-62: Perceptron
- 1970-86: Backpropagation, RBM, RNN
- 1979-98: CNN, MNIST, LSTM, Bidirectional RNN
- 2006: “Deep Learning”, DBN
- 2009: ImageNet + AlexNet
- 2014: GANs
- 2016-17: AlphaGo, AlphaZero
- 2017: 2017-19: Transformers

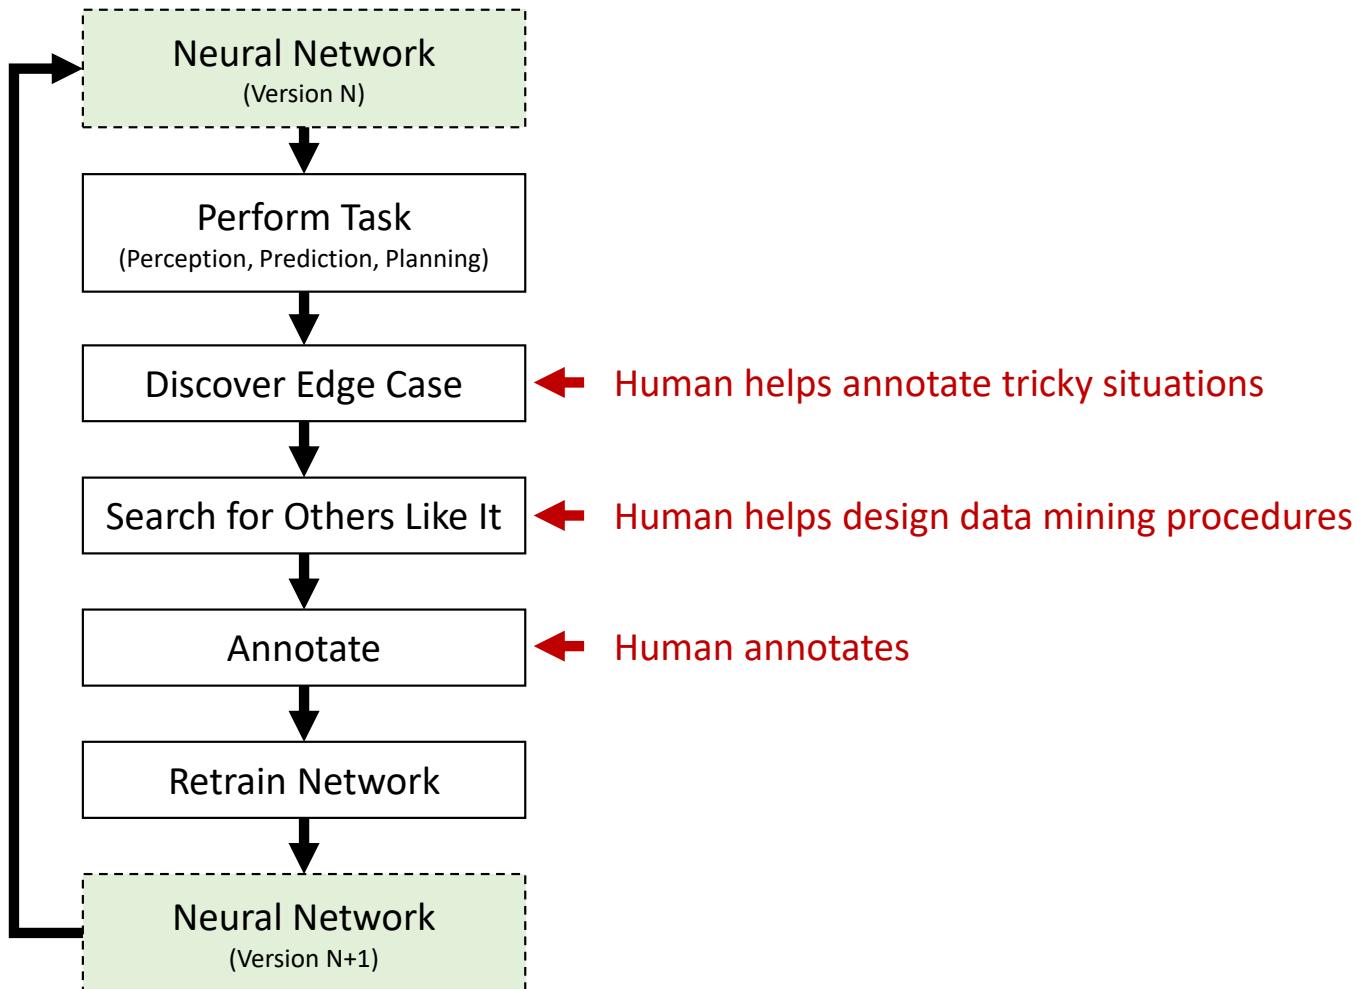
* Dates are for perspective and not as definitive historical record of invention or credit



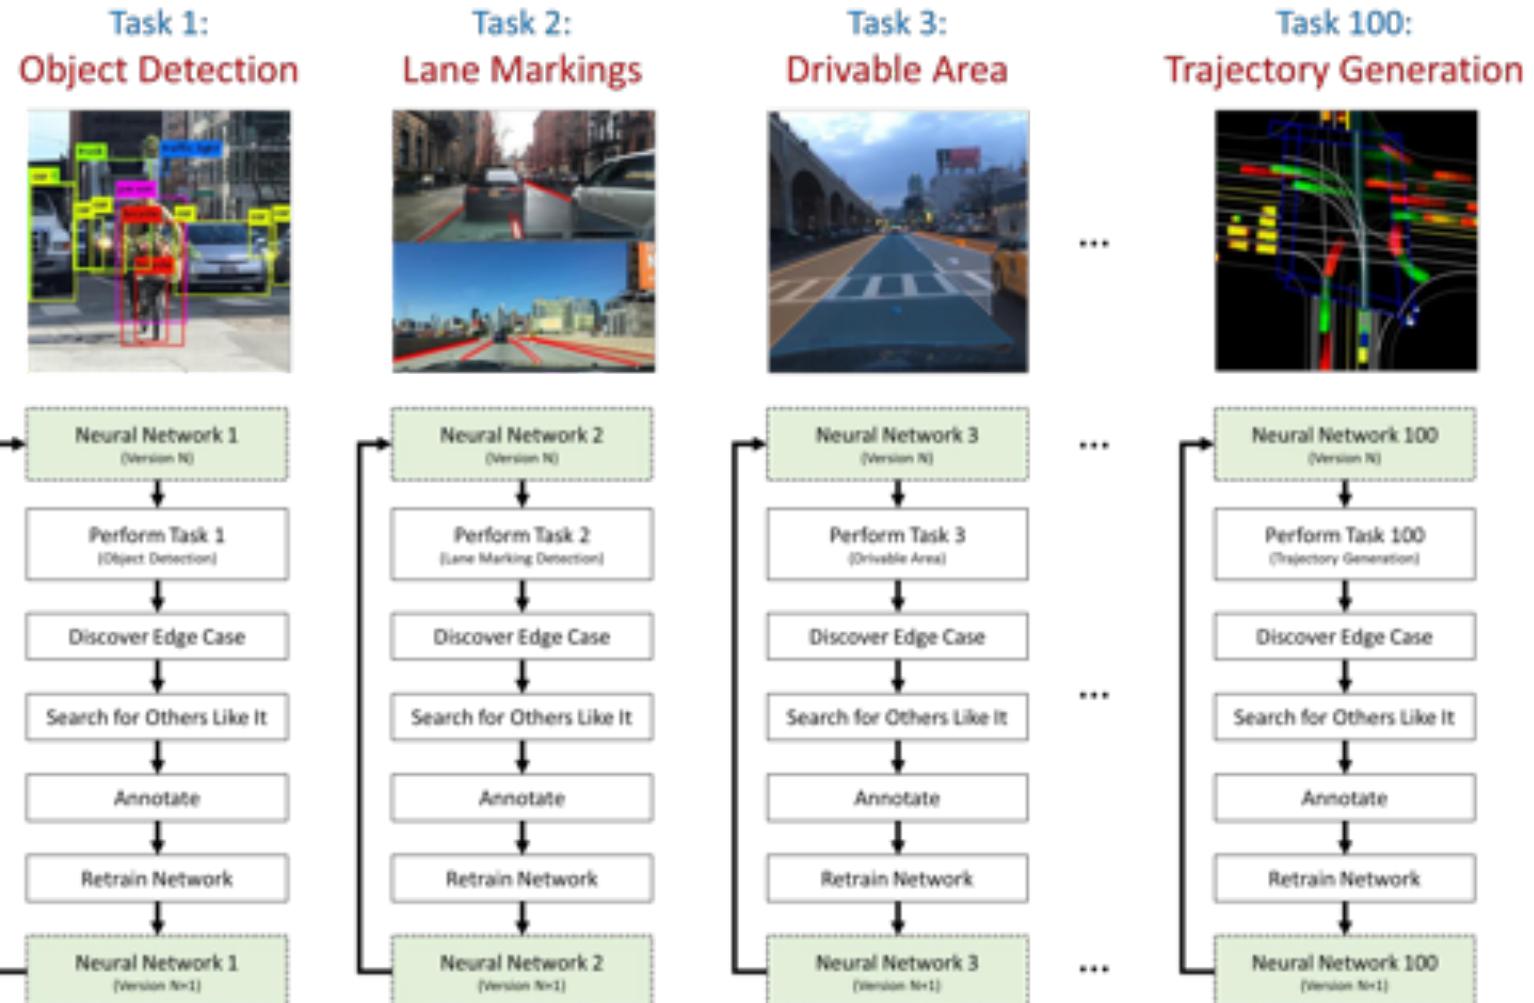
Massachusetts
Institute of
Technology

For the full list of references visit:
<http://bit.ly/deeplearn-sota-2020>

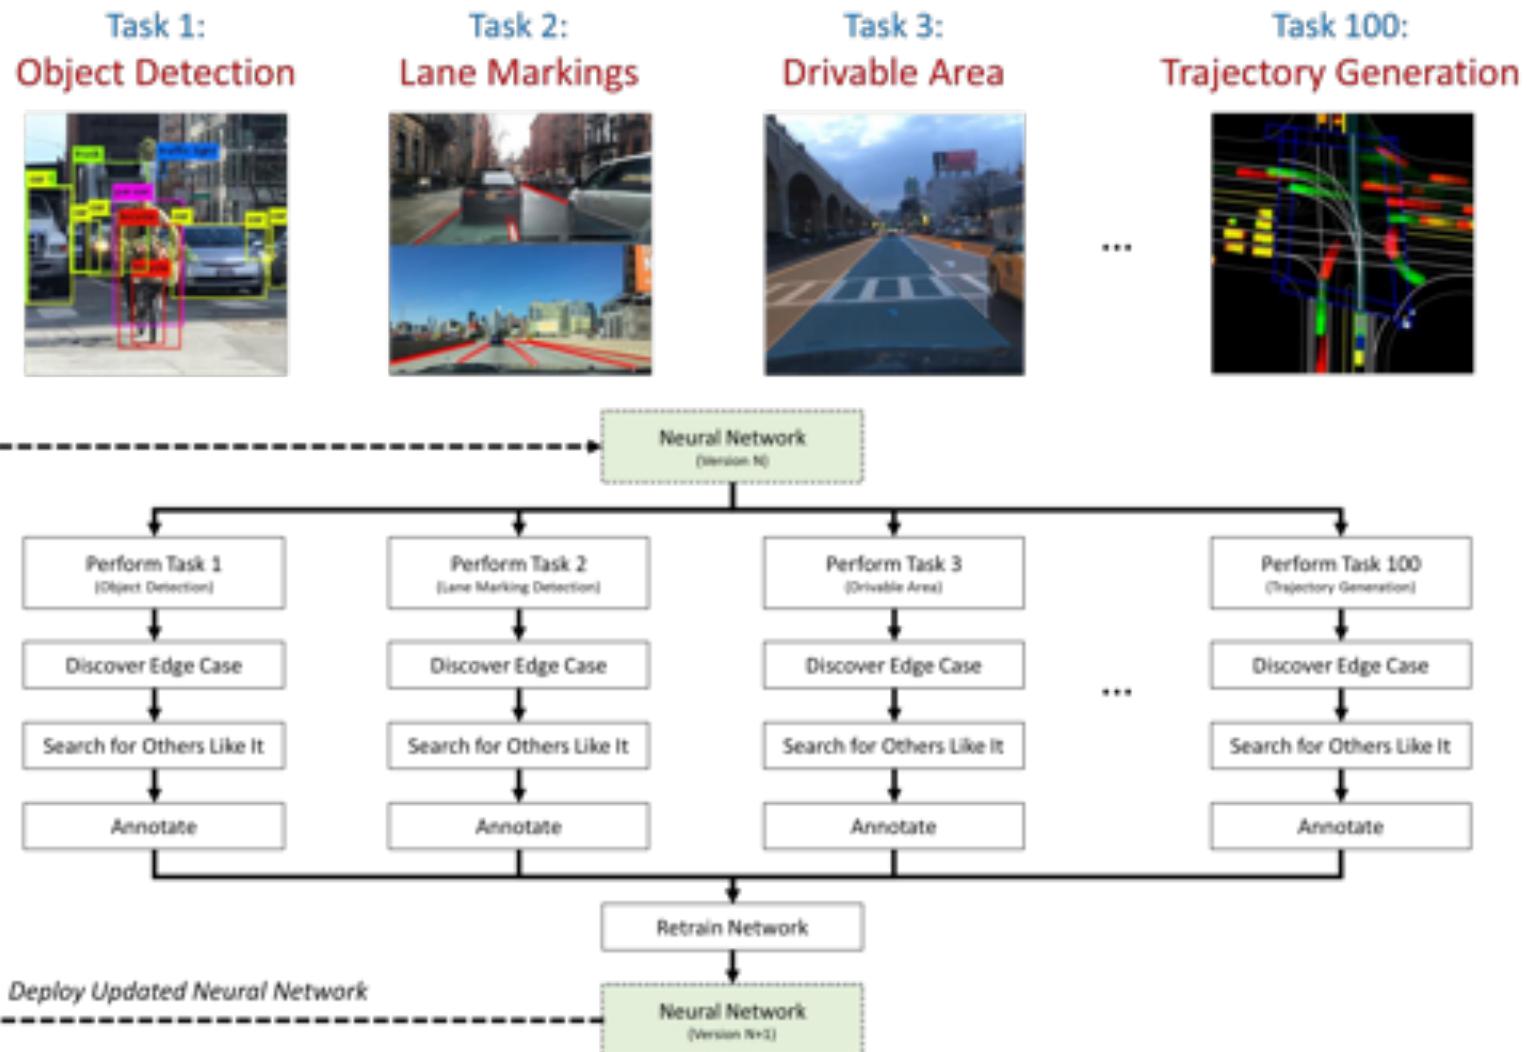
Active Learning Pipeline (aka Data Engine)



Single-Task Learning



Multi-Task Learning



Useful resources

Tensorflow tutorials: <https://www.tensorflow.org/tutorials>

The screenshot shows a web browser window displaying the TensorFlow tutorial website at <https://www.tensorflow.org/tutorials>. The page title is "Data augmentation". The left sidebar contains a navigation menu with sections like "TensorFlow tutorials", "Quickstart for beginners", "Quickstart for experts", "BEGINNER", "ML basics with Keras", "Load and preprocess data", "ADVANCED", "Customization", "Distributed training", and "Images" (which is expanded to show "Convolutional Neural Network", "Image classification", "Transfer learning and fine-tuning", "Transfer learning with TF Hub", "Data Augmentation" - which is selected and highlighted in blue, "Image segmentation", and "Object detection with TF Hub"). Below these are sections for "Text", "Audio", "Structured data", "Generative", "Model Understanding", "Reinforcement learning", and "tf.Estimator". At the bottom of the sidebar is a "Display a menu" button. The main content area starts with an "Overview" section describing data augmentation as a technique to increase training set diversity through random transformations like rotation. It mentions using Keras Preprocessing Layers and tf.image. Below this is a "Setup" section containing Python code for importing matplotlib, numpy, tensorflow, tensorflow_datasets, and tensorflow.keras.layers. The final section is "Download a dataset", which notes the use of the tf.flowers dataset and provides links for TensorFlow Datasets and the load_images tutorial.

Data augmentation

Table of contents

- Overview
- Setup
- Download a dataset
- Use Keras preprocessing layers
- Resizing and rescaling
- ...

Run in Google Colab View source on GitHub Download notebook

Overview

This tutorial demonstrates data augmentation: a technique to increase the diversity of your training set by applying random (but realistic) transformations such as image rotation. You will learn how to apply data augmentation in two ways. First, you will use [Keras Preprocessing Layers](#). Next, you will use [tf.image](#).

Setup

```
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds

from tensorflow.keras import layers
```

Download a dataset

This tutorial uses the [tf.flowers](#) dataset. For convenience, download the dataset using [TensorFlow Datasets](#). If you would like to learn about other ways of importing data, see the [load_images](#) tutorial.

https://www.tensorflow.org/tutorials/images/data_augmentation

Tensorflow datasets: <https://www.tensorflow.org/datasets/overview>

Screenshot of the Know Your Data Catalog interface showing a list of TensorFlow Datasets:

- vgg_face2**: 3,311,286 items. Preview images show faces. Buttons: See dataset, Explore in KYD.
- places365_small**: 2,168,460 items. Preview images show indoor and outdoor scenes. Buttons: See dataset, Explore in KYD.
- open_images_v4**: 1,910,098 items. Preview images show various objects. Buttons: See dataset, Explore in KYD.
- open_images_challenge2...**: 1,884,659 items. Preview images show objects. Buttons: See dataset, Explore in KYD.

Navigation and search bar at the top. Buttons: Catalog, STATS, RELATIONS, ITEM, ABOUT.

Screenshot of the Know Your Data Catalog interface for the **patch_cameliony** dataset:

Showing 327,680 of 327,680 items

Source features: patch_cameliony Source

The PatchCamelyon benchmark is a new and challenging image classification dataset. It consists of 327,680 color images (96 x 96px) extracted from histopathologic scans of lymph node sections. Each image is annotated with a binary label indicating presence of metastatic tissue. PCam provides a new benchmark for machine learning models: bigger than CIFAR10, smaller than Imagenet, trainable on a single GPU.

Less

label

Value	Count
0	163,862
1	163,818

split

Value	Count
test	2,768
train	262,144
validation	2,768

Preview images of histopathology slides labeled 0 or 1. Buttons: Add Filter, Select, Draw, Group by, Sort groups by, Sort items by, Order.

<https://knowyourdata-tfds.withgoogle.com>

Kaggle:

The screenshot shows the 'Dogs vs. Cats' competition page on Kaggle. The left sidebar includes links for Home, Competitions, Datasets, Code, Discussions, Courses, and More. Under 'Recently Viewed', there are links to 'Building Faces out of P...', 'Face landmarks with G...', 'Face Images with Mark...', 'Binary Classification', and 'Intro to Deep Learning'. The main content area displays the competition title 'Dogs vs. Cats' with a subdescription 'Create an algorithm to distinguish dogs from cats.' It shows a thumbnail of a dog and a cat, and a snippet of text: 'In this competition, you'll write an algorithm to classify whether images contain either a dog or a cat. This is easy for humans, dogs, and cats. Your computer will find it a bit more difficult.' Below this is a photo of a dog and a cat. At the bottom, there's a section titled 'The Asirra data set' with a note: 'Web services are often protected with a challenge that's supposed to be easy for people to solve, but...'. Navigation tabs at the bottom include Overview, Data, Code, Discussion, Leaderboard, Rules, and Team.

<https://www.kaggle.com/c/dogs-vs-cats/overview>

The screenshot shows the 'Building Faces Out of Parts' competition page on Kaggle. The left sidebar includes Home, Competitions, Datasets, Code, Discussions, Courses, and More. Under 'Recently Viewed', there is a link to 'Dogs vs. Cats'. The main content area displays a grid of 10 images showing faces with various facial features highlighted by colored bounding boxes (blue, yellow, green). To the right, a sidebar titled 'Version 25 of 25' shows a list of steps: 'copied from Exploring me Images and facial landmarks (>422 - 146)', 'Notebook', 'Building Faces Out Of Parts', 'Load The Data', 'Keep Only The Images With All 15 Keypoints...', 'Show Random Subset Of Images With All 15...', 'Define Crop Boundaries For The Eyes, Nose, ...', 'Plot The Resulting Bounding Boxes Of The Facial Parts', 'Define Target Size Of The Facial Parts', 'Create "Eyes", "Noses", "Mouths" From Keypoint Input [I]', 'Execution Info', 'Log', and 'Comments (0)'. A code editor window on the right contains Python code for calculating target sizes for facial parts based on mean keypoint distances.

<https://www.kaggle.com/selfishgene/building-faces-out-of-parts>

Over 2000 tutorials

Link: <http://bit.ly/36skFE7>

- Topics
 - Machine learning
 - Activation and Loss Functions
 - Bias
 - Perceptron
 - Regression
 - Gradient descent
 - Generative learning
 - Support vector machines
 - Backpropagation
 - Deep Learning
 - Optimization
 - Long Short Term Memory
 - Convolutional Neural Networks
 - Recurrent Neural Nets (RNNs)
 - Reinforcement Learning
 - Generative Adversarial Networks
 - Multi-task Learning
 - NLP
 - Word Vectors
 - Encoder-Decoder
 - TensorFlow
 - PyTorch
-

<https://github.com/kmario23/deep-learning-drizzle>

Links to great youtube lectures

The screenshot shows the GitHub repository page for `kmario23/deep-learning-drizzle`. The repository has 8.3k stars and 2.4k forks. The README.md file contains the following content:

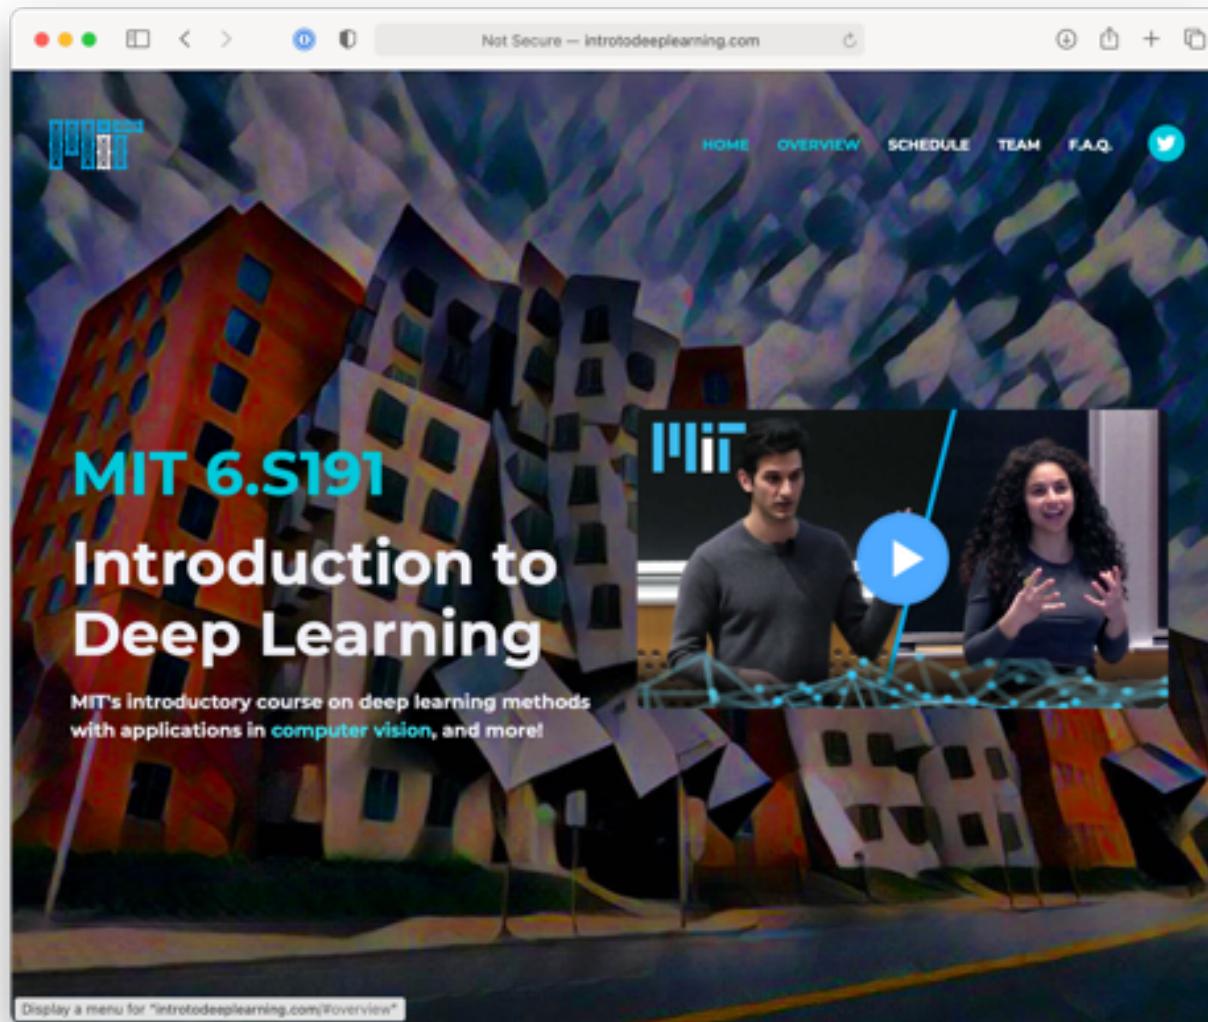
```
Deep Learning Drizzle 🎉🎈

"Read enough so you start developing intuitions and then trust your intuitions and go for it!" 🍻
Prof. Geoffrey Hinton, University of Toronto

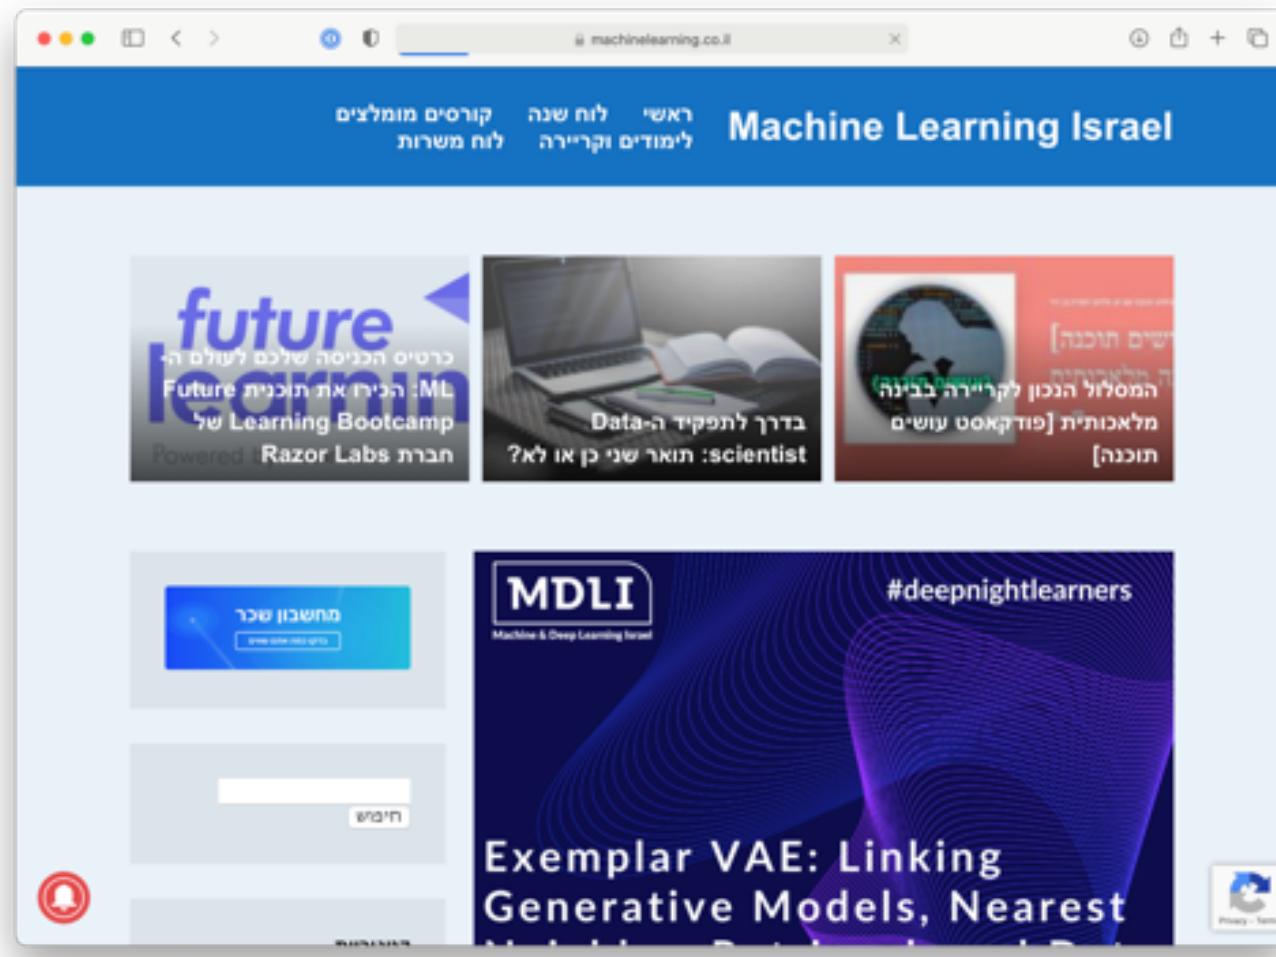
-----
Contents
-----
Deep Learning [Deep Neural Networks] [!] Probabilistic Graphical Models [!]
Machine Learning Fundamentals [!] Natural Language Processing [!]
Optimization for Machine Learning [!] Automatic Speech Recognition [!]

Topics
-----
deep-learning-drizzle.github.io
machine-learning
natural-language-processing
deep-neural-networks
reinforcement-learning computer-vision
deep-learning optimization
machine-translation
deep-reinforcement-learning
medical-imaging speech-recognition
artificial-neural-networks
pattern-recognition
probabilistic-graphical-models
bayesian-statistics
artificial-intelligence-algorithms
visual-recognition
graph-neural-networks
Readme
```

9.	CS231n: CNNs for Visual Recognition	Andrej Karpathy, Stanford University	CS231n	YouTube-Lectures (Academic Torrent)	2016
10.	Neural Networks	Hugo Larochelle, Université de Sherbrooke	Neural-Networks	YouTube-Lectures (Academic Torrent)	2016
11.	CS224d: Deep Learning for NLP	Richard Socher, Stanford University	CS224d	YouTube-Lectures (Academic Torrent)	2016
12.	CS224n: NLP with Deep Learning	Richard Socher, Stanford University	CS224n	YouTube-Lectures	2017
13.	CS231n: CNNs for Visual Recognition	Justin Johnson, Stanford University	CS231n	YouTube-Lectures (Academic Torrent)	2017
14.	Topics in Deep Learning	Ruslan Salakhutdinov, CMU	10707	YouTube-Lectures	F2017
15.	Deep Learning Crash Course	Leo Isikdogan, UT Austin	None	YouTube-Lectures	2017
16.	Deep Learning and its Applications	François Fleuret, Trinity College Dublin	EE4C16	YouTube-Lectures	2017
17.	Deep Learning	Andrew Ng, Stanford University	CS230	YouTube-Lectures	2018
18.	UvA Deep Learning	Efstratios Gavves, University of Amsterdam	UvA-DLC	Lecture-Videos	2018
19.	Advanced Deep Learning and Reinforcement Learning	Many legends, DeepMind	None	YouTube-Lectures	2018
20.	Machine Learning	Peter Bloem, Vrije Universiteit Amsterdam	MLVU	YouTube-Lectures	2018
21.	Deep Learning	François Fleuret, EPFL	EE-59	Video-Lectures	2018



<http://introtodeeplearning.com>



<https://machinelearning.co.il>

<https://github.com/ChristosChristofidis/awesome-deep-learning>

A great list of resources..

The screenshot shows a GitHub repository page for 'Awesome Deep Learning'. The repository has 472 stars and was last updated 6 months ago. The README.md file contains a table of contents with links to various categories: Books, Courses, Videos and Lectures, Papers, Tutorials, Researchers, Websites, Datasets, Conferences, Frameworks, Tools, Miscellaneous, and Contributing.

master

ChaiBapchya Merge pull request #186 from SauravMaheshkar/ma... on Nov 30, 2020 472

README.md 6 months ago

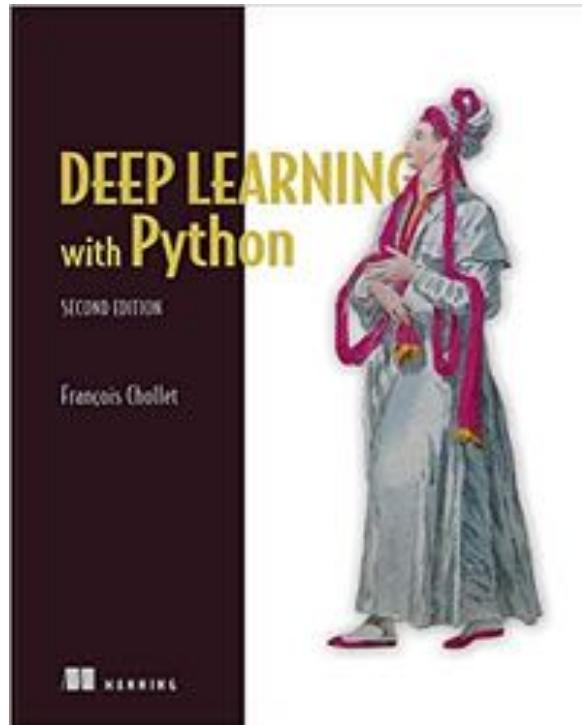
README.md

Awesome Deep Learning

Table of Contents

- Books
- Courses
- Videos and Lectures
- Papers
- Tutorials
- Researchers
- Websites
- Datasets
- Conferences
- Frameworks
- Tools
- Miscellaneous
- Contributing

Display a menu



Highly recommended book (second edition expected later this year..)

<https://www.manning.com/books/deep-learning-with-python-second-edition>

<https://github.com/fchollet/deep-learning-with-python-notebooks>

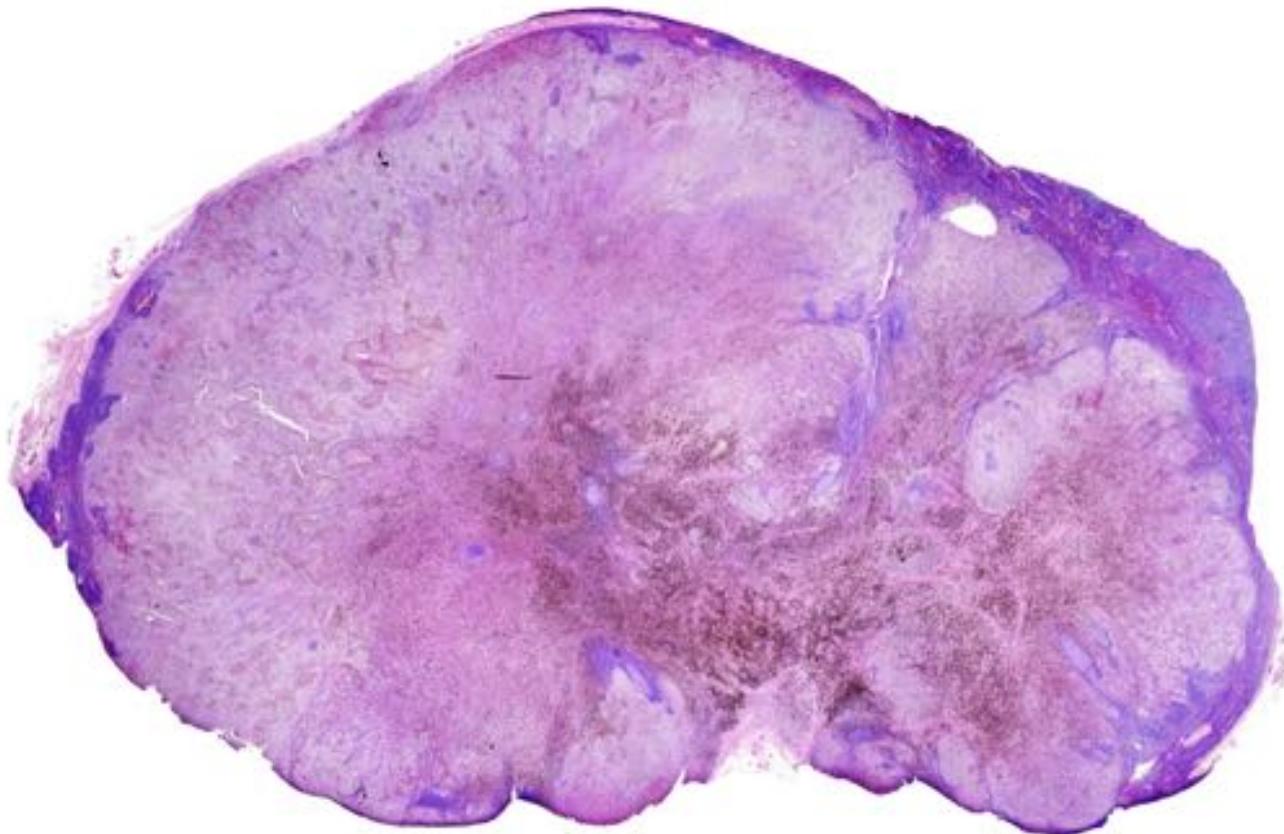
(Great **ipynb** examples code available online)

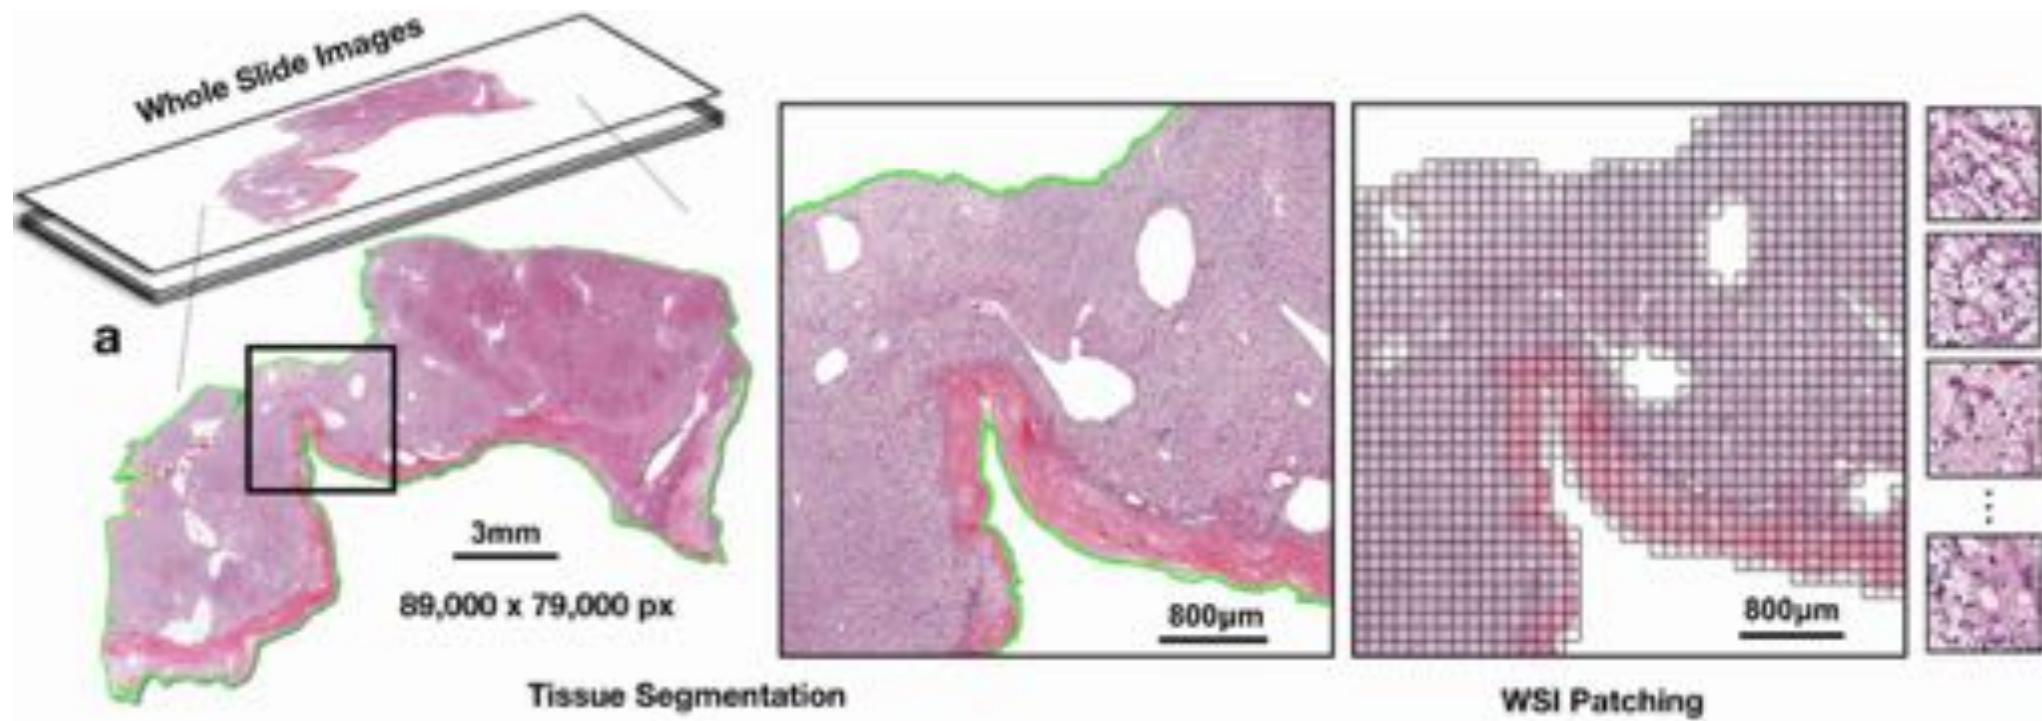
Table of contents

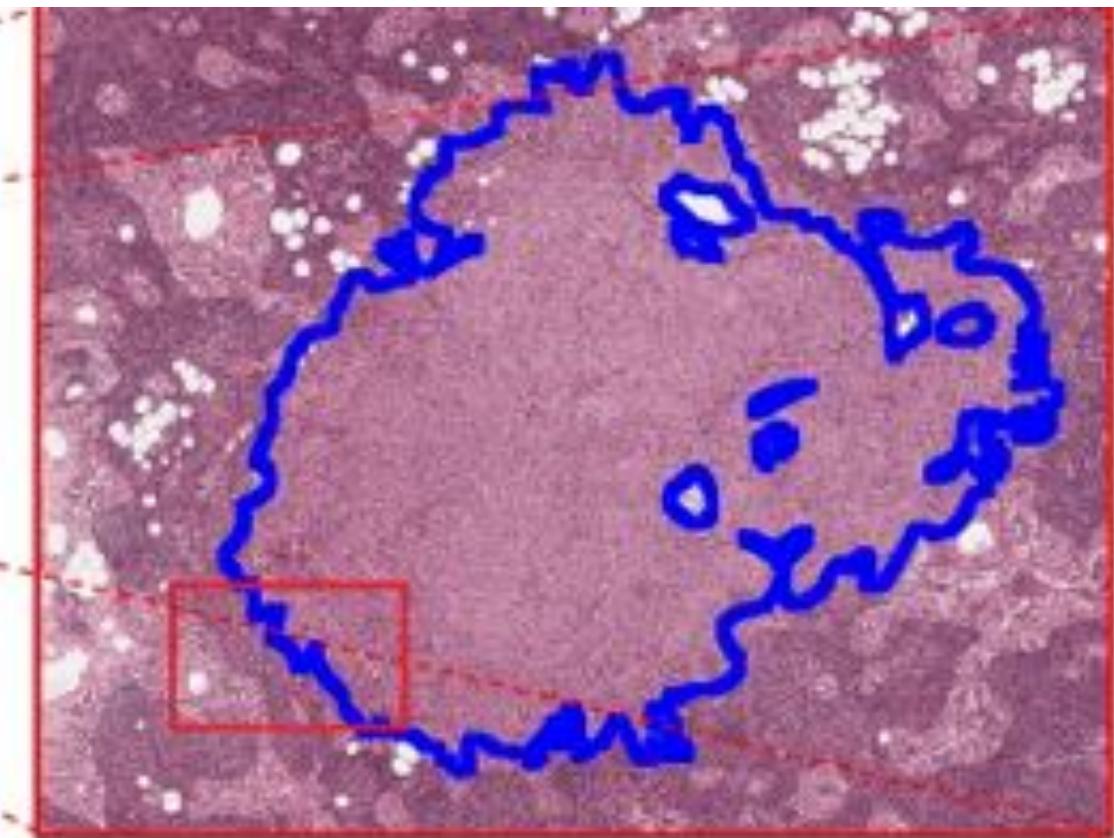
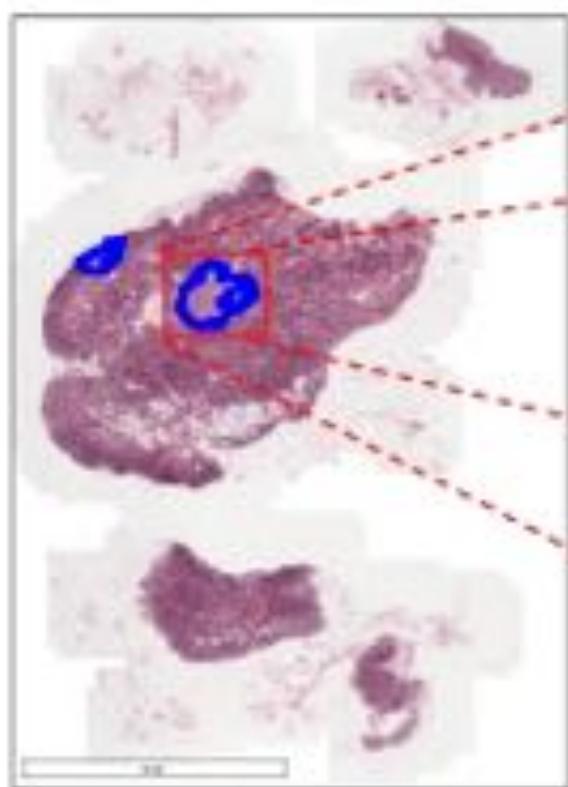
- Chapter 2:
 - 2.1: A first look at a neural network
- Chapter 3:
 - 3.5: Classifying movie reviews
 - 3.6: Classifying newswires
 - 3.7: Predicting house prices
- Chapter 4:
 - 4.4: Underfitting and overfitting
- Chapter 5:
 - 5.1: Introduction to convnets
 - 5.2: Using convnets with small datasets
 - 5.3: Using a pre-trained convnet
 - 5.4: Visualizing what convnets learn
- Chapter 6:
 - 6.1: One-hot encoding of words or characters
 - 6.1: Using word embeddings
 - 6.2: Understanding RNNs
 - 6.3: Advanced usage of RNNs
 - 6.4: Sequence processing with convnets
- Chapter 8:
 - 8.1: Text generation with LSTM
 - 8.2: Deep dream
 - 8.3: Neural style transfer
 - 8.4: Generating images with VAEs
 - 8.5: Introduction to GANs

Final Project

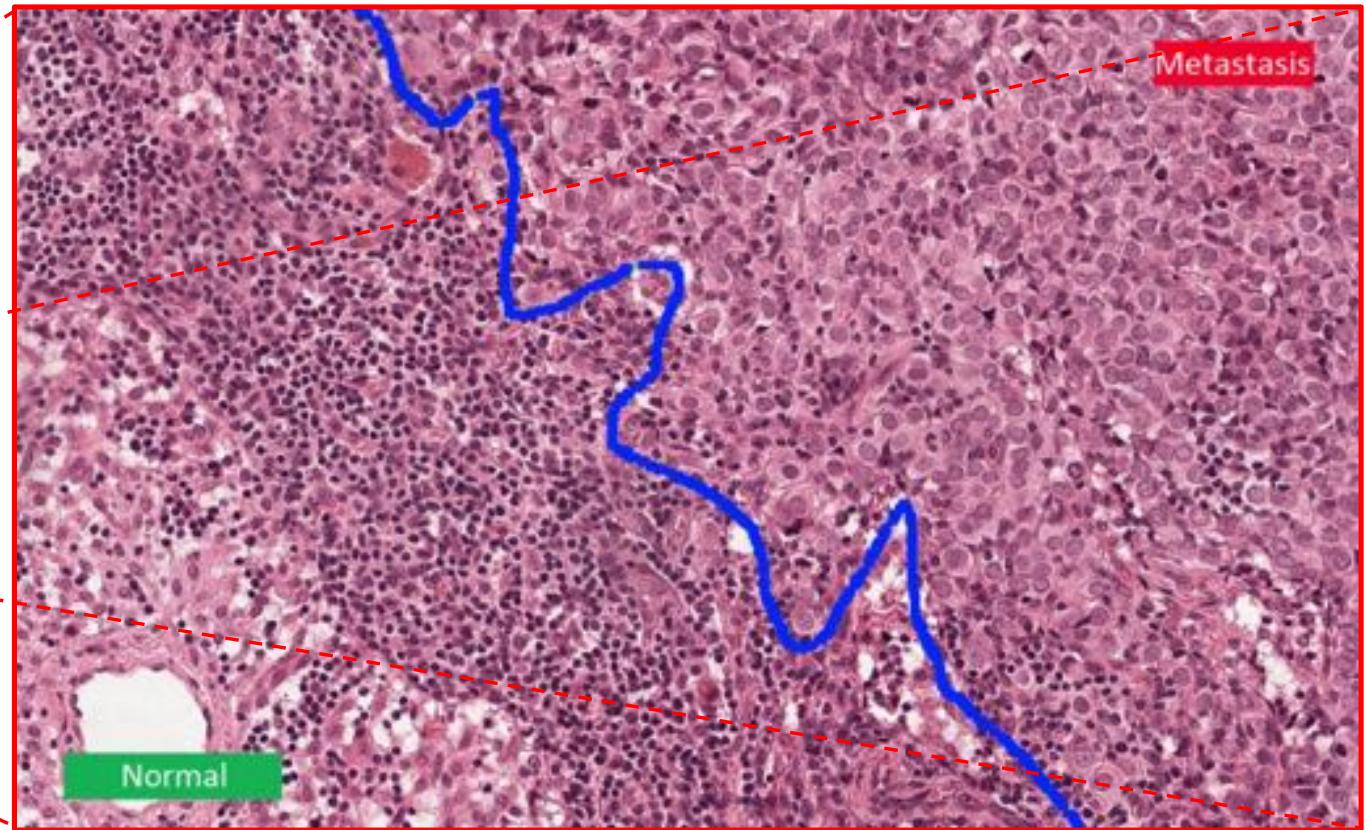
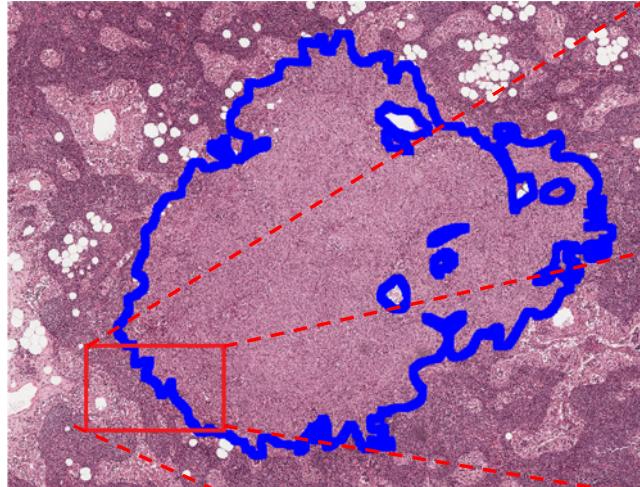
Classification of histopathology images

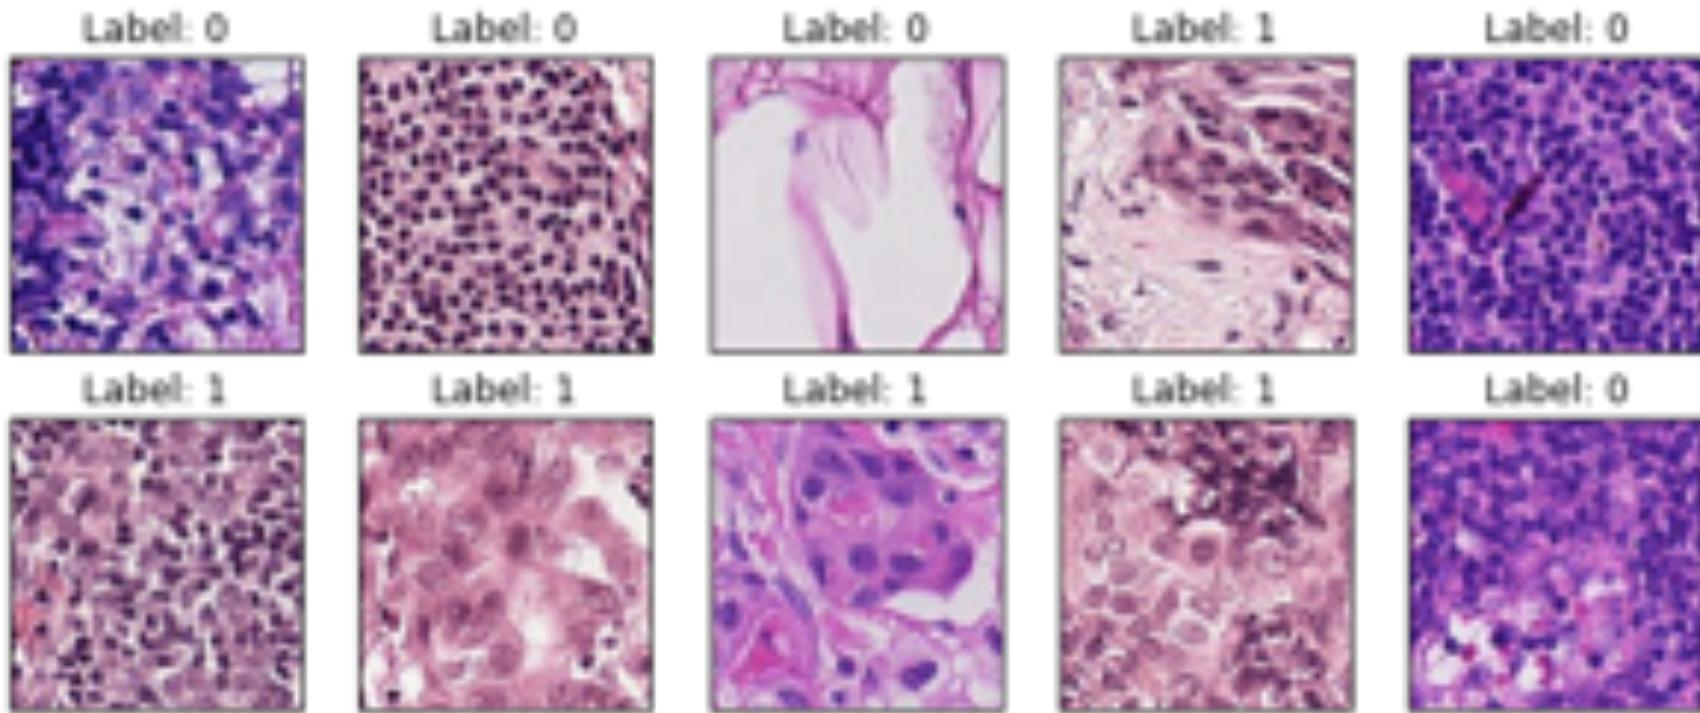






The data in this challenge contains a total of 400 whole-slide images (WSIs)





Colorectal histology dataset:

Classification of textures in colorectal cancer histology.

Each example is a 150 x 150 x 3 RGB image of one of 8 classes

The screenshot shows the 'Catalog' tab of the 'colorectal_histology' dataset in the 'Know Your Data' interface. It displays 5,000 items, each being a 150 x 150 x 3 RGB image of one of 8 classes. The interface includes filters for 'label' (adipose, complex, debris, empty, lympho, mucosa, stroma, tumor), sorting options, and a 'Cloud Vision' section for automatically computed text labels. A sidebar provides details about the dataset, stating it contains 5,000 items and is used for classifying textures in colorectal cancer histology.

https://knowyourdata-tfds.withgoogle.com/#tab=STATS&dataset=colorectal_histology

(The dataset for the final project)

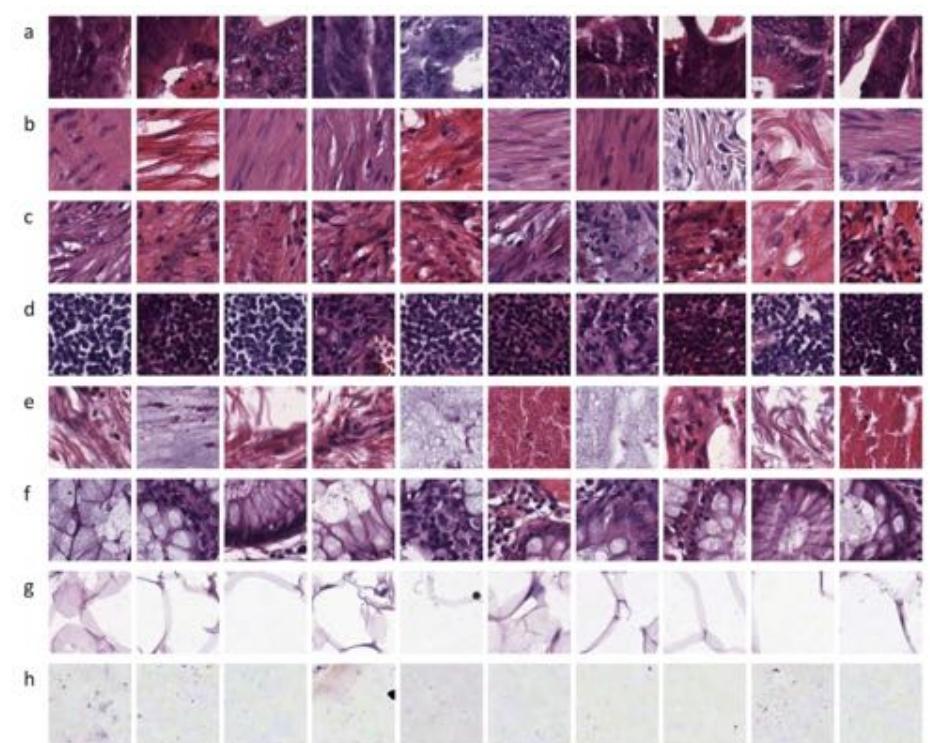


Figure 1. Representative images from our dataset. Here, the first 10 images of every tissue class in our dataset are shown. They represent the wide variation of illumination, stain intensity and tissue textures present in routine histopathological images. Images were extracted from 10 independent samples of colorectal cancer (CRC) primary tumours. (a) tumour epithelium, (b) simple stroma, (c) complex stroma (stroma that contains single tumour cells and/or single immune cells), (d) immune cell conglomerates, (e) debris and mucus, (f) mucosal glands, (g) adipose tissue, (h) background.

SCIENTIFIC REPORTS



OPEN

Multi-class texture analysis in colorectal cancer histology

Jakob Nikolas Kather^{1,2}, Cleo-Aron Weis¹, Francesco Bianconi³, Susanne M. Melchers⁴, Lothar R. Schad², Timo Gaiser¹, Alexander Marx¹ & Frank Gerrit Zöllner²

Received: 02 March 2016

Accepted: 25 May 2016

Published: 16 June 2016

Automatic recognition of different tissue types in histological images is an essential part in the digital pathology toolbox. Texture analysis is commonly used to address this problem; mainly in the context of estimating the tumour/stroma ratio on histological samples. However, although histological images typically contain more than two tissue types, only few studies have addressed the multi-class problem. For colorectal cancer, one of the most prevalent tumour types, there are in fact no published results on multiclass texture separation. In this paper we present a new dataset of 5,000 histological images of human colorectal cancer including eight different types of tissue. We used this set to assess the classification performance of a wide range of texture descriptors and classifiers. As a result, we found an optimal classification strategy that markedly outperformed traditional methods, improving the state of the art for tumour-stroma separation from 96.9% to 98.6% accuracy and setting a new standard for multiclass tissue separation (87.4% accuracy for eight classes). We make our dataset of histological images publicly available under a Creative Commons license and encourage other researchers to use it as a benchmark for their studies.

https://www.researchgate.net/publication/303998214_Multi-class_texture_analysis_in_colorectal_cancer_histology/link/5762810e08ae0c8c2553b63b/download

Know Your Data Catalog colorectal_histology_large

2 of 2 matches Contains Q color Done

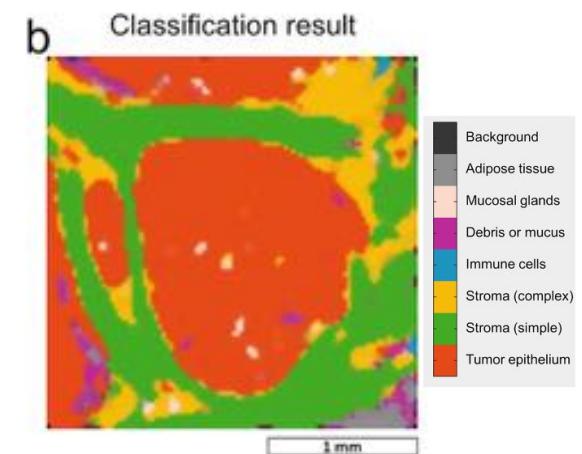
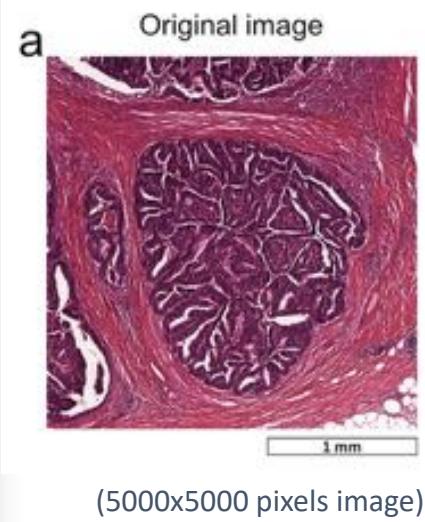
ID: CRC-Prim-HE-07_APPLICATION.tif

Add Filter Showing 10 of 10 items

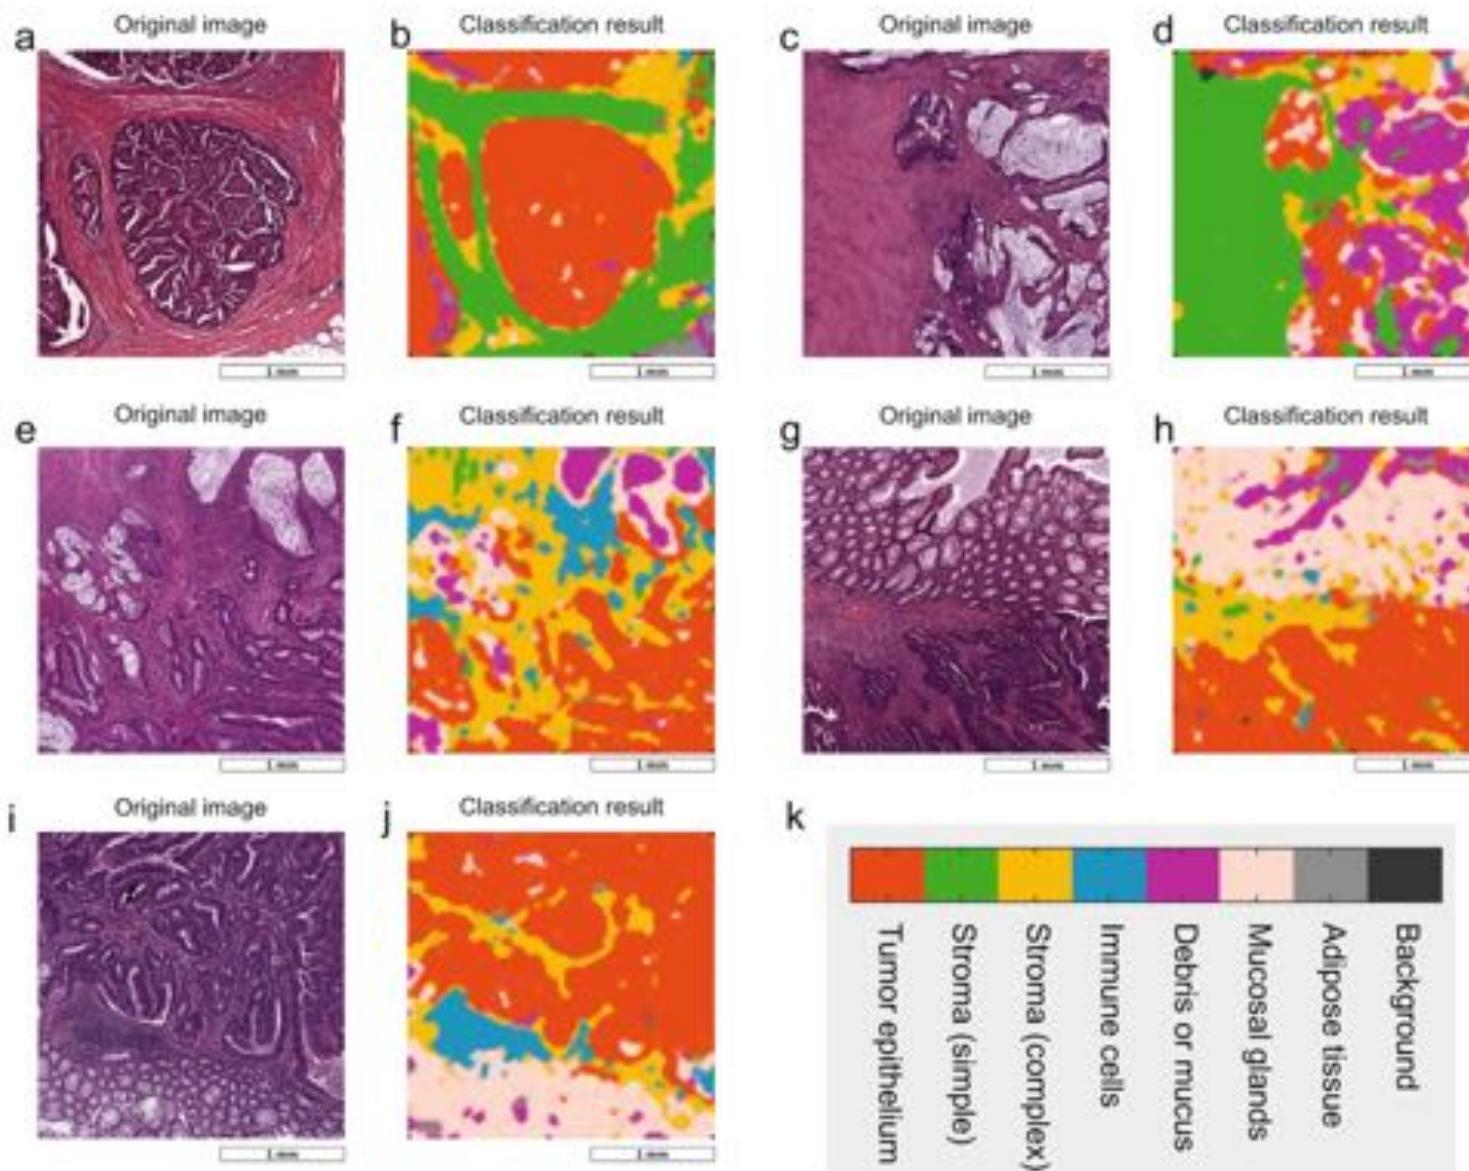
Select Draw Group by Sort groups by Sort items by Order

Feature	
image_content_metadata	
aspect_ratio	1
format	PNG
megapixel_resolution	23.84
mode	RGB
pixel_height	5000
pixel_width	5000
image_quality	-11, 16

Display a menu



https://knowyourdata-tfds.withgoogle.com/#tab=STATS&dataset=colorectal_histology_large



Project goals

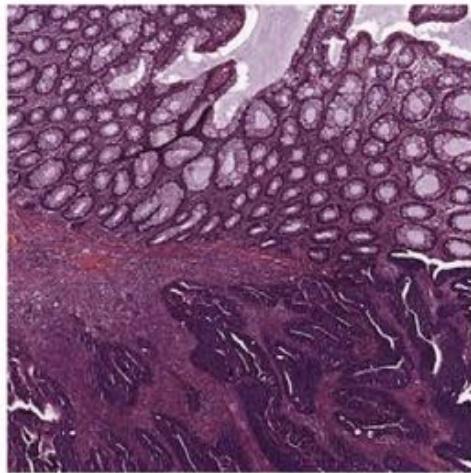
- Train a ConvNet classifier
(either design your own or use a pretrained network + transfer-learning/fine-tuning)
- Use train/validation protocols
(e.g., 90% data for training, 10% for validation)
- Try several schemes, and compare your results in a table
(e.g., different architectures, optimizers, preprocessing & data-augmentation, ...)

Project goals

- Using your best model:
 - Plot its confusion matrix.
 - Use its *one-before-last* layer (its “features vector” layer), and **PCA** or **t-SNE** to visualize the labelled-data distribution in **2-D**.
 - Pick each of the large 5000x5000 pixels images, slice them into smaller 150x150 patches, and use your model to classify each patch. Visualize the classification results (on-top or side-by-side of the large 5000x5000 images) in two ways:
 - Color each patch (in the large image) using the predicted class label (8 colors for 8 classes).
 - Color each patch using the probability value of the “tumor ep.” class (a ‘heatmap’ result).

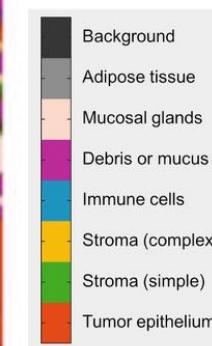
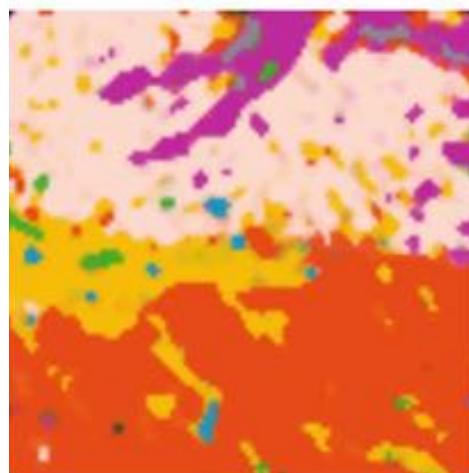
Visualization examples:

Original image

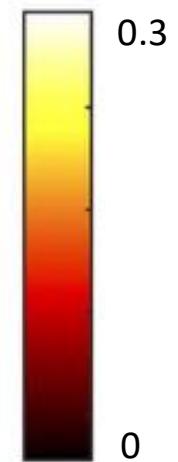
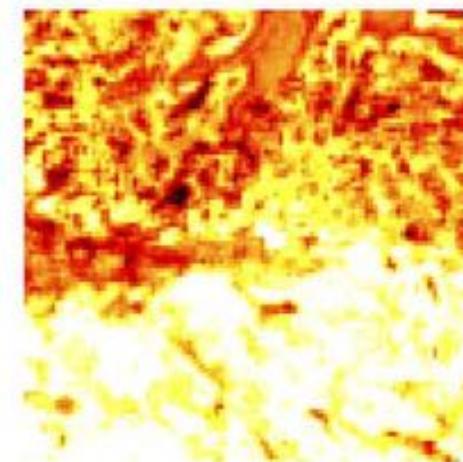


5000x5000 pixels

Classification result



Tumor epithelium



Hint: you can use some overlap (e.g., 50%) when dividing the large image into patches, to get a nicer results.

הנחיות הגשה

- הגשה בלבד או בזוגות.
- **מועד אחרון להגשת הפרויקט: 3 יולי 2021**
- יש להגיש קוד **markup מסודר**, כולל הערות, וכל השלבים השונים.
- יש להסבירCells markup כדי להסביר בכל שלב מה ביצעתם (עברית/אנגלית)
- יש להציג גרפים ותוצאות מסודרות (בגודל סביר, כולל כתוב לצירים השונים וכו') לכל אורך הפרויקט.
- לדאוג **לנקות מהקובץ הסופי** את כל הקוד/תוצאות הלא רלוונטיים.
- בנוסף, יש להגיש **פוסטר** (בפורמט PPT), לפי הנחיות באתר:

<https://www.hit.ac.il/sciences/projects>

קוד אטי

- ניתן להעזר בחברי הקורס, ובחומרים ברשותם כ摹בן, **כל עוד העבודה המוגשת בסוף מבוססת על עבודה עצמאית ומוקנית של חברי הקבוצה.**
- יש לצרף קישורים לחומרים ברשות שנעזרתם בהם במהלך העבודה.

Add markup sections to your final *ipynb*:

Part1

In this part, we are going to modify the original video by adding the at the end of half a minute of original video after being rotated.

We rotated the video to test the function of saving the path.

```
def translate(dx,dy):
    M = np.eye(3)
    M[:,2] = [dx,dy,1]
    return M

def rotate(angle):
    M = np.eye(3)
    sin, cos = np.sin(angle), np.cos(angle)
    M[1:2,:2] = [[cos,-sin],[sin,cos]]
    return M
```

The function for rotate the image

```
def RotateImage (frame, angle):
    h,w,_ = frame.shape

    T1 = translate(w/2, h/2)
    T2 = translate(-w/2, -h/2)

    angle =angle/180*2*np.pi
    M = rotate(angle)
    P = (T1.dot(M)).dot(T2)
    frame_warp = cv2.warpPerspective(frame.copy(), P, (w,h),borderValue=(255,255,255))

    return frame_warp
```

Describe your functions:

```
handle_canny_and_refresh_good_features(apply_canny, threshold, F0, feature_params)
```

Used to apply Canny algo (optional) and find corners using Shi-Tomasi.

apply_canny - Whether to apply Canny alg.

threshold - Canny's threshold parameter.

F0 - The frame to modify.

feature_params - Parameters for the Shi-Tomasi alg.

```
def handle_canny_and_refresh_good_features(apply_canny, threshold, F0, feature_params):
    # Use Canny to find edges.
    if apply_canny:
        F0_gray = cv2.Canny(F0, threshold, 8, 1)

    # Make sure to return the frame gray.
    else:
        F0_gray = cv2.cvtColor(F0, cv2.COLOR_BGR2GRAY)
    pts0 = cv2.goodFeaturesToTrack(F0_gray, mask = None, **feature_params)

    # Squeeze to pairs of coords.
    if pts0 is not None and len(pts0) > 0:
        pts0 = pts0.squeeze() # (n,1,2) -> (n,2)

    return F0_gray, pts0
```

```
overlay_image_alpha(img, img_overlay, pos, alpha_mask)
```

An aux function to display image on top of an image.

Taken from here: <https://stackoverflow.com/a/45118011/3450476>

And modified it a bit to fit the project needs (removed alpha-channel manipulation so the images won't be displayed semi-transparent)

Add comments:

```
while True:
    ret, frame = cap.read()
    if ret==True:

        # check if we are currently underwater using hue values (underwater mainly blue)
        img_hsv = cv2.cvtColor(frame, cv2.COLOR_BGR2HSV)
        hue,saturation,values = cv2.split(img_hsv)
        (values,counts) = np.unique(hue,return_counts=True)
        ind=np.argmax(counts)
        output_frame = frame.copy()

        if 90 <= values[ind] <= 110:
            # we are underwater

            # mask lane bounds - clip to / swimming range !
            bounds_mask = np.zeros((h, w), dtype=np.uint8)
            cv2.fillPoly(bounds_mask, points, (255))
            clipped = cv2.bitwise_and(frame,frame,mask = bounds_mask)

            # mask lane by hsl - dark blue
            hsv_cropped = cv2.cvtColor(clipped, cv2.COLOR_BGR2HSV)
            lower = np.array([100, 100, 100])
            upper = np.array([120, 255, 255])
            color_mask = cv2.inRange(hsv_cropped, lower, upper)
            imask = color_mask>0
            masked_frame = np.zeros_like(frame, np.uint8)
            masked_frame[imask] = frame[imask]

            # convert to grayscale
            gray = cv2.cvtColor(masked_frame, cv2.COLOR_BGR2GRAY)

            # convert to black and white using thershold
            _,binary = cv2.threshold(gray,30,255,cv2.THRESH_BINARY)

            # erode noise
            eroded = cv2.erode(binary,kernel,iterations = 1)
```



Efficient Convolutional Neural Networks for Cloud Detection in Satellite Images

John Merriman Sholar

CS 231N: Convolutional Neural Networks for Visual Recognition



Problem

Overview

- Goal: identify clouds in satellite images
- Motivation: removing ("masking out") clouds important for downstream analysis for satellite imagery.
- Task: image segmentation with classes including three clouds, three clouds, shadow, water, snow, and land.
- Evaluation: per-pixel classification accuracy, CE loss, qualitative analysis of output, classification speed.

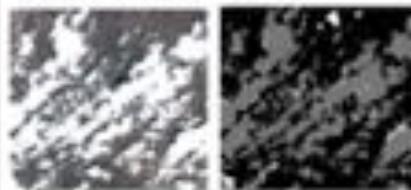
Subtasks

- We consider several permutations of the problem definition.
- Given full-resolution image segmentation using all 13 classes, 2-class, or only using RGB.
 - Output classes we consider 2-class, 3-class, and 6-class classification problems, defined as follows:

2 Class	3 Class	6 Class
Cloud, Non-Cloud	Cloud, Cloud, Shadow	Cloud, Cloud, Shadow, Water, Snow, Land

Data

- 80 Sentinel-2 satellite images = 10,000 by 10,000 px. each.
- Segmented area = 120,000 tiles of size 224 by 224 px.
- Sentinel-2 images have 13 spectral bands, surface from 3 (RGB), RGB frequencies range from 400nm to 490nm, with 30 nm range from 2,100nm to 2,400nm, with more information (e.g. infrared, short-wave infrared).
- Satellite imagery is very different from "traditional" imagery (top-down, spatial variations, lack of overall context, noise), making transfer learning or pre-trained classifiers less useful.
- Sentinel-2 ships with a proprietary cloud mask (false), which provides noisy but useful training data.



Baseline Methods

Precious Approaches

- Precious approaches are decision trees and linear classifiers, emphasizing speed.

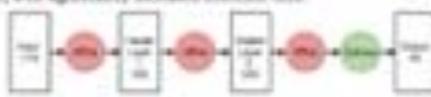
Pixel-Level Decision Trees

- A variation of previous research by Hollingshead et al.
- Extremely fast, but erroneous for problems more advanced than the binary problem.
- Used for a form of transfer learning: no ground truth data for CNNs.

- By applying decision trees, using the Sentinel-2 cloud mask as an extremely noisy label.

Pixel-Level MLP

- An expansion on the pixel-level classification using decision trees. Ultimately no more accurate than decision trees, likely due to the same inability to take advantage of spatial correlations, with significantly increased inference time.



Methods

Fully Convolutional Networks (Long et al.)

- A modification of the AlexNet (Krizhevsky et al.)
- Architecture which replaces the fully connected layers in favor of fully convolutional layers, and with a single final activation. (Below are: architecture, loss, and accuracy)



Deconvolutional Neural Networks (Ostaf et al.)

- Parallel convolutional and deconvolutional structure.
- Convolutional first half initialized using IJCVRC-pretrained VGG-16

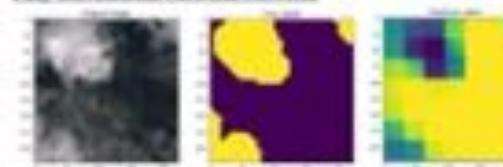


Results

Pixel-Level Decision Trees



Fully Convolutional Networks (Sholar et al.)



Deconvolutional Neural Networks (Ostaf et al.)

- We are unsuccessful in using deconvolutional neural networks. This is likely a result of several factors. First, for the RGB-only task, the VGG-16 ConvNet is pretrained on IJCVRC data, which does not parallel Sentinel-2 data. For the 13-band task, training may simply require more computational resources.

Inference Speed and Accuracy

	Inference Accuracy (ITL Testset)	Inference Speed
Decision Trees	0.60	10 million pixels/second
Fully Conv. Networks	0.62	1 million pixels/second
Deconv. Networks	0.60	#/sec pixels/second

Discussion and Future Work

- Pixel-level decision trees achieve visually satisfactory results and high numerical accuracy for the binary problem, but fail on more difficult classes, such as shadow.
- Fully convolutional neural networks produce coarse output, but could be refined by using a series of deconvolutional layers – likely the most stable next step.
- Deconvolutional networks have strong potential for this task, but must be pretrained on IJCVRC data, which is significantly different from satellite imagery. Future work might include training a DNN and reusing it on Sentinel-2 data.
- Satellite images are extremely large; processing time is key. Future work might include application of SegNet (Shen et al.) architectures to Sentinel-2.

Photoshop 2.0: Generative Adversarial Networks for Photo Editing

Hanmin Kim, Michelle Zhang, Brandon Yang
CS231N Spring 2017

Background

- Generative Adversarial Networks (GANs) generate images from a min-max game between generator and discriminator.
- Recent research has focused on improving GAN architectures for training and image quality, such as WGAN.
- We focus on building an architecture to directly train autoencoders for GANs to better understand the learned latent space representation of these networks.
- Previous research into conditional GANs conditions GAN training on labeled data for more specific image generation.
- Preliminary results in the DCGAN paper illustrate the potentials of latent space traversal and selective dropout for tuning image generation and semantic understanding.

Dataset

- The celebA dataset contains over 200,000 face images of various celebrities.
- Labeled with 40 different binary attributes for facial features, including hair color, gender, relative age, and whether or not the person is smiling.

Models

- To train a GAN auto-encoder we first trained the DCGAN model for 25 epochs and used the trained generator weights as our encoder.
- We experimented with three approaches to encoder networks with the L2 similarity metric:
 - Fully connected encoder model baseline
 - Deep convolutional encoder
 - Transfer learning encoder with weights from the trained discriminator from the DCGAN model.
- We then experimented with different loss functions to produce more realistic encodings: L1, L2, and SSIM.

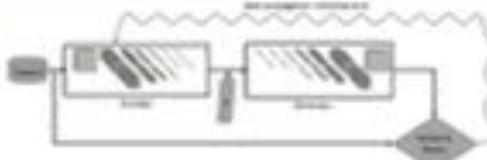


Figure 1: GAN Autoencoder and Generator network overview



Figure 2: Encoder network architectures

Qualitative Results

Training the extended transfer learning encoder:



Using GAN autoencoder to create smiling image



Evaluation



Figure 3: Training loss curves for transfer learning encoder

	L2 Loss	L1 Loss	SSIM Loss
FC Encoder	468.239	-	-
DeepConv Encoder	525.860	-	-
Transfer Learning Encoder	643.906	2995.321	0.293
Extended Transfer Learning Encoder	781.279	-	-

Figure 4: Loss results

Conclusions and Future Work

- GAN autoencoders can be effectively created by transfer learning from the discriminator network.
- Manipulating encoded images in the latent vector space with simple vector operations to generate images.
- Our model seems limited by inherent biases in the generator and dataset (female faces), which can perhaps be addressed with different generator approaches.
- In the future, we want to explore better quantitative metrics for generated image quality and possible end-to-end architectures for image processing.
- Explore different operations in the latent space and selective dropout and effects on generated images.

Bibliography

1. J.J. Kindermans, J. Puget-Jones, M. Mirza, B. Xu, D. Warde-Farley, S. Osuri, A. Courville, and Y. Bengio. Generative adversarial networks. 2014.
2. G. Huang. Photo-editing with generative adversarial networks. Apr. 2017.
3. F. Fleuret. Neural face images. *Pattern Recognition Letters* 30(10):1470–1476, 2009.
4. A. P-L. Lucas, S. A. Sandercock, and C. Whistler. Autoencoding faces with a generative latent space. *CvPR*, abs/1112.0948, 2011.
5. Z. Liu, P. Luo, X. Wang, and Y. Tang. Deep learning face attributes in the wild. In *Proceedings of International Conference on Computer Vision (ICCV)*, Dec. 2015.
6. A. Radford, L. Metz, and S. Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. *CvPR*, abs/1511.06434, 2015.
7. Z. Weihs. Deep vector.

Recurrent Neural Network for Politician's Text Generation

Tomer Vainstein

21329 Deep Learning, Spring 2020

Abstract

- Politicians often write posts (short texts) to express their agenda and promote themselves in modern social media platforms such as Facebook and Twitter.
- By scraping thousands of Facebook posts from several key politicians in Israel, and using their texts as a corpus, I will attempt to train a neural network that takes sequences of texts (e.g. sentences) and predicts the next word after each sequence, resulting in a model which can generate full sentences by learning from a corpus of politicians' posts.
- Previous work showed good results under domain-specific corpus such as Books from Project Gutenberg, Tweets of politicians and News Headlines.

Model

- The chosen model for the project was based on LSTM network (variation of RNN).
- Each model I experimented with contained an Embedding layer to convert words into vectors of fixed dimension, and an Activation FC layer (softmax).
- I tried different models containing Single/Dual/Bidirectional LSTM layer of 100/200/400 units.
- After the N hidden layers I tried to apply a Dropout of 10-25% to prevent overfitting.

Conclusions

I learned that generating "good" texts is a common result of text generator network implementation using LSTM, and modern solutions such as Transformers (GPT-2, BERT) are showing more realistic results.

Evaluation

- The model was compiled to calculate categorical cross entropy as loss function with Adam optimizer and learning rate of 0.0001.
- The model converged after 250 epochs, with 70% accuracy.

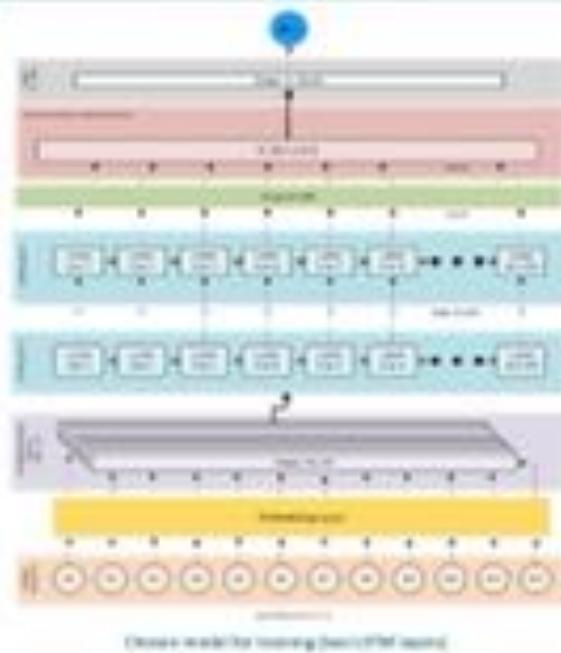


Qualitative Results



Dataset

- The dataset consist of 3,612 Facebook posts from 7 Israeli politicians which I scraped using a web-crawler I wrote.
- After preprocessing, the training data had 88,857 sequences of texts, each had a label (predicting next word).
- There are 24,248 unique words/tokens which indicates the number of classes/labels for classification.



References

- Shawn居士, Beginner's Guide to Text Generation using LSTM, August 2016.
- Aditya Chitta, Simple Text Generation (how to teach a machine to "speak"), Medium, April 2018.
- Zhang Wei, Recurrent Neural Networks: A Practical Guide, Medium, October 2018.

Thank you, and...

GOOD LUCK !